

# 3: The Processor (Single Cycle)

#### Acknowledgement:

- Text and companion slides by Harris and Harris
- Text by Patterson and Hennessey and companion slides by Mary Jane Irwin

#### Note:

 Not all slides will be covered in the lecture. The rest are left as a self-learning exercise.

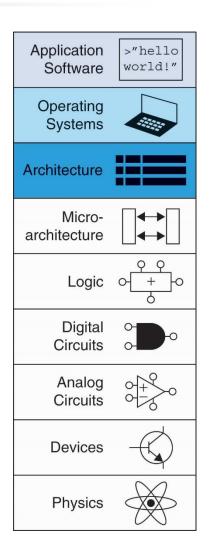


- ARM architecture overview
- ARM7 (ARMv3) instruction encoding
- Microarchitecture : Datapath design
- Microarchitecture : Controller design
- Extensions for other formats / instructions



### Recap: Architecture vs Microarchitecture

- Architecture: programmer's view of computer
  - Defined by instructions & operand locations
  - Assembly language: human-readable format of instructions
  - Machine language: computer-readable format (1's and 0's)
  - Assembly language -> Machine language conversion is done by the assembler
    - one to one correspondence (except for pseudo-instructions)
- Microarchitecture: how to implement an architecture in hardware



### Why Learn ARM?

- Developed in the 1980's by Advanced RISC Machines now called ARM Holdings
- More than 10 billion ARM processors sold/year
- Almost all cell phones and tablets have multiple ARM processors
- Well over 75% of humans use products with an ARM processor
- Used in servers, cameras, robots, cars, pinball machines,... etc.
- ARM provides
  - architecture licenses
    - Companies design their own microarchitectures [and SoCs using these microarchitectures]
    - e.g., Qualcomm Kryo [Snapdragon SoC], Apple Avalanche, Blizzard [A15 SoC]

as well as

- microarchitecture licenses
  - Companies use microarchitectures from ARM [and SoCs using these microarchitectures]
  - e.g., Cortex M3 [LPC 1769 SoC], Cortex-A710 [MediaTek Helio Dimensity 9000 SoC]

### **ARM Families and Architectures**

- Architectures evolve. Complete list of families and architectures at <a href="https://en.wikipedia.org/wiki/List of ARM microarchitectures">https://en.wikipedia.org/wiki/List of ARM microarchitectures</a>
- CG2028/EE2028/EE2024 was based on ARMv7-M\* ISA
  - Newer (beware of legal issues if you try to implement)
  - Executes Thumb (16-bit instructions, 32 bit data) and Thumb 2 (32-bit)
  - Thumb an encoding which allows better code density (with some limitations such as limited number of usable registers)
  - Thumb 2 32 bit instructions to extend Thumb functionality
  - ARMv7-M = Microcontroller; ARMv7-A = Application (phones etc.)
- In this module, we will be building an ARMv3 ISA based processor
  - Original 32-bit instruction set, implemented in ARM7 microarchitecture
  - Older (~1994, so you are unlikely to run into legal issues. DO NOT put it up on public repos though, to be on the safe side)
  - Somewhat similar to ARMv7-M at assembly level, though very different at machine level



### ARM / RISC Design Principles

- Regularity supports design simplicity
  - Constant length instructions (and data) and consistent instruction format
  - Same number of operands (two sources and one destination)
  - Ease of encoding and handling in hardware
- Make the common case fast
  - ARM includes only simple, commonly used instructions
  - Hardware to decode and execute instructions kept simple, small, and fast
  - More complex instructions (that are less common) performed using multiple simple instructions
- Smaller is faster
  - ARM includes only a small number of registers
- Good design demands good compromises
  - Number of instruction formats kept small. Multiple instruction formats allow flexibility... but complicates hardware

### Architecture: Registers, Instruction Formats

- ARM has 16 general purpose registers, with some special functions or typical uses
- We will be looking at 3 instruction formats
  - Data-processing
    - ADD, SUB, CMP, AND, ORR, MOV (LSL, ROR etc. are variants of MOV)
  - Memory
    - LDR, STR, LDRB, STRB etc.
  - Branch
    - B, BL

Load Byte. To load a *Byte* from the location (LDR loads a *Word*)

Register Name	<b>Function / Typical Use</b>
R0	Argument / return value / temporary variable
R1-R3	Argument / temporary variables
R4-R11	Saved variables
R12	Temporary variable
R13 (SP)	Stack Pointer
R14 (LR)	Link Register
R15 (PC)	Program Counter

### **Architecture: Conditional Execution**

- Every instruction is conditionally executed based on N, Z, C, V flags.
   Think of flags as 1-bit registers (actually, they are bits within xPSR)
  - Example : ADDEQ R1, R2, R3 affects the value of R1 only when Z = 1
- Flags are set by instructions with suffix S
  - Example : ADDS affects flags, ADD doesn't
  - Exceptions : CMP, CMN, TST, TEQ which are used only to set flags (result is discarded)

cond	Mnemonic	Name	CondEx
0000	EQ	Equal	Z
0001	NE	Not equal	$ar{Z}$
0010	CS / HS	Carry set / Unsigned higher or same	С
0011	CC / LO	Carry clear / Unsigned lower	Ē
0100	MI	Minus / Negative	N
0101	PL	Plus / Positive of zero	$\overline{N}$
0110	VS	Overflow / Overflow set	V
0111	VC	No overflow / Overflow clear	$ar{V}$
1000	HI	Unsigned higher	Σ̄C
1001	LS	Unsigned lower or same	Z OR $\bar{\mathcal{C}}$
1010	GE	Signed greater than or equal	$\overline{N \oplus V}$
1011	LT	Signed less than	$N \oplus V$
1100	GT	Signed greater than	$\bar{Z}(\overline{N \oplus V})$
1101	LE	Signed less than or equal	$Z OR (N \oplus V)$
1110	AL (or none)	Always / unconditional	ignored

Interpretation based on SUBS/CMP

# **Data-processing Instruction Format**

### Operands

<OP Rd, Rn, Src2> OP => ADD, SUB, MOV,...

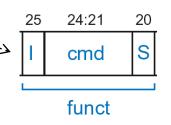
- Rn : first source register
- Src2: second source register or immediate
- Rd : destination register

#### **Data-processing**

	31:28	27:26	25:20	19:16	15:12	11:0
	cond	op	funct	Rn	Rd	Src2
•	4 bits	2 bits	6 bits ``	4 bits	4 bits	12 bits

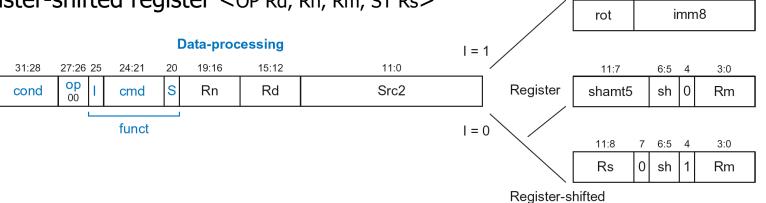
#### Control fields

- cond : specifies conditional execution
- op : the operation code or opcode
  - op = "00" for data-processing (DP) instructions
- funct : the function/operation to perform
  - funct is composed of cmd, I-bit, and S-bit



### **Data-processing Instruction Format**

- cmd: specifies the specific data-processing instruction <OP Rd, Rn, Src2>
  - cmd = "0100" for ADD; cmd = "0010" for SUB; ...
  - See page A3-9 of the ARM Architecture Reference Manual for the complete cmd code list
- I-bit
  - I = 0: Src2 is a register. I = 1: Src2 is an immediate
- S-bit: 1 if instruction sets condition flags
  - S = 0: SUB R0, R5, R7. S = 1: ADDS R8, R2, R4 or CMP R3, #10
- Src2 can be
  - Immediate <OP Rd, Rn, #Imm8\_rot> ST = Shift Type = LSL / LSR / ASR / ROR
  - Register with imm. shift <OP Rd, Rn, Rm, ST #shamt5>
  - Register-shifted register < OP Rd, Rn, Rm, ST Rs>



**Immediate** 

Register

11:8

7:0

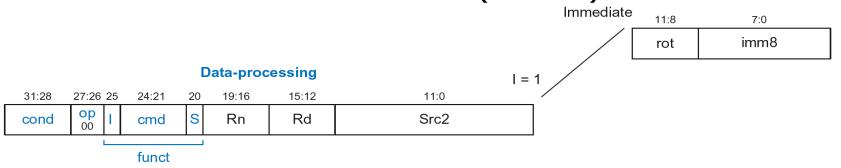
10



### DP Instruction with Immediate Src2

<OP Rd, Rn, #Imm8\_rot>

- Immediate encoded as
  - Imm8: 8-bit unsigned immediate
  - rot : 4-bit rotation value
- 32-bit constant is: Imm8 ROR ( $rot \times 2$ )



The assembly language programmer need not break down the 32-bit constant into imm8 and rot - the assembler will do this. If the 32-bit constant is such that it can't be encoded using imm8 and rot, the assembler will flag an error. For constants that fit within 8 bits, rot = 0.

Note: This is very different from ARM v7-M / Thumb where *Src2* / *Operand2* is **Imm8m**, formed by **left-shifting** an 8-bit value by **any number of bits**, or a bit pattern of one of the forms 0xXYXYXYXY, 0x00XY00XY or 0xXY00XY00



### DP Instruction with Immediate Src2: e.g.

<OP Rd, Rn, #Imm8\_rot>

### SUB R2, R3, #0xFF0

- cond = "1110" (14) for unconditional execution
- op = "00" (0) for data-processing instructions
- cmd = "0010" (2) for SUB
- Src2 is an immediate so I=1

• 
$$Rd = 2$$
,  $Rn = 3$ 

ROL by 4 => ROR by 
$$(32-4) = 28$$
  
 $rot \times 2 = 28 => rot = 14$ 

Imm8 must be rotated right by 28 to produce 0xFF0, so rot = 14

31:28	27:26	25	24:21	20	19:16	15:12	11:8	7:0
1110 <sub>2</sub>	002	1	00102	0	3	2	14	255
cond	ор	Ι	cmd	S	Rn	Rd	rot	imm8
1110	00	1	0010	0	0011	0010	1110	11111111

SUB R2, R3, #0xFF0 => 0xE2432EFF



# DP Instruction with Register *Src2*

<OP Rd, Rn, Rm, ST #shamt5>

Rm : the second source operand

shamt5: the amount Rm is shifted

sh : the type of shift

funct

Shift Type	sh
LSL	00
LSR	01
ASR	10
ROR	11

#### **Data-processing**

31:28	27:26 25	24:21	20	19:16	15:12	11:0		11:7	6:5	4	
cond	op 00	cmd	S	Rn	Rd	Src2	Register	shamt5	sh	0	
							,				

I = 0

3:0 **Rm** 



### DP Instruction with Register Src2: e.g.

<OP Rd, Rn, Rm, ST #shamt5>

- ORR R8, R1, R3, LSR #2
  - Operation: R8 = R1 OR (R3 >> 2)
  - cond = "1110" (14) for unconditional execution
  - op = "00" (0) for data-processing instructions
  - cmd = "1100" (12) for ORR
  - Src2 is a register so I=0
  - Rd = 8, Rn = 1, Rm = 3
  - shamt5 = 2, sh = "01" (LSR)

#### **Data-processing**

31:28	27:26 25	24:21	20	19:16	15:12	11:0		11:7	6:5 4	3:0
cond	op 00	cmd	S	Rn	Rd	Src2	Register	shamt5	sh 0	Rm
		funct				1=	0			

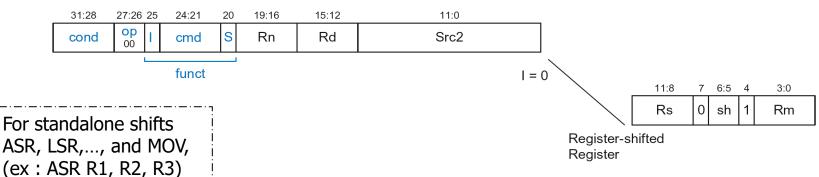
# DP with Register-shifted Register Src2: e.g.

<OP Rd, Rn, Rm, ST Rs>

- EOR R8, R9, R10, ROR R12
  - Operation: R8 = R9 XOR (R10 ROR R11)
  - cond = "1110" (14) for unconditional execution
  - op = "00" (0) for data-processing instructions
  - cmd = "0001" (1) for EOR
  - Src2 is a register so I=0
  - *Rd* = 8, *Rn* = 9, *Rm* = 10, *Rs* = 12
  - sh = "11" (ROR)

cmd="1101" and Rn=0

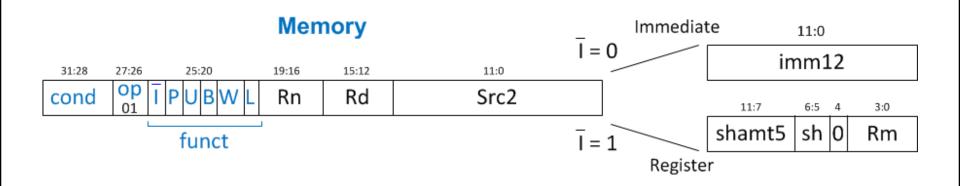
#### **Data-processing**





### **Memory Instruction Format**

- Encodes LDR, STR, LDRB, STRB
  - op : "01"
  - funct : 6 control bits
  - Rn : base register
  - Rd : destination (load), source (store)
  - Src2 : offset
    - immediate <OP Rd, [Rn, #Imm12]>
    - register (optionally shifted) <OP Rd, [Rn, Rm, ST #shamt5]>



## **Memory Instruction Format**

### funct

■ I : Immediate bar

P : Preindex

U : Add

B : Byte

W: Writeback

L : Load

P	W	Indexing Mode
0	1	Not supported
0	0	Postindex
1	0	Offset
1	1	Preindex

L	В	Instruction
0	0	STR
0	1	STRB
1	0	LDR
1	1	LDRB

LDR R3, [R5, #16]!

- R3 = mem[R5+16]
- R5 = R5 + 16

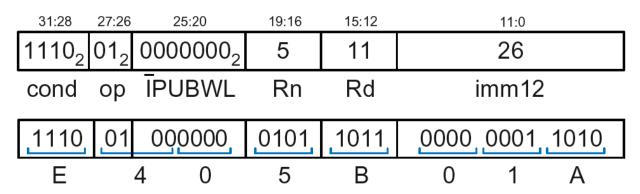
Value	$\overline{I}$	U
0	Immediate offset in <i>Src2</i>	Subtract offset from base
1	Register offset in Src2	Add offset to base

### Memory Instr. with Immediate Src2: e.g.

<OP Rd, [Rn, #Imm12]>

- STR R11, [R5], #-26
  - Operation: mem[R5] <= R11; R5 = R5 26</p>
  - cond = "1110" (14) for unconditional execution
  - op = "01" (1) for memory instruction
  - funct = "0000000" (0)
  - I = 0 (immediate offset), P = 0 (postindex), U = 0 (subtract), B = 0 (store word), W = 0 (postindex), L = 0 (store)
  - *Rd* = 11, *Rn* = 5, *Imm12* = 26

#### **Field Values**

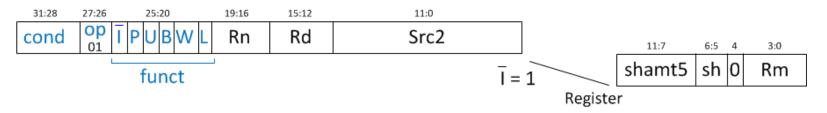


### Memory Instr. with Scaled Reg. Src2: e.g.

STR R9, [R1, R3, LSL #2]

- <OP Rd, [Rn, Rm, ST #shamt5]>
- Operation: mem[R1 + (R3 << 2)] <= R9</li>
- cond = "1110" (14) for unconditional execution
- op = "01" (1) for memory instruction
- funct = "111000" (56)
- I = 1 (register offset), P = 1 (offset indexing),
   U = 1 (add), B = 0 (store word), W = 0 (offset indexing),
   L = 0 (store)
- Rd = 9, Rn = 1, Rm = 3, shamt = 2, sh = "00" (LSL)

#### Memory



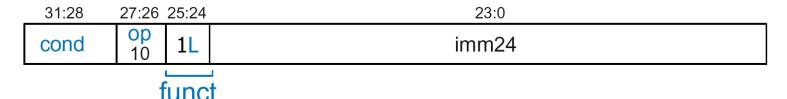
### **Branch Instruction Format**

Encodes B and BL

<OP Imm24>

- op = "10"
- funct = "1L": L = 1 for BL, L = 0 for B
- Imm24: 24-bit immediate encoding Branch Target Address (BTA)
  - BTA = Next PC when branch taken
  - Imm24 = # of words BTA is away from current PC+8
  - Note: (# of words) = (# of bytes / 4) = (# of bytes >> 2)

#### **Branch**



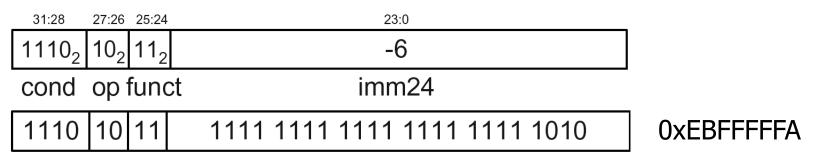


### Branch Instruction: e.g.

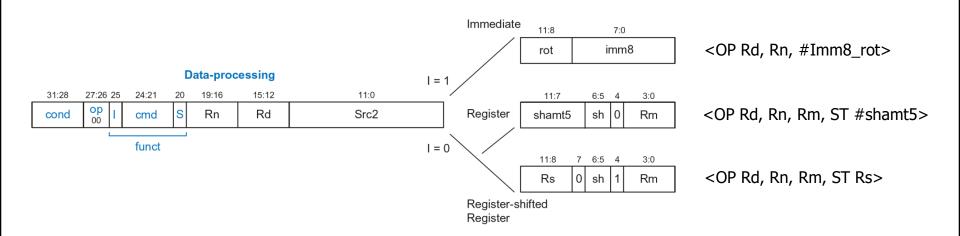
```
0x8040 TEST: LDRB R5, [R0, R3] \leftarrow BTA
0x8044
             STRB R5, [R1, R3]
0x8048
                 R3, R3, #1
             ADD
0x804C
             MOV PC, LR
0x8050
                               ← PC
             BL
                  TEST
0x8054
             LDR R3, [R1], #4
0x8058
                  R4, R3, #9 ← PC+8
             SUB
```

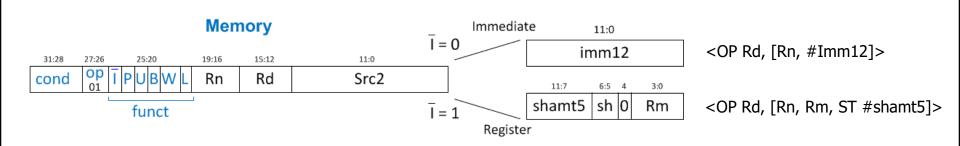
- PC = 0x8050
- PC + 8 = 0x8058
- BTA =  $0 \times 8040$
- TEST label is 6 instructions
  before PC+8, so Imm24 = -6
  Offset = 0x8040-0x8058 =
  - -0x18 bytes = -0x6 words

#### **Field Values**



### **Review: Instruction Formats**



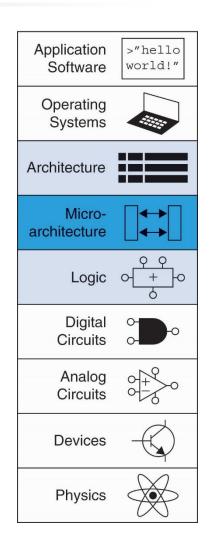




<OP Imm24>



- Processor
  - Datapath: functional blocks
  - Control: control signals
- Basic styles
  - Single-cycle: Each instruction executes in a single cycle
  - Multicycle: Each instruction is broken up into series of shorter steps
  - Pipelined: Each instruction broken up into series of steps & multiple instructions execute at once



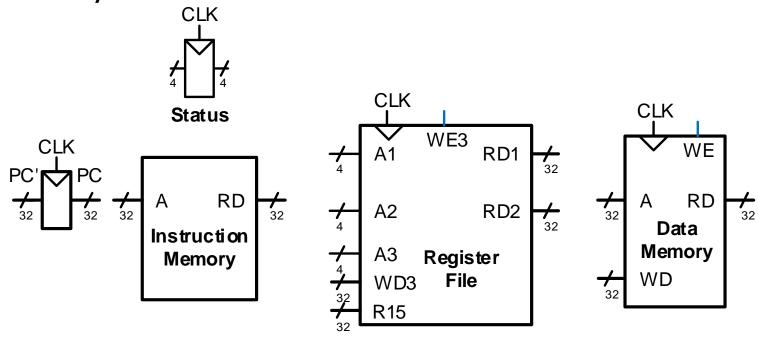


### **ARM Processor: Getting Started**

- We will get started with a subset of ARM instructions
- Data-processing instructions
  - ADD, SUB, AND, ORR
  - with register and immediate Src2, but no shifts
- Memory instructions
  - LDR, STR
  - with positive immediate offset
- Branch instructions
  - B

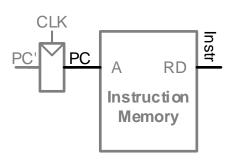
### **Architectural State Elements**

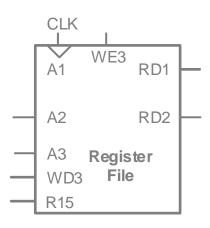
- Determines everything about a processor
- Architectural state
  - 16 registers (including PC) -> Register File (RF)
  - Status register
- Memory

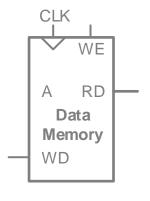




STEP 1: Fetch instruction

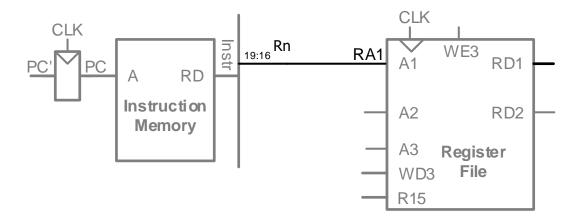




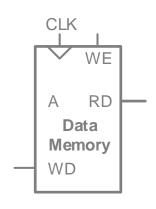


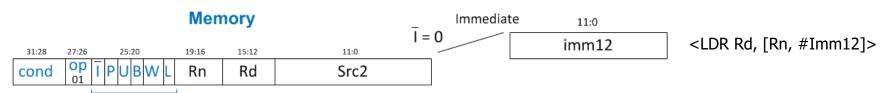
# Single-Cycle Datapath: LDR Reg Read

STEP 2: Read source operands from RF



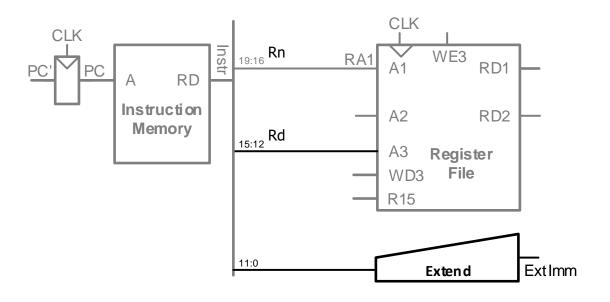
funct



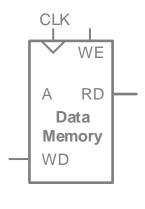


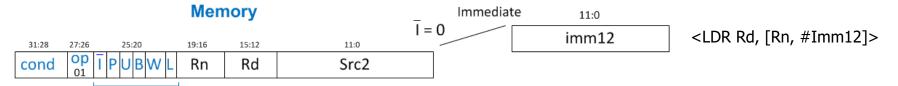
# Single-Cycle Datapath: LDR Immed.

#### STEP 3: Extend the immediate



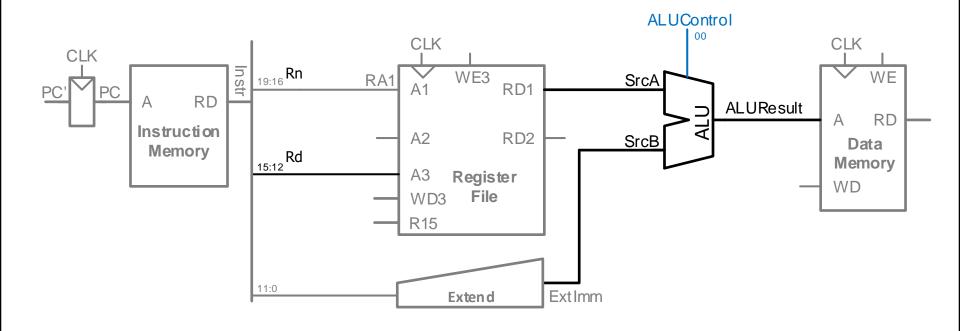
funct

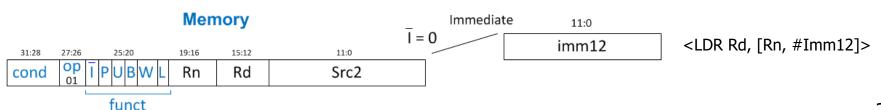




### Single-Cycle Datapath: LDR Address

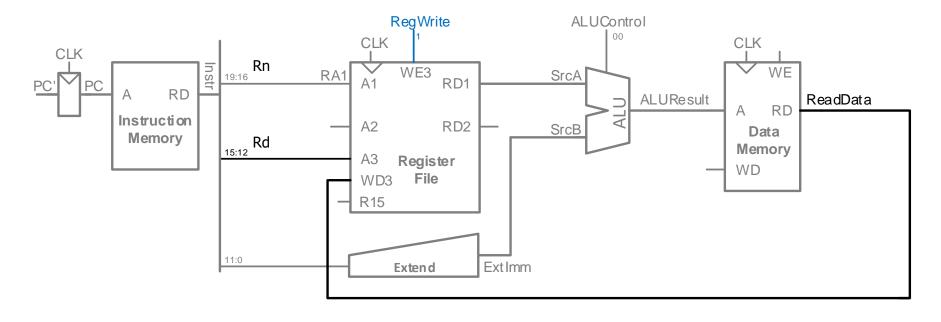
STEP 4: Compute the memory address

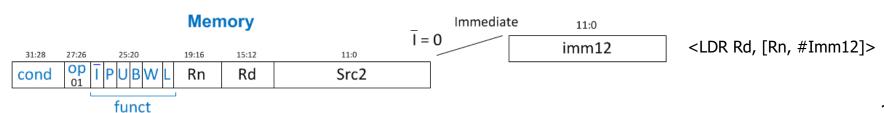




# Single-Cycle Datapath: LDR Mem Read

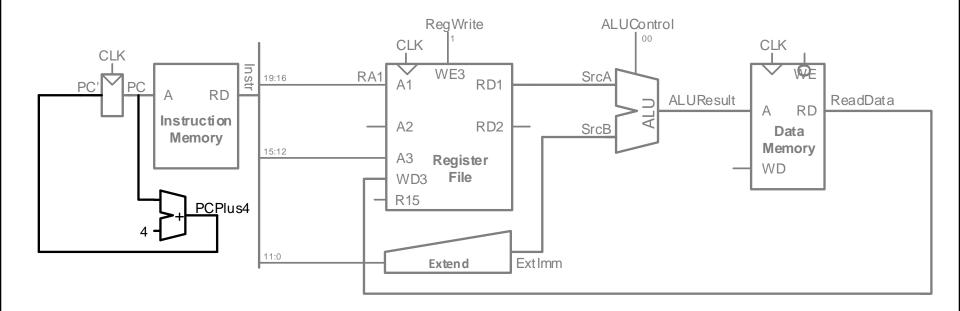
 STEP 5: Read data from memory and write it back to register file





### Single-Cycle Datapath: PC Increment

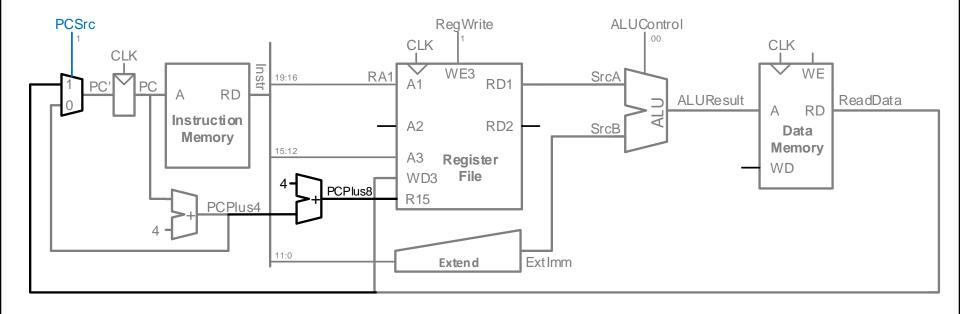
STEP 6: Determine address of next instruction





### Single-Cycle Datapath: PC Source/Dest

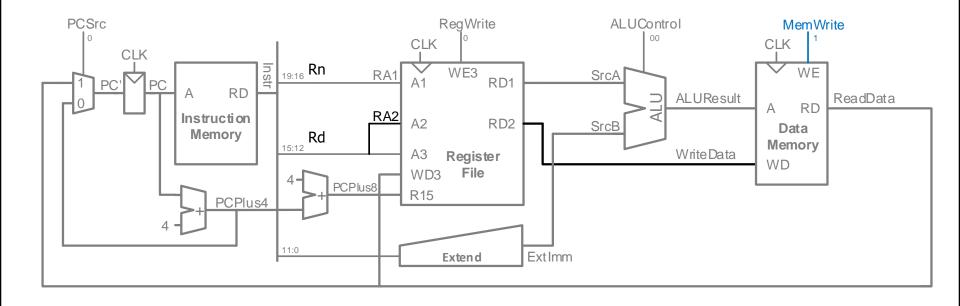
- PC can be source/destination of instruction
  - Source: R15 (PC+8) available in Register File
  - Destination: Be able to write result to PC

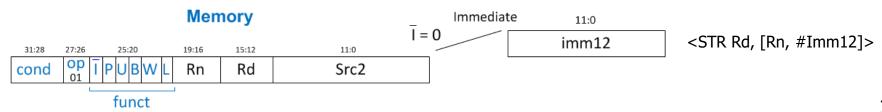


R15 is not "stored" in the register file. If A1 or A2 is 15, RD1 or RD2 gets the value of PC after being passed through two adders (i.e., PC+8)

# Single-Cycle Datapath: STR

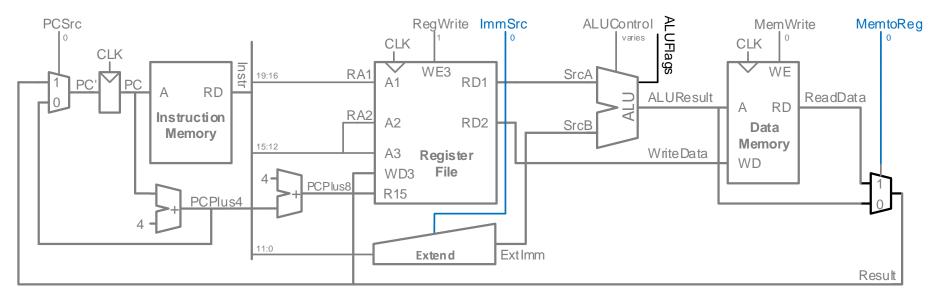
- Write data in RD to memory
  - Note that Rd is a source operand for STR

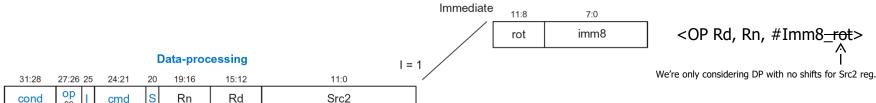




## Single-Cycle Datapath: Data Processing

- With immediate Src2
  - Read from Rn and Imm8 (ImmSrc chooses the zero-extended Imm8 instead of Imm12)
  - Write ALUResult to register file (Rd)

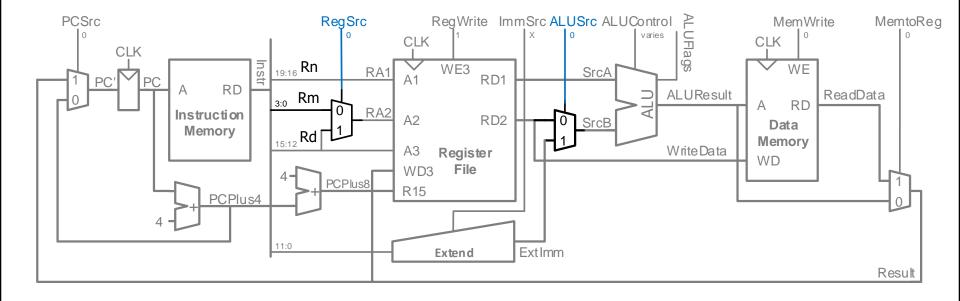


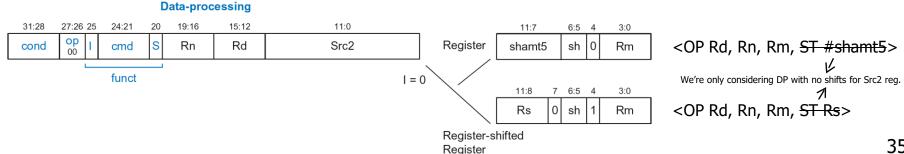


34

## Single-Cycle Datapath: Data Processing

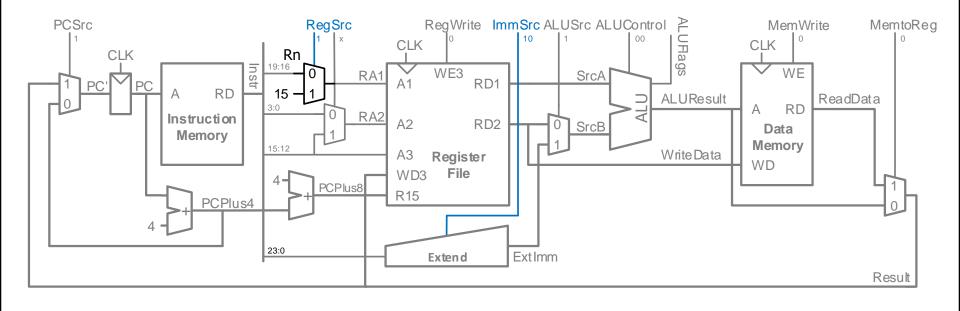
- With register *Src2* 
  - Read from *Rn* and *Rm* (instead of *Imm8* )



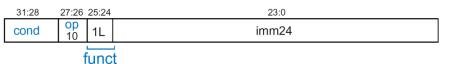


## Single-Cycle Datapath: Branch

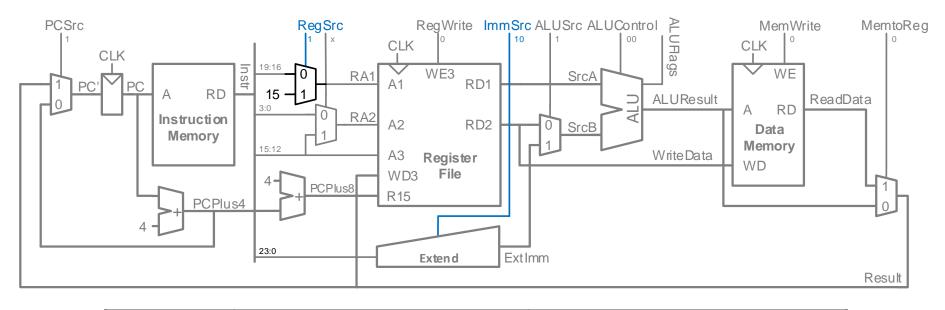
- Calculate branch target address
  - $\blacksquare BTA = (ExtImm) + (PC + 8)$
  - ExtImm = (sign-extended Imm24) << 2</p>



#### **Branch**



# Single-Cycle Datapath: ExtImm



ImmSrc <sub>1:0</sub>	ExtImm	Description
00	( 24{`0'}, Instr <sub>7:0</sub> )	Zero-extended Imm8
01	( 20{`0'}, Instr <sub>11:0</sub> )	Zero-extended Imm12
10	( 6{Instr <sub>23</sub> }, Instr <sub>23:0</sub> , 2{'0'} )	(Sign-extended Imm24)<<2

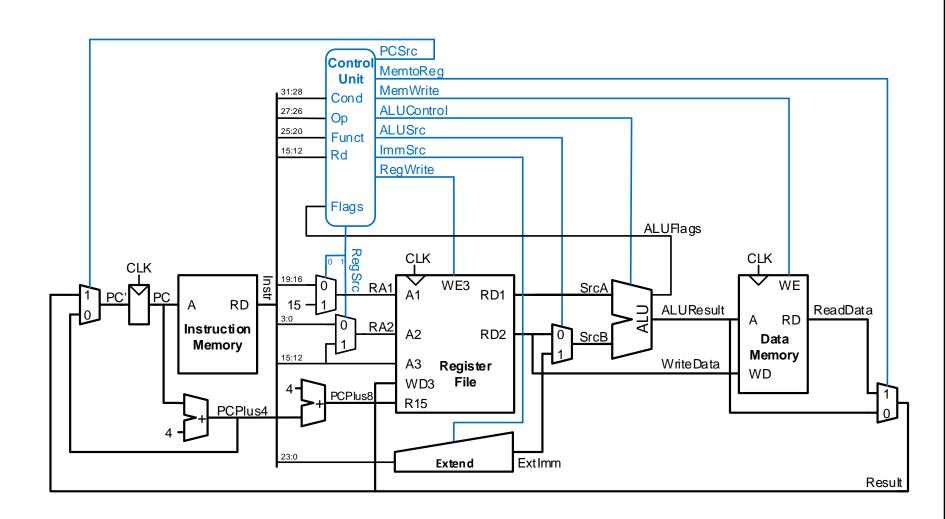
How would you extend LDR and STR to support negative offsets?. Do you need sign extension?

Q: Do you need a shifter to shift Imm24?

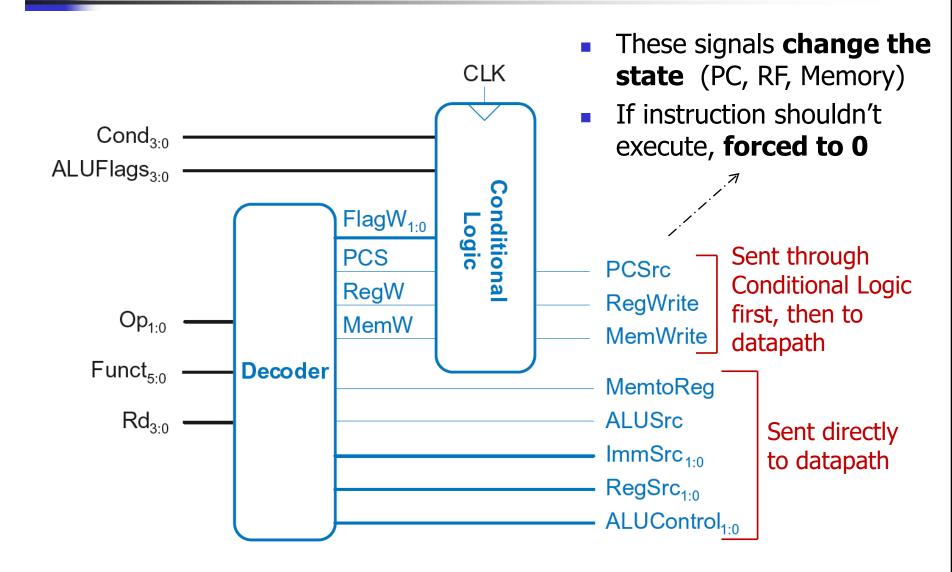
Ans: NO!. A shifter is required only if the shift amount is variable.

Shifting by a constant requires no logic gates -> can be accomplished by appropriate wiring!

# Single-Cycle ARM Processor with Control

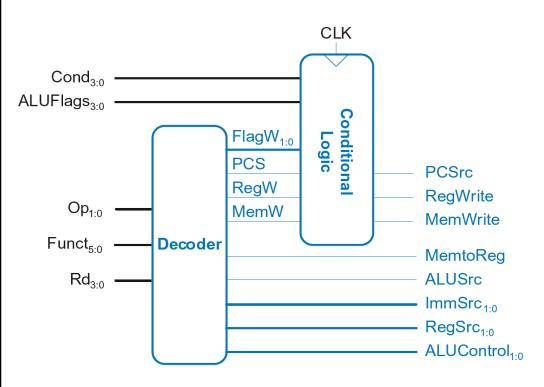


# Single-Cycle Control





# Single-Cycle Control

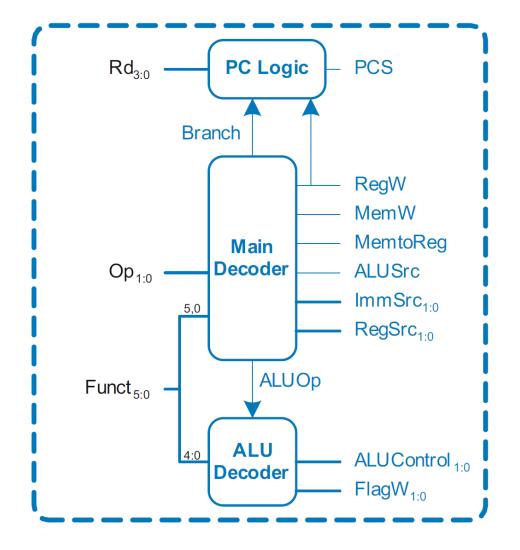


- FlagW<sub>1:0</sub>: Flag Write signal, asserted when ALUFlags should be saved (i.e., on instruction with S=1)
- ADD, SUB update all flags (NZCV)
- AND, ORR only update NZ flags
- So, two bits needed:
  - FlagW<sub>1</sub> = 1: NZ saved (ALUFlags<sub>3:2</sub> saved)
  - FlagW<sub>0</sub> = 1: CV saved (ALUFlags<sub>1:0</sub> saved)



### Submodules

- Main Decoder
- ALU Decoder
- PC Logic





## Control Unit: Main Decoder

Op	Funct <sub>5</sub>	Funct <sub>o</sub>	Туре	Branch	MemtoReg	MemW	ALUSrc	ImmSrc	RegW	RegSrc	dOnTV
00	0	X	DP Reg	0	0	0	0	XX	1	00	1
00	1	X	DP Imm	0	0	0	1	00	1	X0	1
01	X	0	STR	0	X	1	1	01	0	10	0
01	X	1	LDR	0	1	0	1	01	1	X0	0
10	X	X	В	1	0	0	1	10	0	X1	0

To make full sense of this slide, read it together with slides 22, 38 and 41

### Control Unit: ALU Decoder

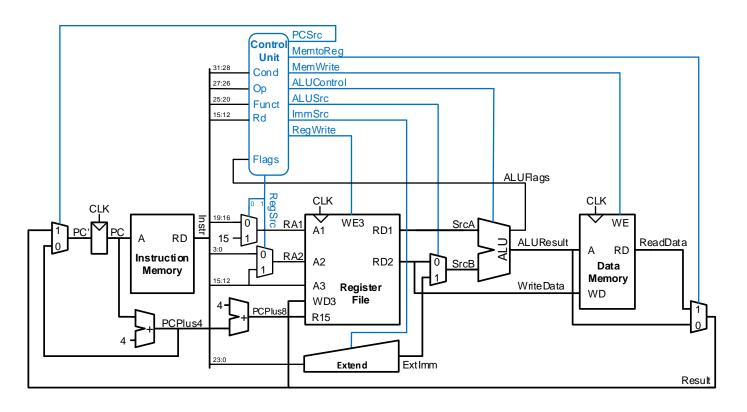
ALUOp	Funct <sub>4:1</sub> (cmd)	Funct <sub>0</sub> (S)	Туре	ALUControl <sub>1:0</sub>	FlagW <sub>1:0</sub>
0	X	X	Not DP	00	00
1	0100	0	ADD	00	00
		1			11
	0010	0	SUB	01	00
		1			11
	0000	0	AND	10	00
		1			10
	1100	0	ORR	11	00
		1			10

- $FlagW_1 = 1$ : NZ ( $Flags_{3:2}$ ) should be saved
- FlagW<sub>0</sub> = 1: CV (Flags<sub>1:0</sub>) should be saved\*

<sup>\*</sup> In reality, AND & ORR affects C as well. Omitted here for simplicity as Src2 does not supports shifts

# Single-Cycle Control: PC Logic

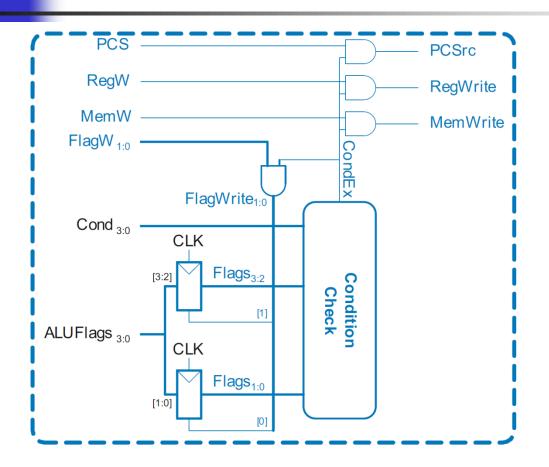
• PCS = 1 if PC is written by an instruction or branch (B): PCS = ((Rd == 15) & RegW) | Branch



If instruction is executed: PCSrc = PCS

Else PCSrc = 0 (i.e., PC = PC + 4)

# Single-Cycle Control: Conditional Logic



Recall:

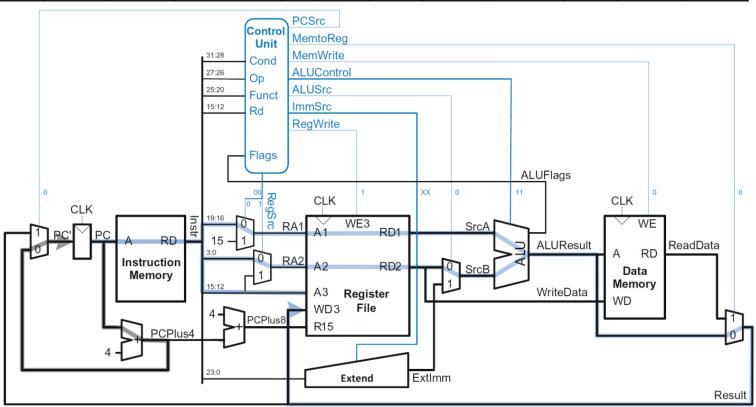
ADD, SUB update all flags AND, OR update NZ only. So Flags status register has two write enables.

Flags<sub>3:0</sub> =  $\{N,Z,C,V\}$ 

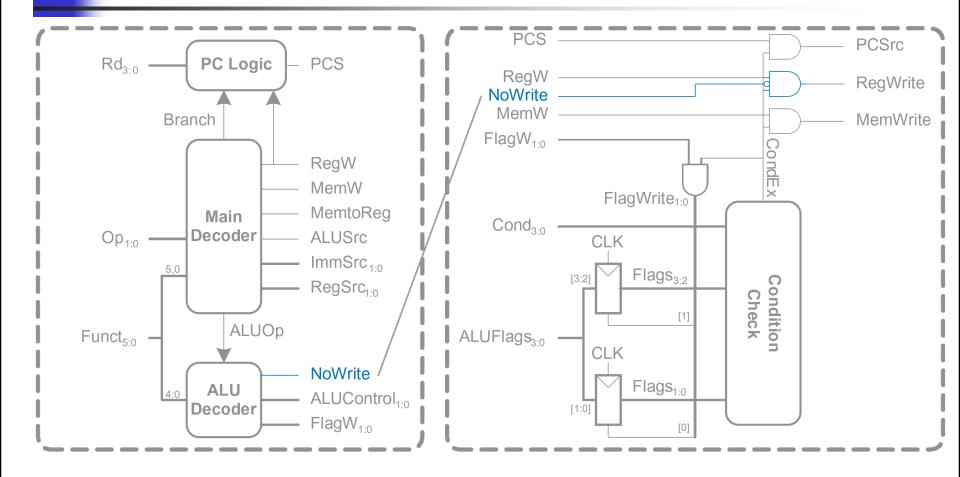
Depending on condition mnemonic ( $Cond_{3:0}$ ) and condition flags ( $Flags_{3:0}$ ) the instruction is executed (CondEx = 1)

# Example: ORR

Op	Funct <sub>s</sub>	Funct <sub>o</sub>	Type	Branch	MemtoReg	MemW	ALUSrc	ImmSrc	RegW	RegSrc	ALUOp
00	0	X	DP Reg	0	0	0	0	XX	1	00	1



# **Extended Functionality: CMP**

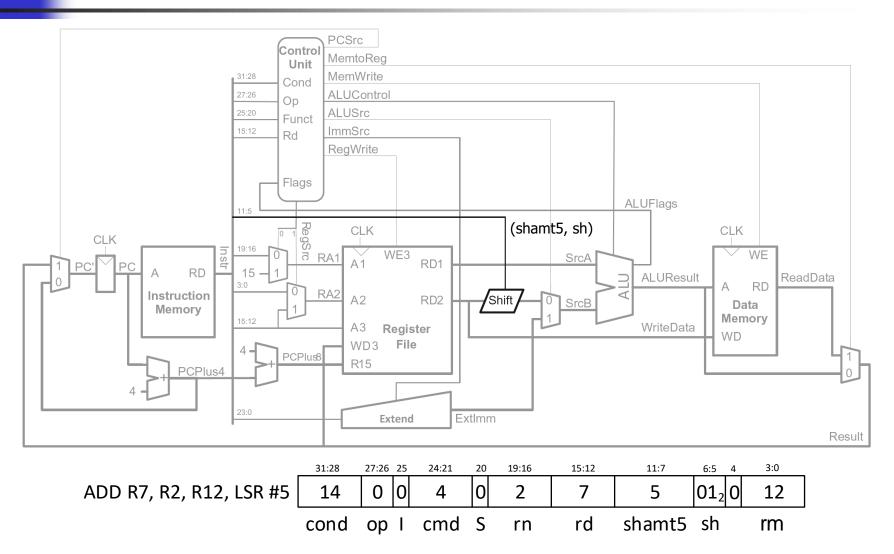


- Recall: CMP does subtraction, sets the flags and discards the result
- No change to datapath needed for implementation!

# **Extended Functionality: CMP**

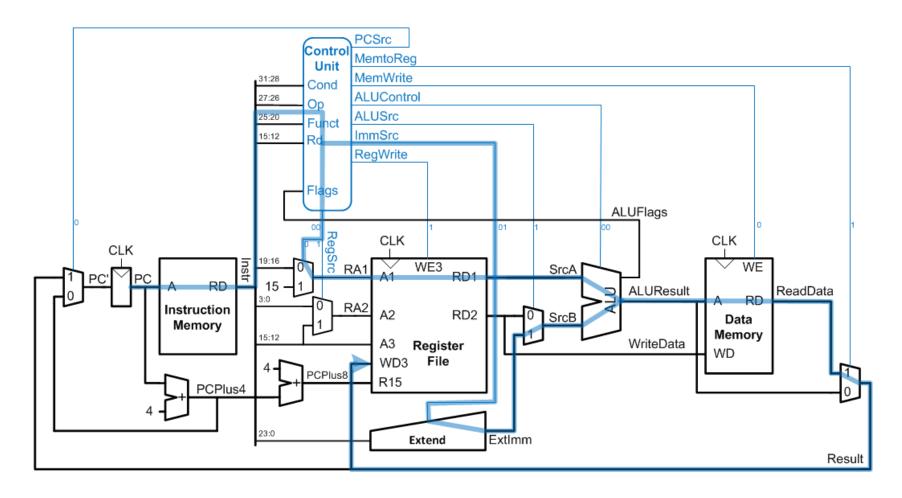
ALUOp	Funct <sub>4:1</sub> (cmd)	Funct <sub>0</sub> (S)	Туре	ALUControl <sub>1:0</sub>	FlagW <sub>1:0</sub>	NoWrite
0	X	X	Not DP	00	00	0
1	0100	0	ADD	00	00	0
		1			11	0
	0010	0	SUB	01	00	0
		1			11	0
	0000	0	AND	10	00	0
		1			10	0
	1100	0	ORR	11	00	0
		1			10	0
	1010	1	СМР	01	11	1

# Extended Functionality: Shifted Register



No change to the controller!

# Single-Cycle Performance



T<sub>C</sub> (and hence, performance) limited by critical path (LDR)

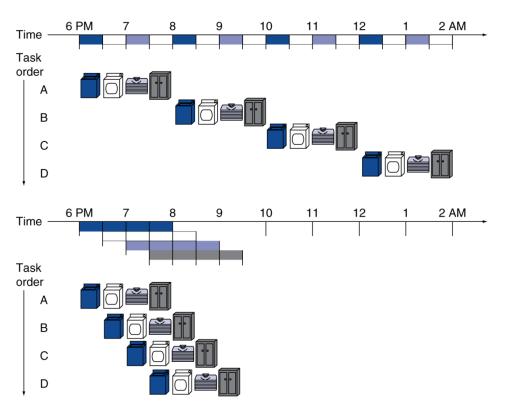


# Single Cycle Design Summary

- Single-cycle fetch, decode and execute each instructions in one clock cycle
  - (+) simple
  - (–) no datapath resource can be used more than once per instruction, so some must be duplicated
    - separate memories for instruction and data
    - 3 adders/ALUs
  - (–) cycle time limited by longest instruction (LDR)
- How Can We Make It Faster?
  - Pipelining (Chapter 6)
  - Superscalar (Chapter 7)
  - Multiprocessor systems (Chapter 8)

# Pipelining Analogy. Stay Tuned for More!

- Pipelined laundry: overlapping execution
  - Parallelism improves performance



#### Four loads:

- Speedup= 8/3.5 = 2.3
- Non-stop:
  - Speedup= 2n/(0.5n + 1.5)≈ 4 = # of stages

# Recap: Endianness

- Each data byte has unique address
- 32-bit word = 4 bytes, so word address increments by 4
- How to number bytes within a word?
  - Little-endian: byte numbers start at the little (least significant) end
  - Big-endian: byte numbers start at the big (most significant) end
  - It doesn't really matter which addressing type is used
    - except when two systems share data

