



# 5 : The Processor (Pipelined)

---

Rajesh Panicker,  
NUS

CG3207

Acknowledgement :

- Text by Harris and Harris and companion materials
- Text by Patterson & Hennessy and companion materials by Mary Jane Irwin, PSU

# Instruction Critical Paths


- What is the clock cycle time assuming negligible delays for muxes, control unit, sign extend, PC read/write, wires, setup and hold times except:
  - Instruction and Data Memory (200 ps)
  - ALU (120 ps)
  - Adders (75 ps)
  - Register File access (reads – 100 ps, writes – 60 ps)

Instr.	I Mem	Reg Rd	ALU Op	D Mem	Reg Wr	Total
DP	200	100	120		60	480
LDR	200	100	120	200	60	680
STR	200	100	120	200		620
B	200	100	120			420

Note : Addition for PC+8 is done in parallel with lmem access and takes even less time. So adders are not in the critical path

# Single Cycle Disadvantages

- Uses the clock cycle inefficiently – the clock cycle must be timed to accommodate the **slowest** instruction
  - especially problematic for more complex instructions like floating point multiply



- May be wasteful of area since some functional units (e.g., adders) must be duplicated since they can not be shared during a clock cycle




# How Can We Make It Faster?

---

- Start fetching and executing the next instruction before the current one has completed
  - **Pipelining** – (all?) modern processors are pipelined for performance
- Remember *the* performance equation
$$\text{CPU time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$
- Under *ideal* conditions and with a large number of instructions, the speedup from pipelining is approximately equal to the number of pipe stages
  - A five stage pipeline is nearly five times faster because the CCT is nearly five times faster

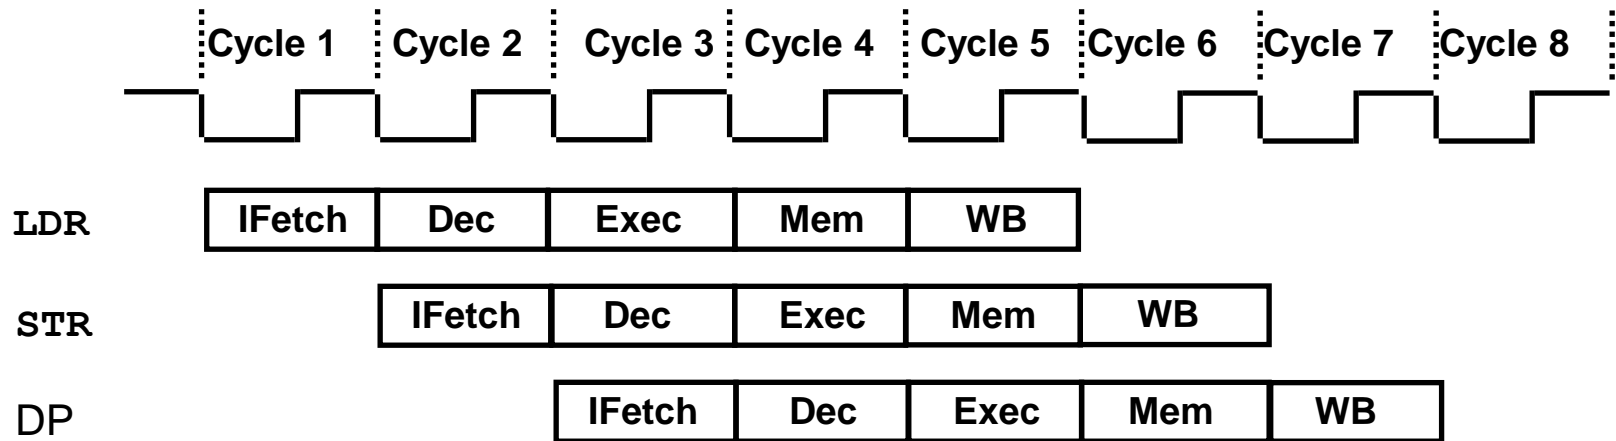
# The Five Stages of Load Instruction



- Fetch: Instruction Fetch and Update PC
- Dec: Registers Fetch and Instruction Decode
- Exec: Execute DP-type; calculate memory address
- Mem: Read/write the data from/to the Data Memory
- WB: Write the result data into the register file

# A Pipelined ARM Processor

- Start the **next** instruction before the current one has completed
  - improves **throughput** - total amount of work done in a given time
  - instruction **latency** (execution time, delay time, response time - time from the start of an instruction to its completion) is *not* reduced

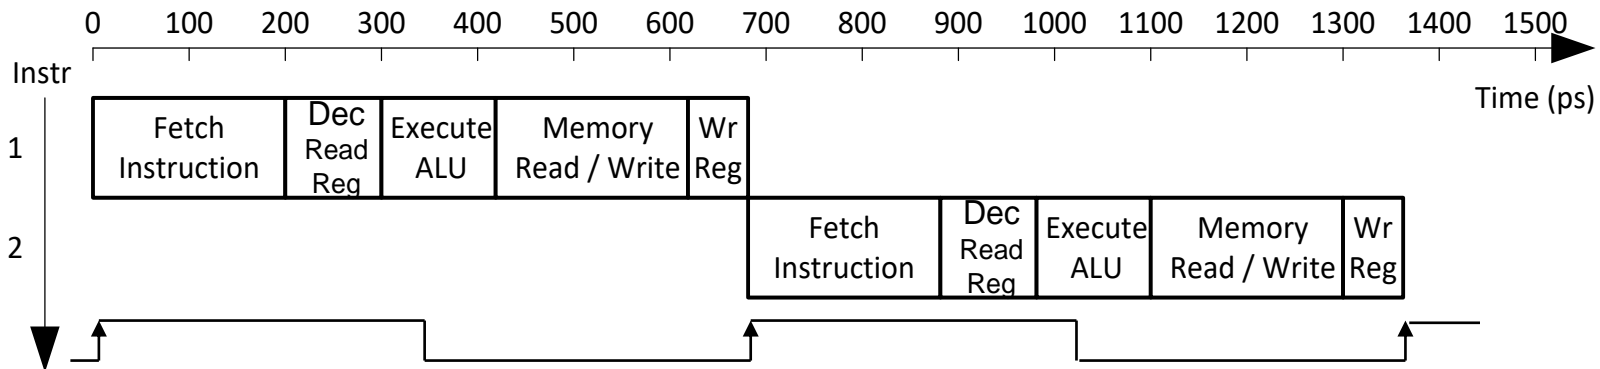


- clock cycle (pipeline stage time) is limited by the **slowest** stage
  - for some stages don't need the whole clock cycle (e.g., WB)
  - for some instructions, some stages are **wasted** cycles (i.e., nothing is done during that cycle for that instruction)

# Single Cycle versus Pipeline

## Single-Cycle

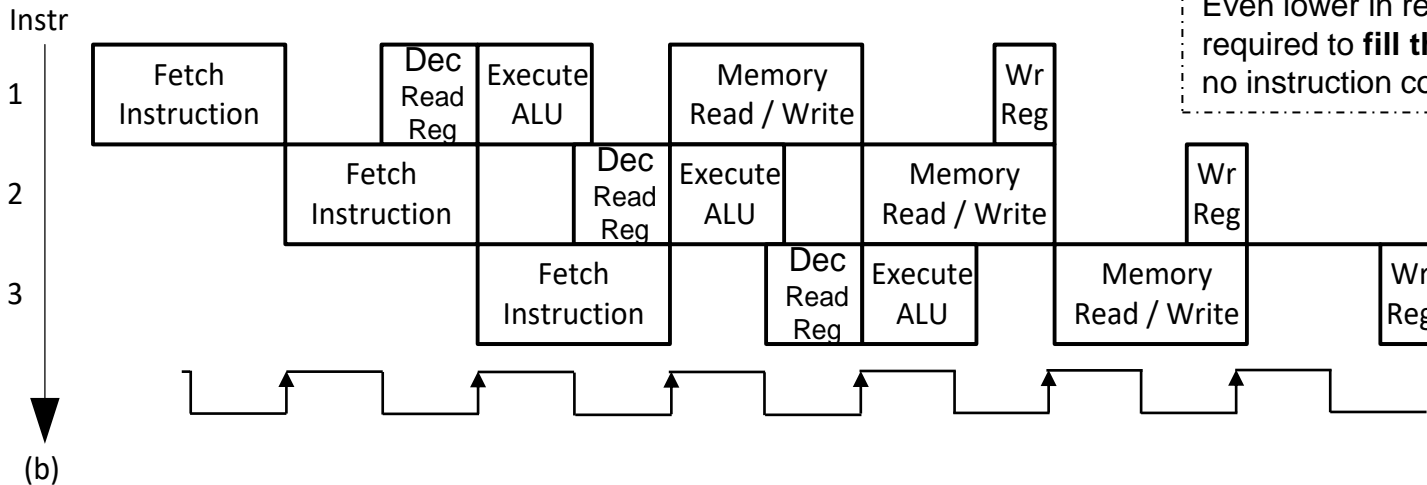
#instructions/sec =  $1/680\text{ps} = 1.47$  Billion



Speedup =  $5/1.47 = 3.4$   
Pipeline stages not balanced =>  
Speedup less than ideal value of 5.

## Pipelined

#instructions/sec =  $1/200\text{ps} = 5$  Billion

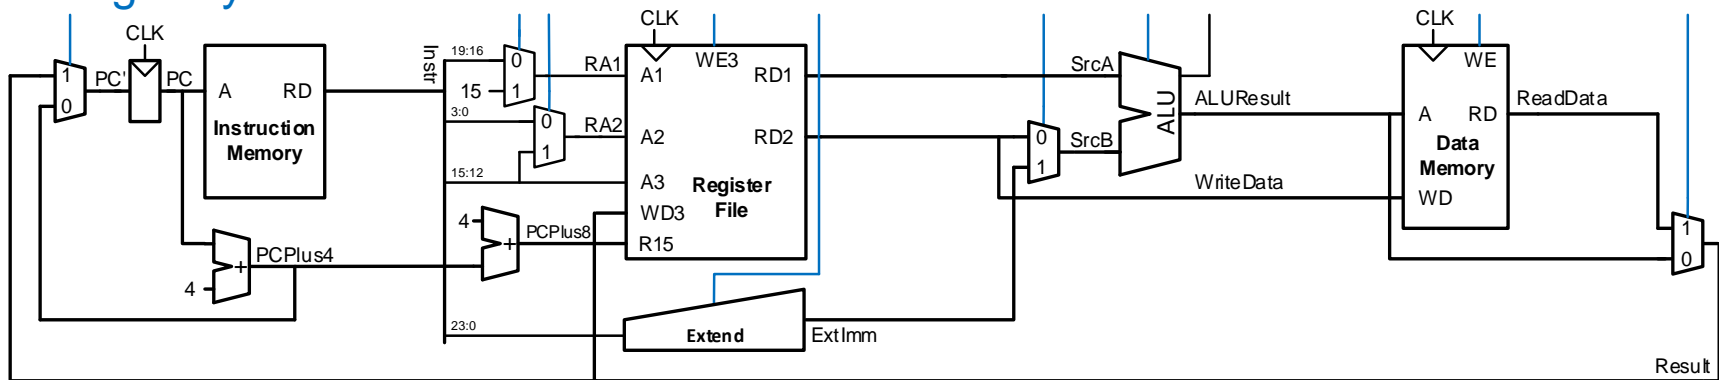


Even lower in reality - some cycles required to **fill the pipeline** during which no instruction completes execution

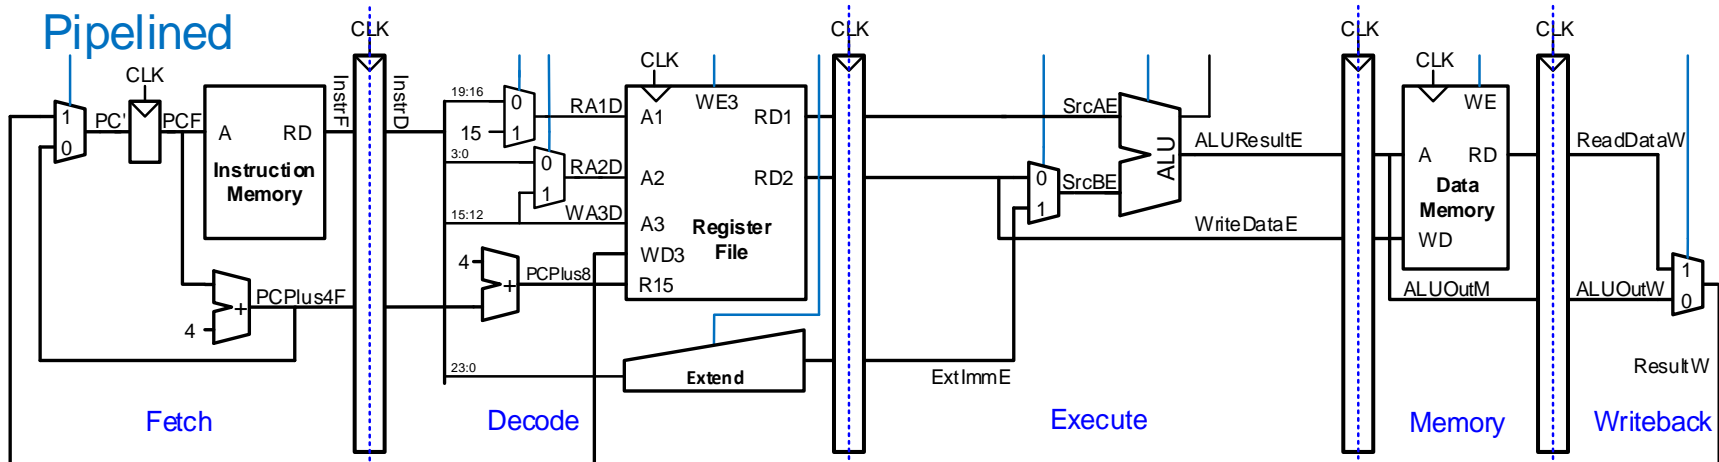
# ARM Pipeline Datapath Additions/Mods

- State registers between each pipeline stage to isolate them

## Single-Cycle

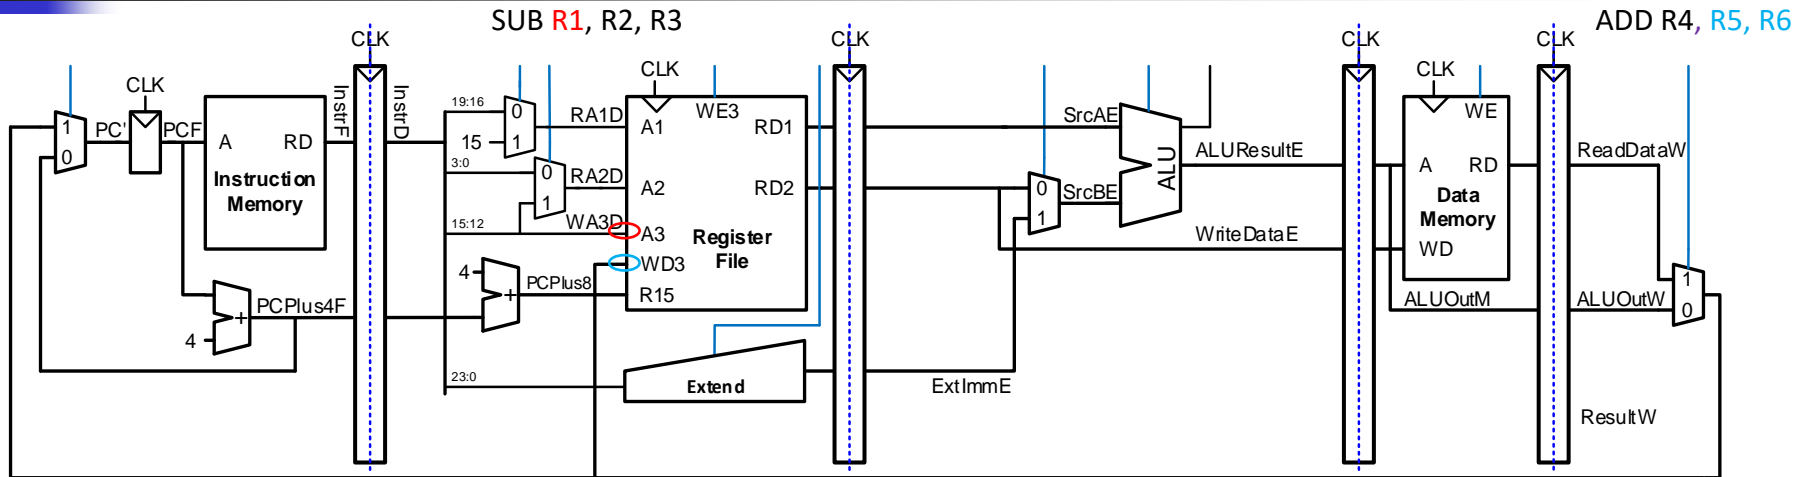


## Pipelined



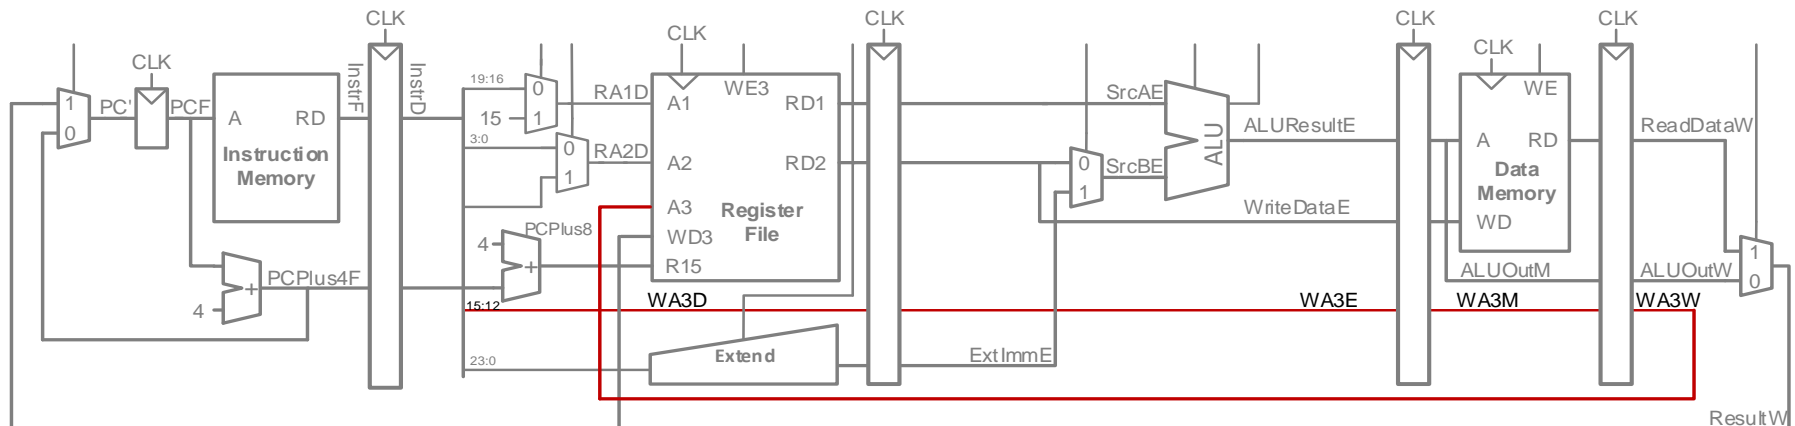


# Corrected Pipelined Datapath



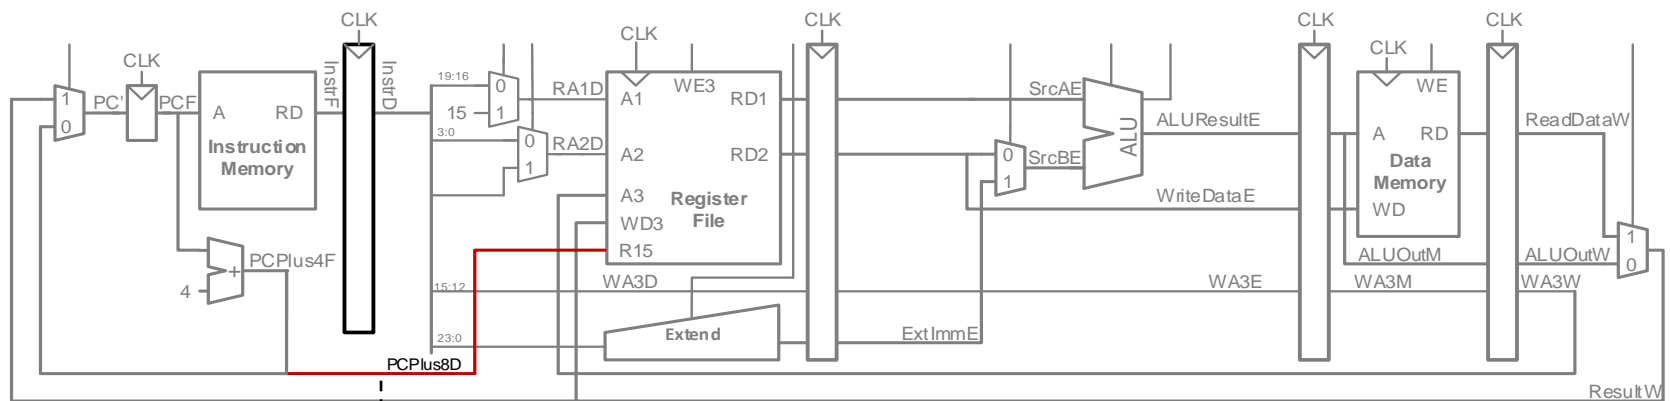
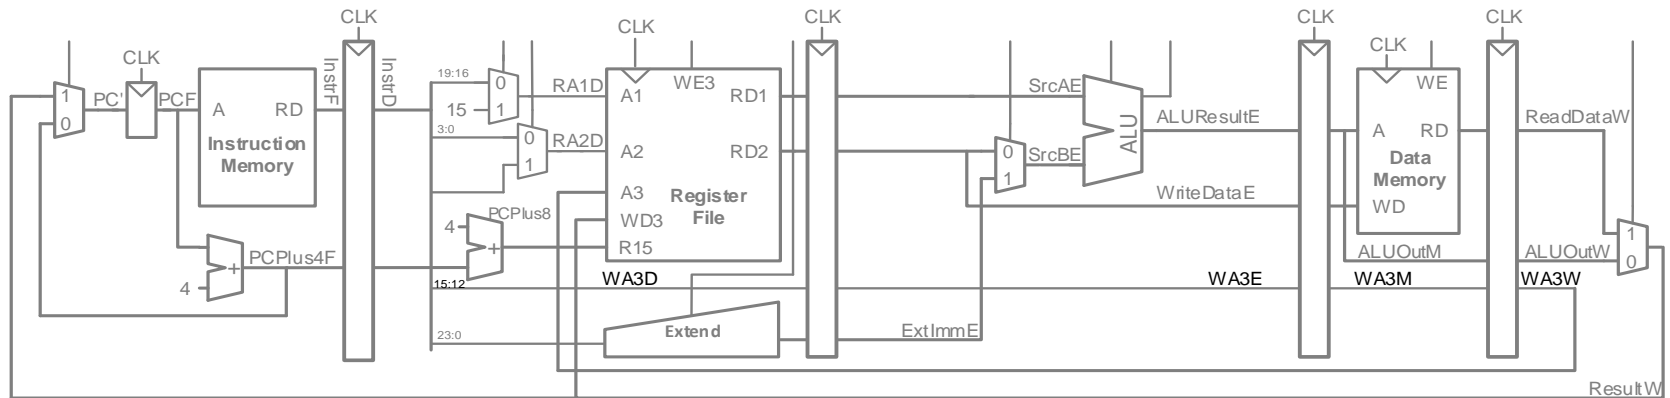
Effectively, ADD R1, R5, R6 ☹️

- WA3** must arrive at same time as *Result* (all the information related to an instruction should move together, i.e., should be delayed equally)



# Optimized Pipelined Datapath

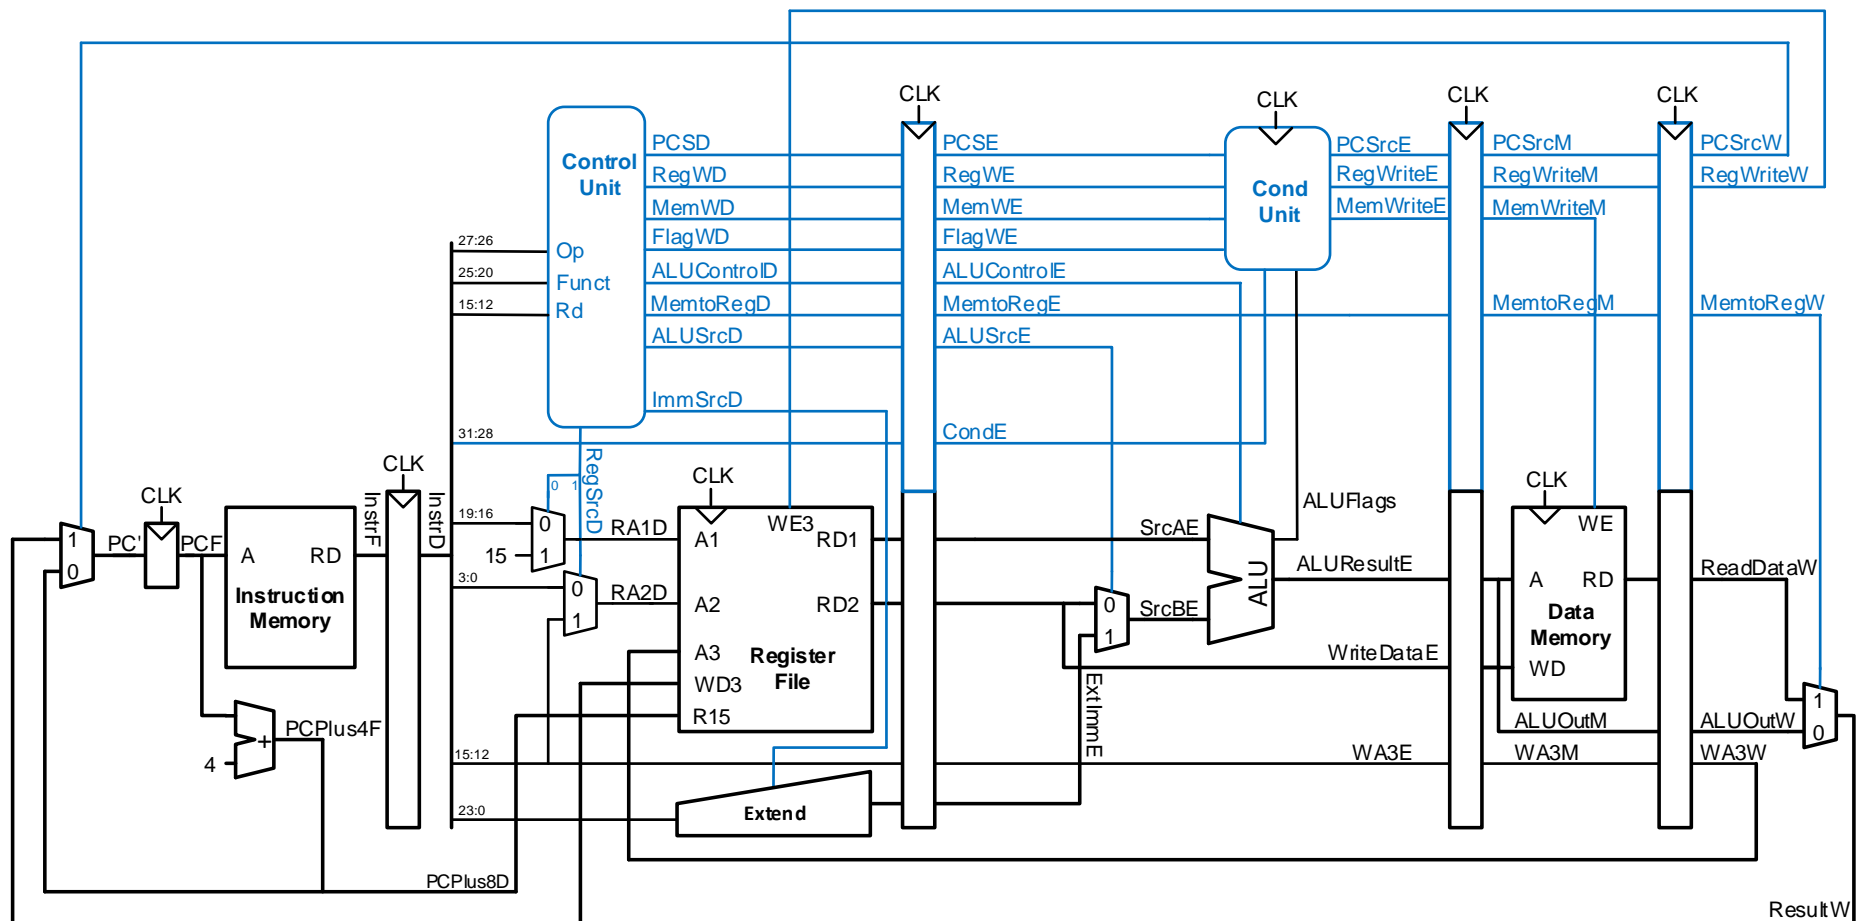
- Remove adder by using  $PCPlus4F$  after  $PC$  has been updated to  $PC+4$



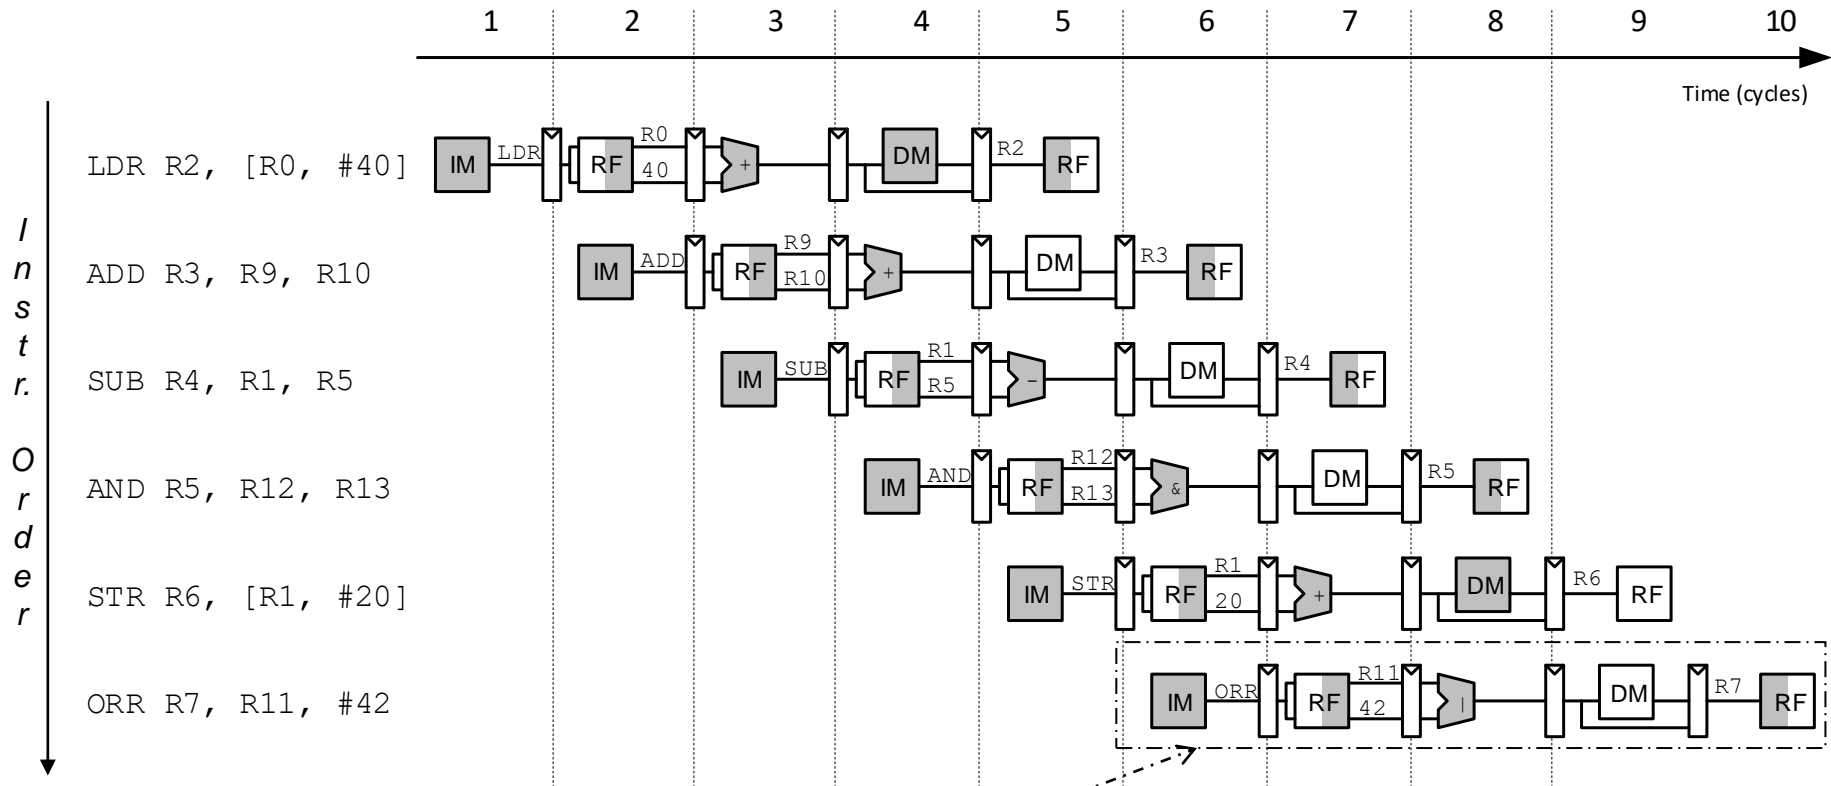
Verify that this works! (Hint :  $PC_{Fetch} = PC_{Decode} + 4$ )

# Pipelined Processor Control

- Same control unit as single-cycle processor
- Control delayed to proper pipeline stage



# Why Pipeline? For Performance!



Abstract pipeline diagram

- Once the pipeline is full, (ideally) one instruction is completed every cycle, so CPI = 1

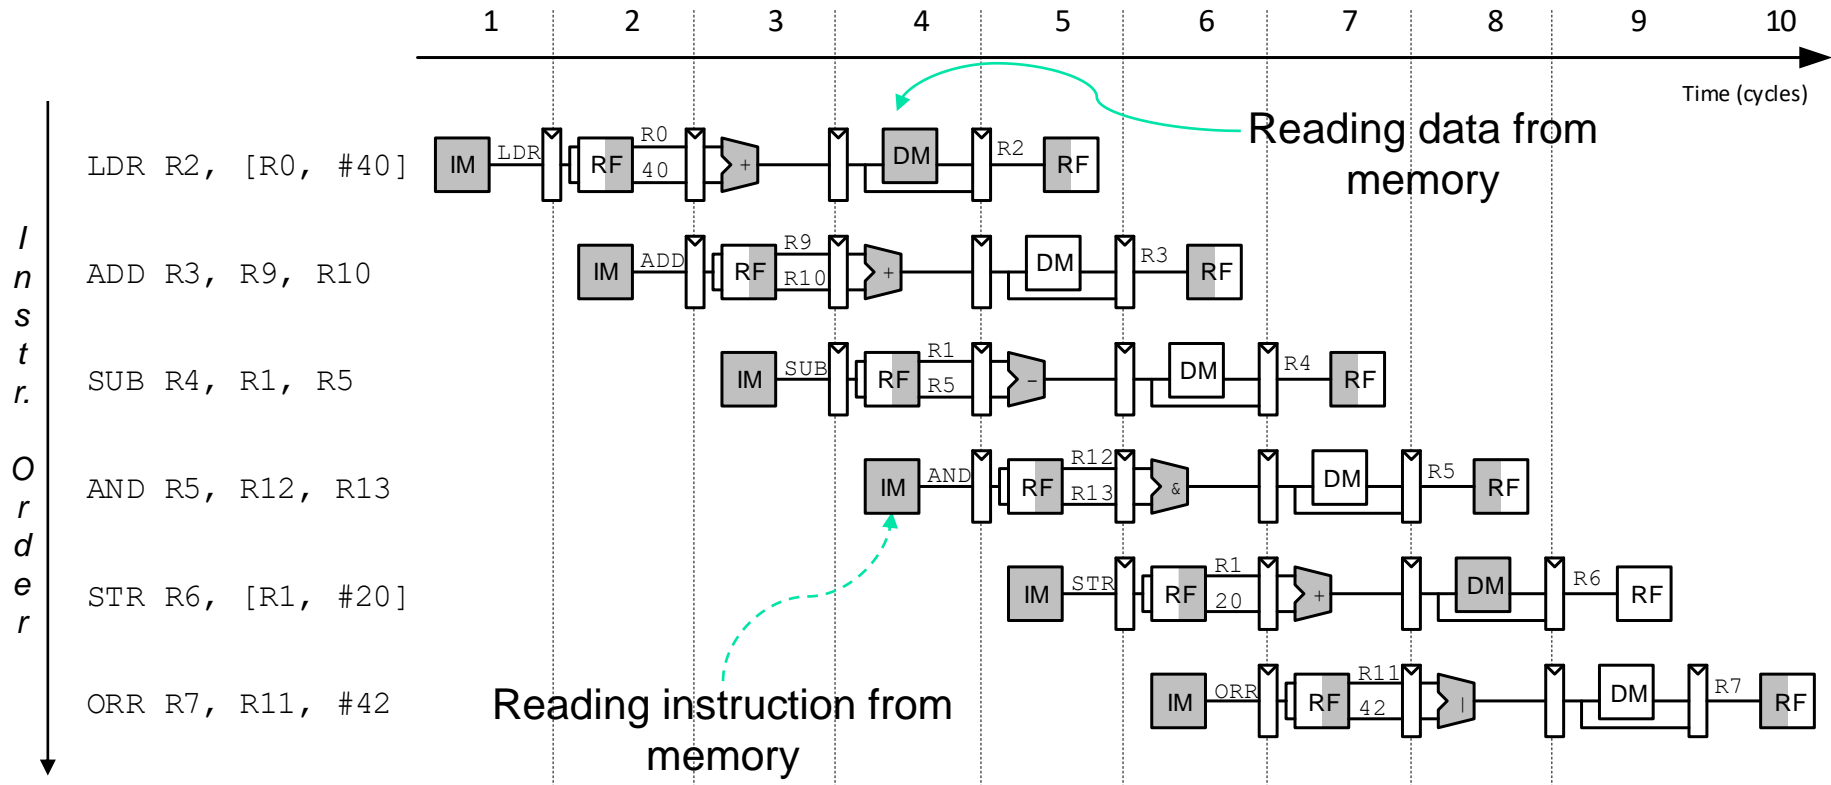


# Can Pipelining Get Us Into Trouble?

---

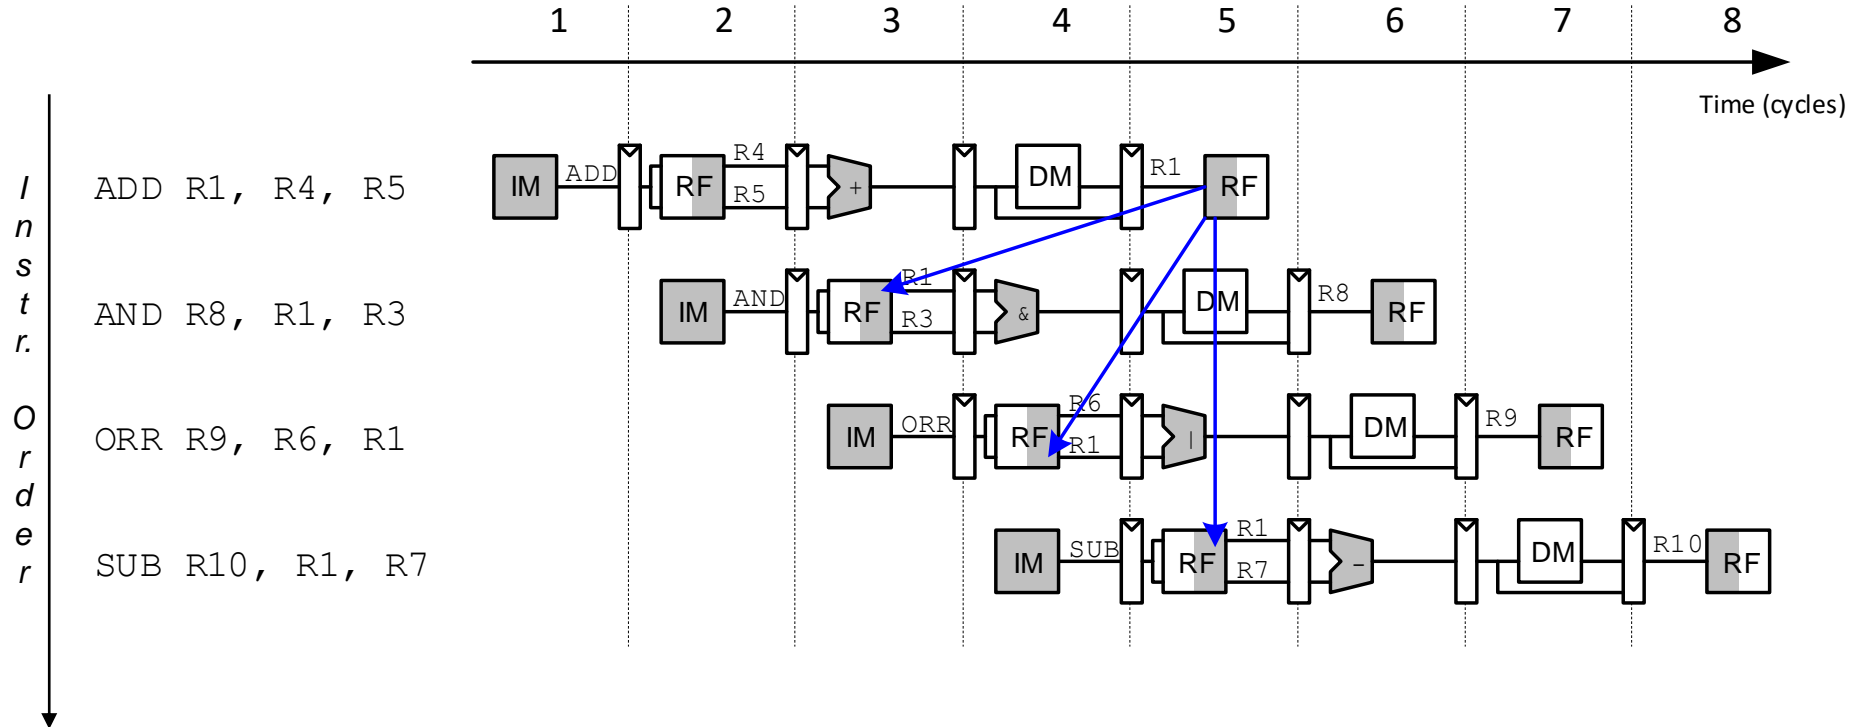
- Yes: **Pipeline Hazards**
  - **structural hazards**: attempt to use the same resource by two different instructions at the same time
  - **data hazards**: attempt to use data before it is ready
    - An instruction's source operand(s) are produced by a prior instruction still in the pipeline
  - **control hazards**: attempt to make a decision about program control flow before the condition has been evaluated and the new PC target address calculated
    - Branch instructions, writes to R15, exceptions
- Can usually resolve hazards by **waiting**
  - Compiler (**software**) or pipeline control (**hardware**) must detect the hazard
  - and take action to resolve hazards

# A Single Memory Would Be a Structural Hazard



- Fix with separate instr and data memories (IM and DM)
  - Or at least separate caches

# Data Hazard



- Also known as RAW (**read after write**) hazard or true data dependency
- Occurs very frequently in practice -> represents the flow of information in the program



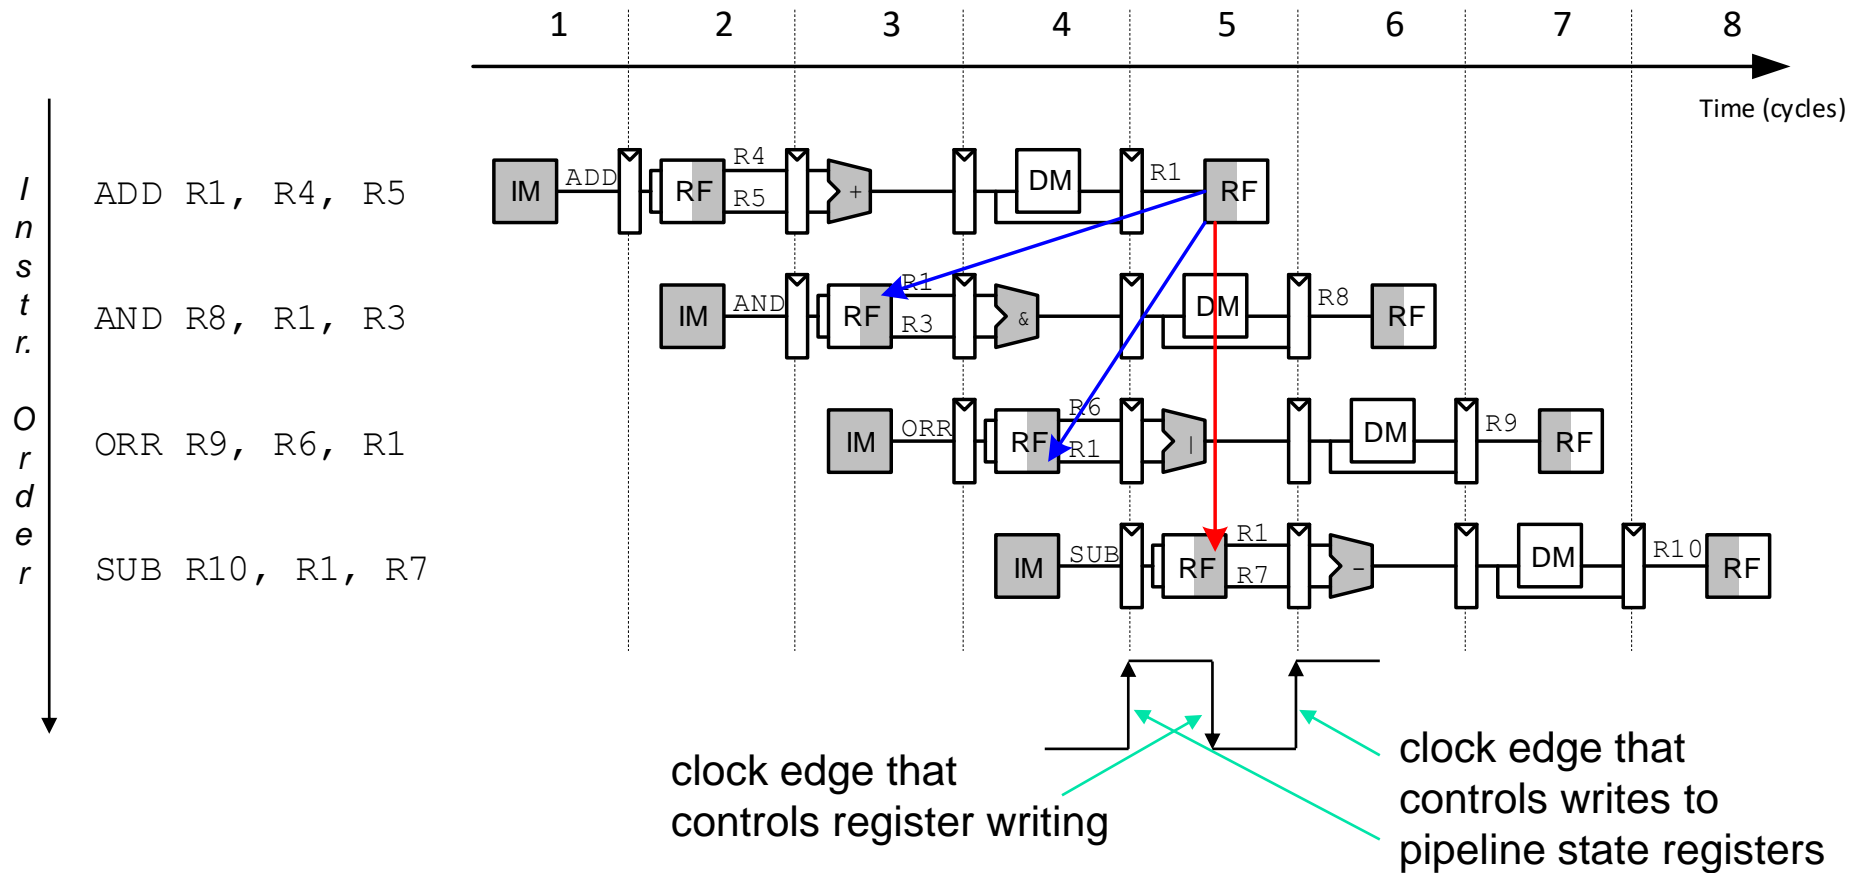
# Handling Data Hazards

---

- Write register file and pipeline registers at different edges of the clock
- Insert `NOPS` in code at compile time
  - The instruction `0x00000000`, which stands for `ANDEQ R0, R0, R0` which has no effect
  - The assembly instruction `NOP` will most likely expand to `MOV R0, R0` which is encoded `0xE1A00000`
- Rearrange code at compile time
- Forward data at run time
- Stall the processor at run time

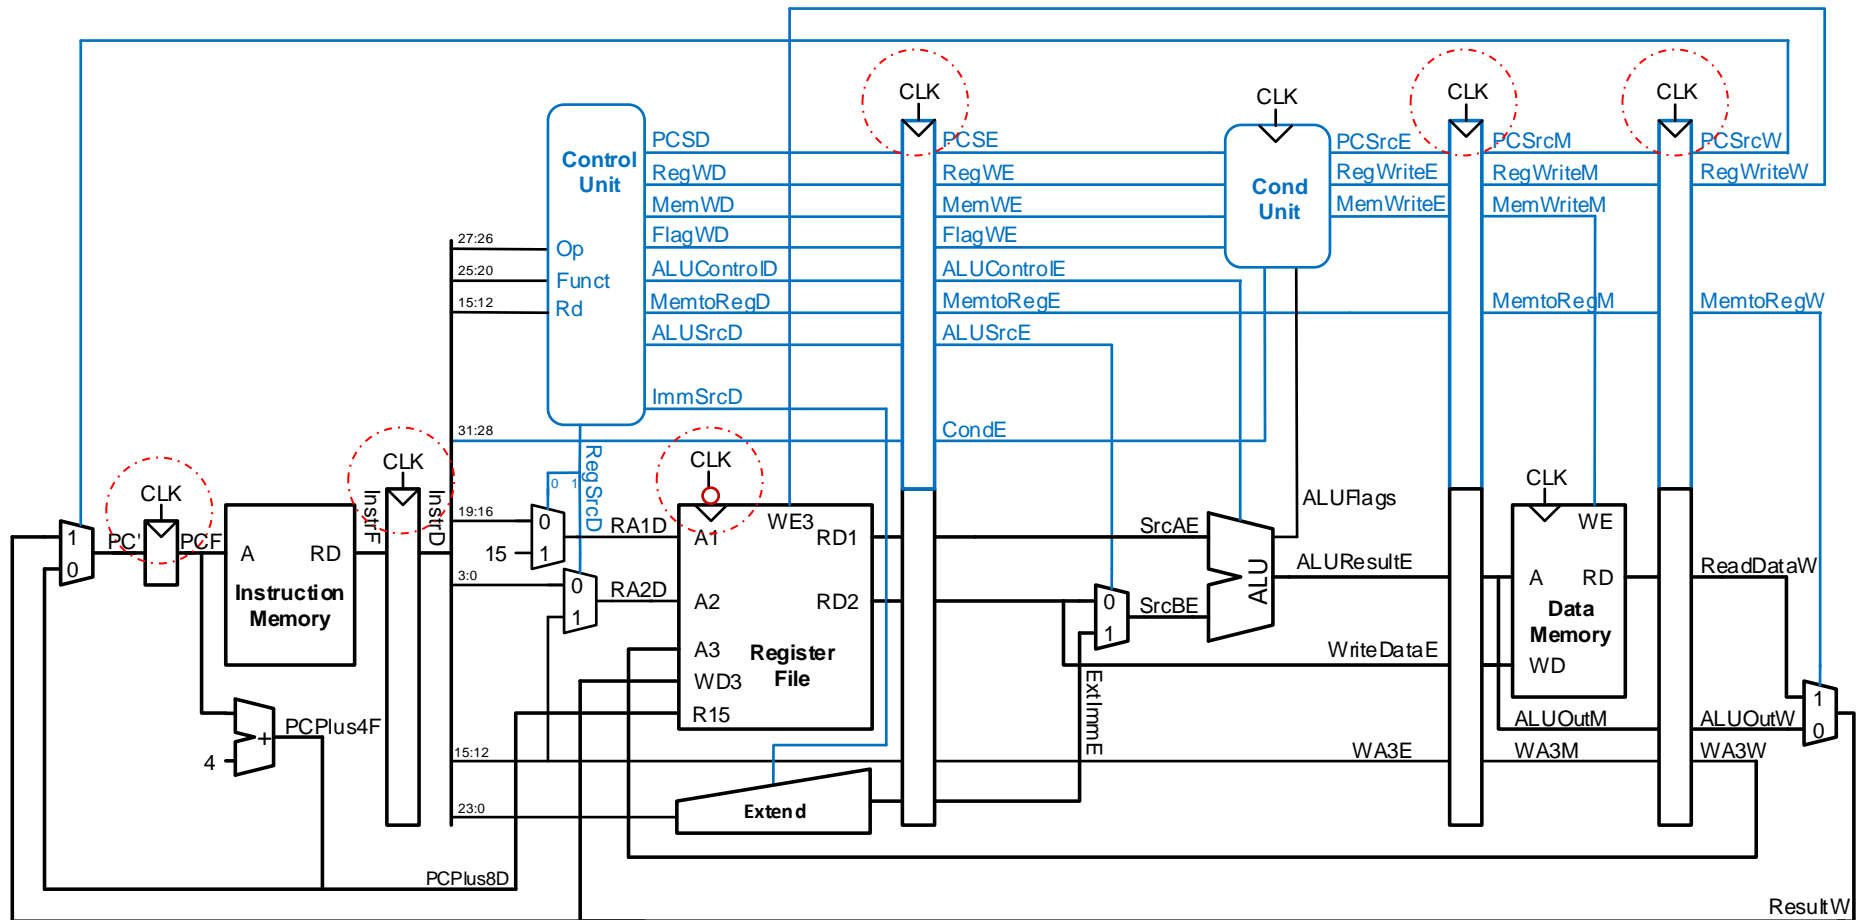


# Data Hazard : Using Different Clock Edges



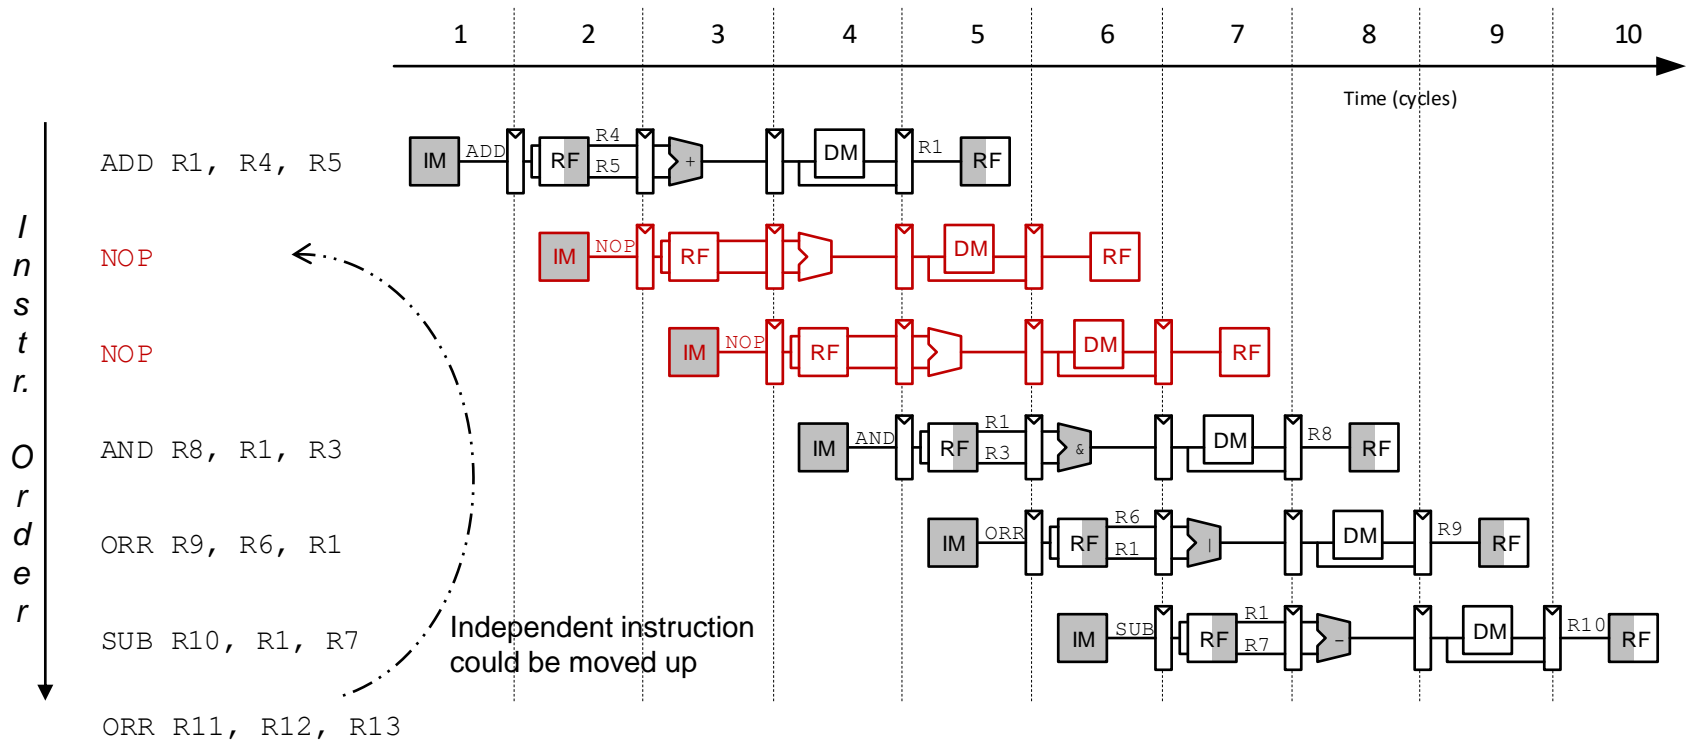
- Reading in the second half of the cycle and writing in the first half helps for instructions which are spaced by at least 2 instructions

# Data Hazard : Using Different Clock Edges

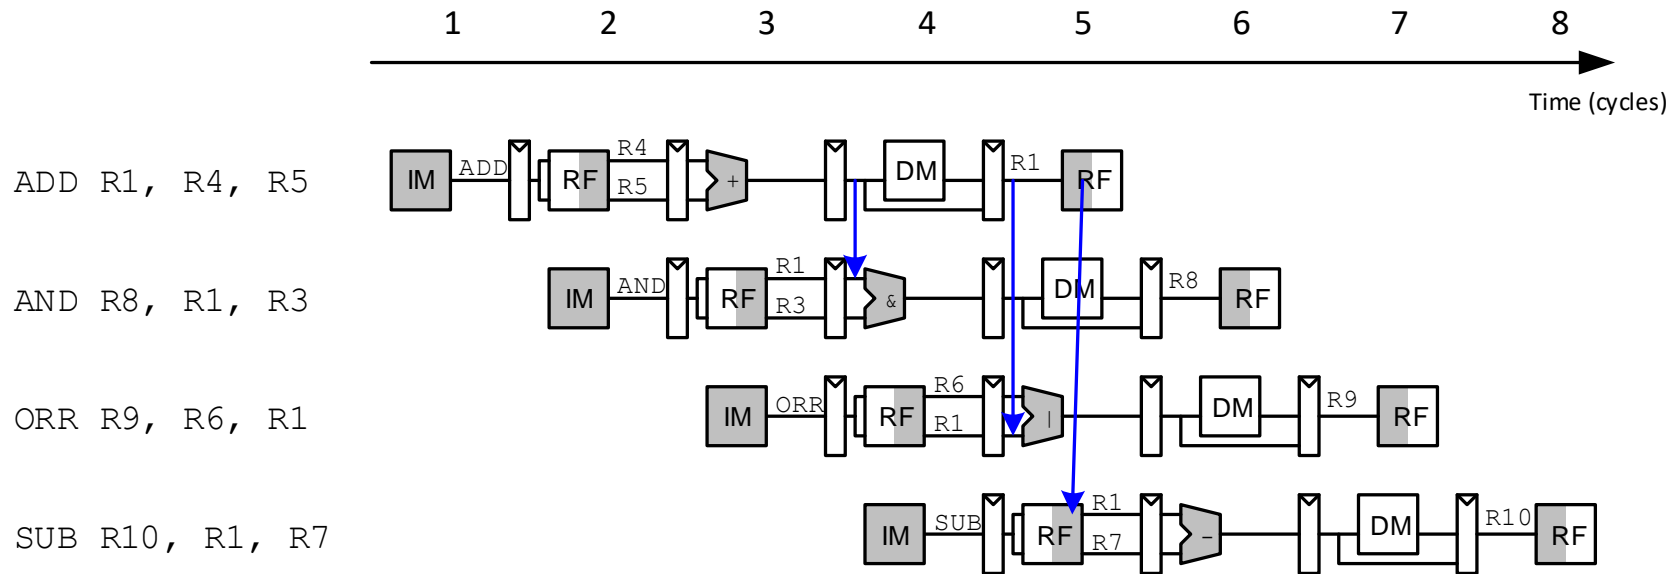


# Data Hazard : Inserting NOPs

- Insert enough NOPs for result to be ready
  - NOPs waste time => CPI is **greater** than the ideal value of 1
  - NOPs waste code memory => makes the code bulkier / bloated
  - Compiler needs to know the microarchitecture => code is not very portable
- Or move independent useful instructions forward
  - Might not be possible all the time



# Data Hazard : Data Forwarding

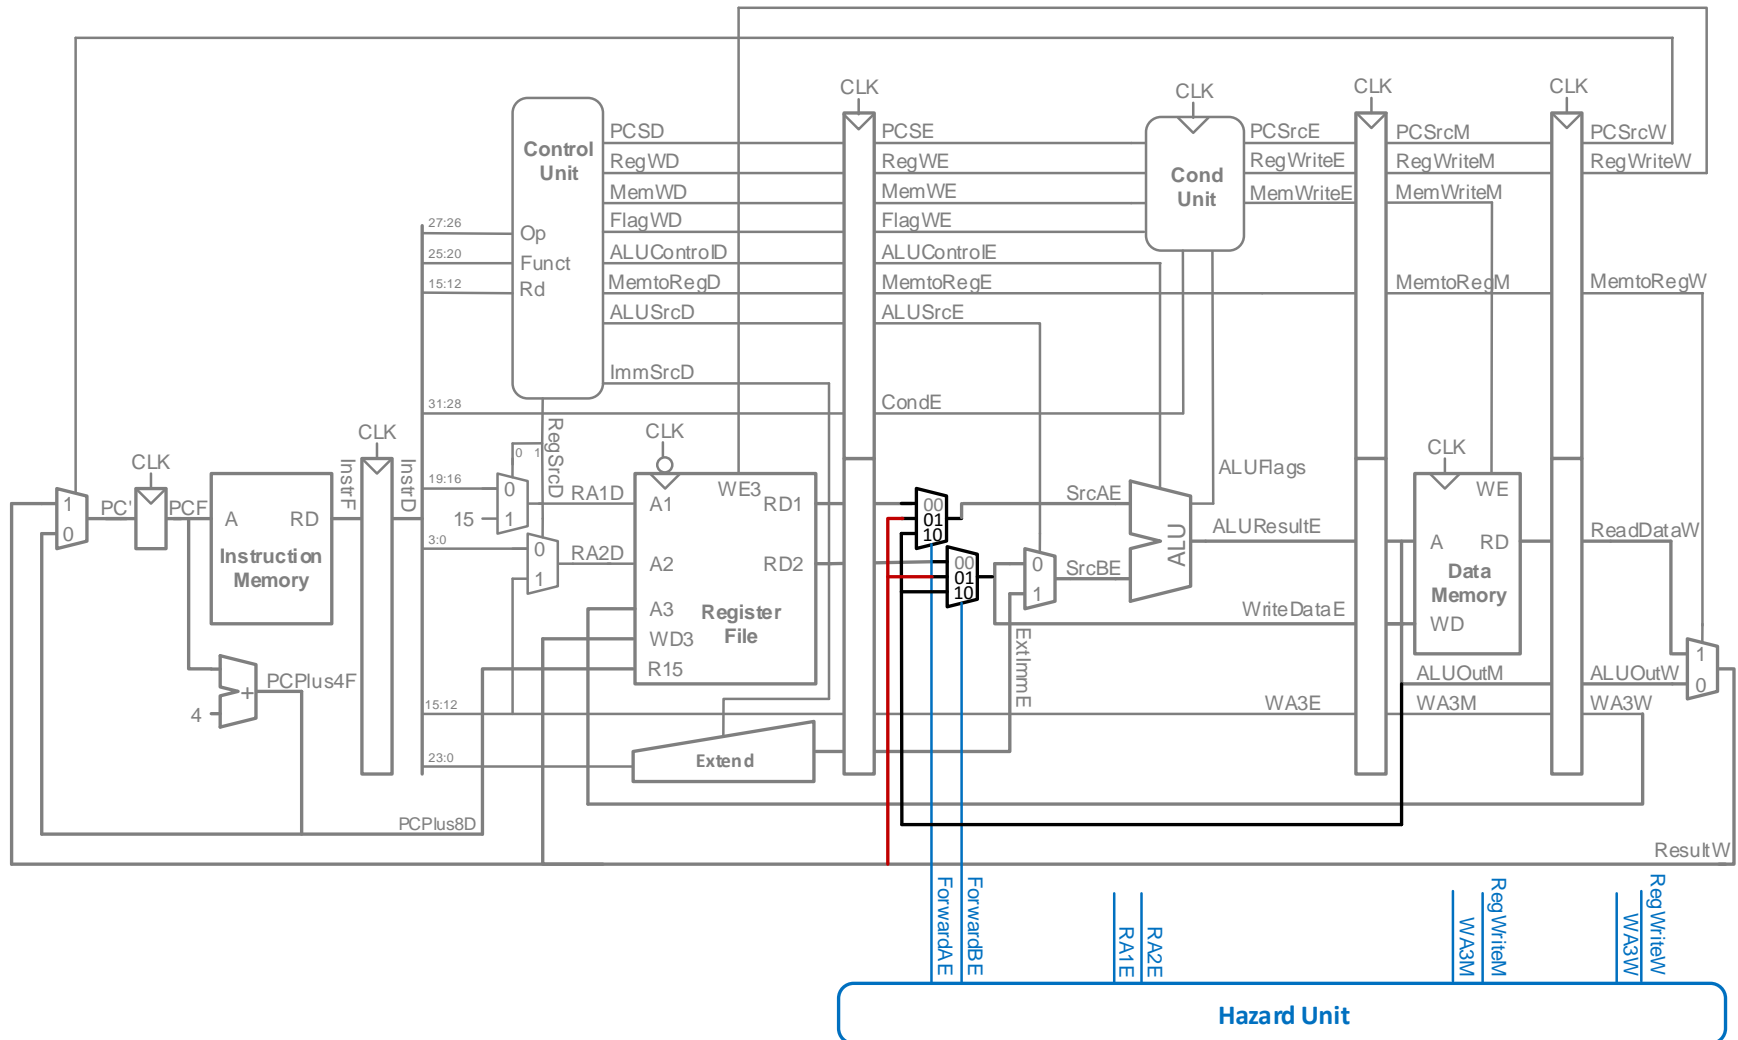


- Check if register read by the instruction which is currently in Execute stage matches register written by the instruction which is currently in Memory or Writeback stage
- If so, forward result

# Data Forwarding ...

- Take the result from the earliest stage / point (latest in the code) that it exists in **any** of the **later** (earlier in the code) pipeline state registers and forward it to the functional units (e.g., ALU) that need it that cycle
  - Take it from where it is present, pass it to where it is needed
- For ALU functional unit: the inputs can come from any pipeline register rather than just from E by
  - adding multiplexers to the inputs of the ALU
  - connecting ALUOutM and ResultW to SrcAE and SrcBE ALU mux inputs
  - adding the proper control hardware to control the new muxes
- Other functional units may need similar forwarding logic (e.g., the Data Memory)
- Can achieve a CPI of 1 even in the presence of data dependencies
- Data forwarding is also called **Bypassing**
  - The usual register file write and then read process for data access is bypassed
  - Register is still written, for future reads

# Data Forwarding : Circuitry



# Data Forwarding : Conditions

Source register for instruction in **Execute** stage matches destination register for instruction in **Memory** stage?

Match\_**1**E\_M = (RA1E == WA3M)

Match\_**2**E\_M = (RA2E == WA3M)

Each comparison  
requires a 4-bit  
comparator. Total  
4 x 4-bit comparators

Source register for instruction in **Execute** stage matches destination register for instruction in **Writeback** stage?

Match\_**1**E\_W = (RA1E == WA3W)

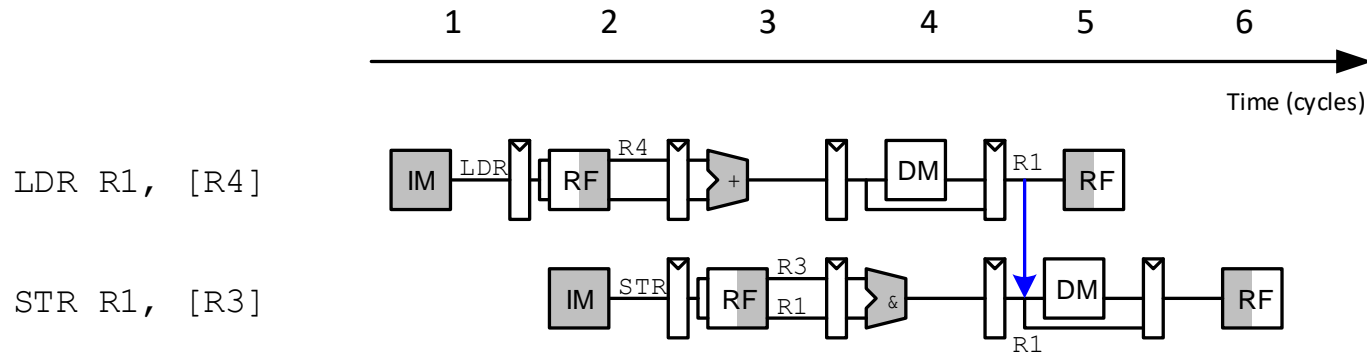
Match\_**2**E\_W = (RA2E == WA3W)

Forward only if the instruction  
in M / W really writes to a  
register. LDR / DP with  
condition not true, B, STR etc.  
do not write to any register

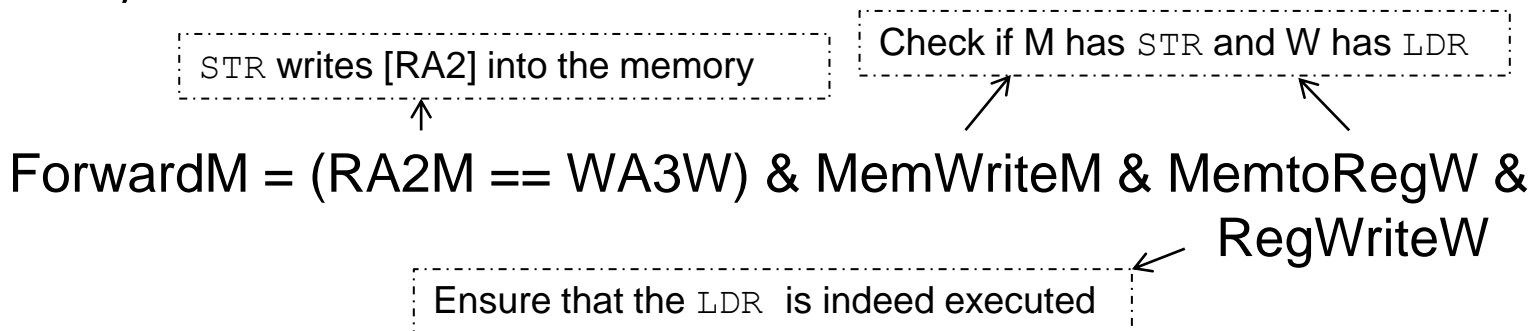
If it matches, forward result:

if	(Match_ <b>1</b> E_M & RegWriteM)	Forward <b>A</b> E = 10;
else if	(Match_ <b>1</b> E_W & RegWriteW)	Forward <b>A</b> E = 01;
else		Forward <b>A</b> E = 00;
Forward <b>B</b> E same but with Match <b>2</b> E		

# Data Forwarding (Memory-Memory copy)

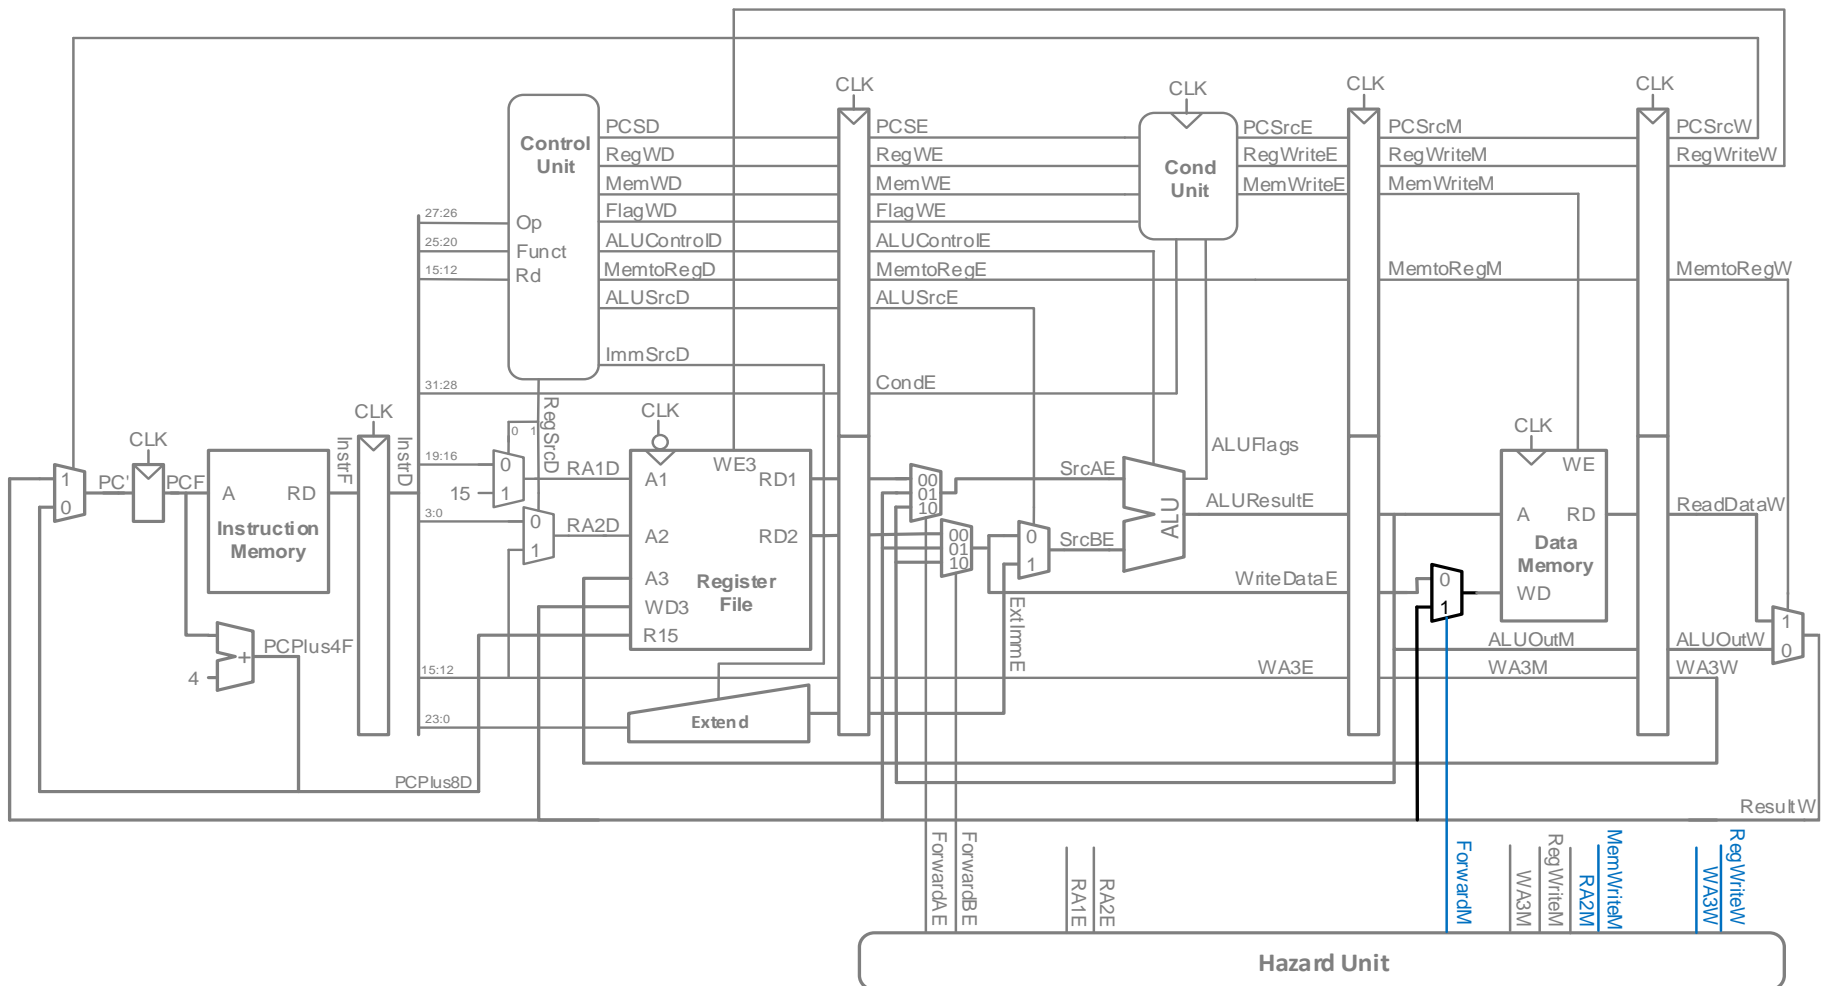


- Check if the register used in Memory stage (i.e.,  $Rd$  of `STR`) by the `STR` instruction matches register written by `LDR` in Writeback stage (i.e.,  $Rd$  of `LDR`)
- If so, forward result

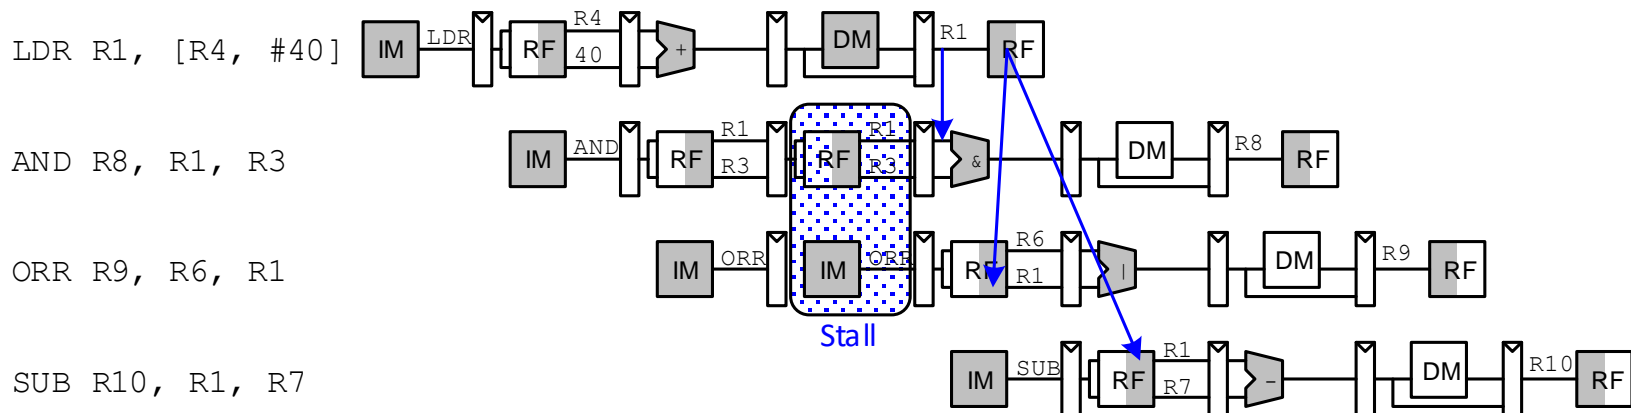
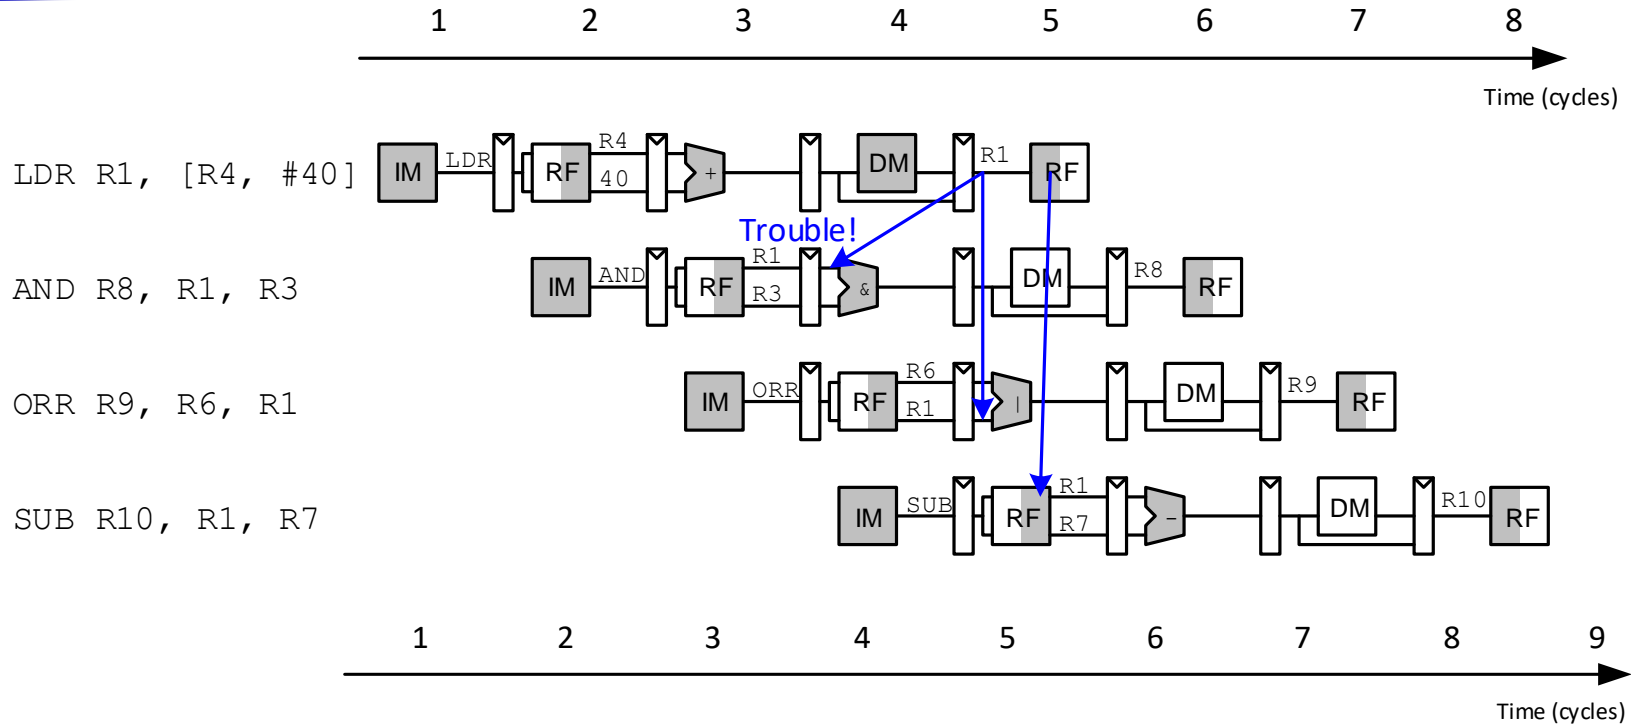




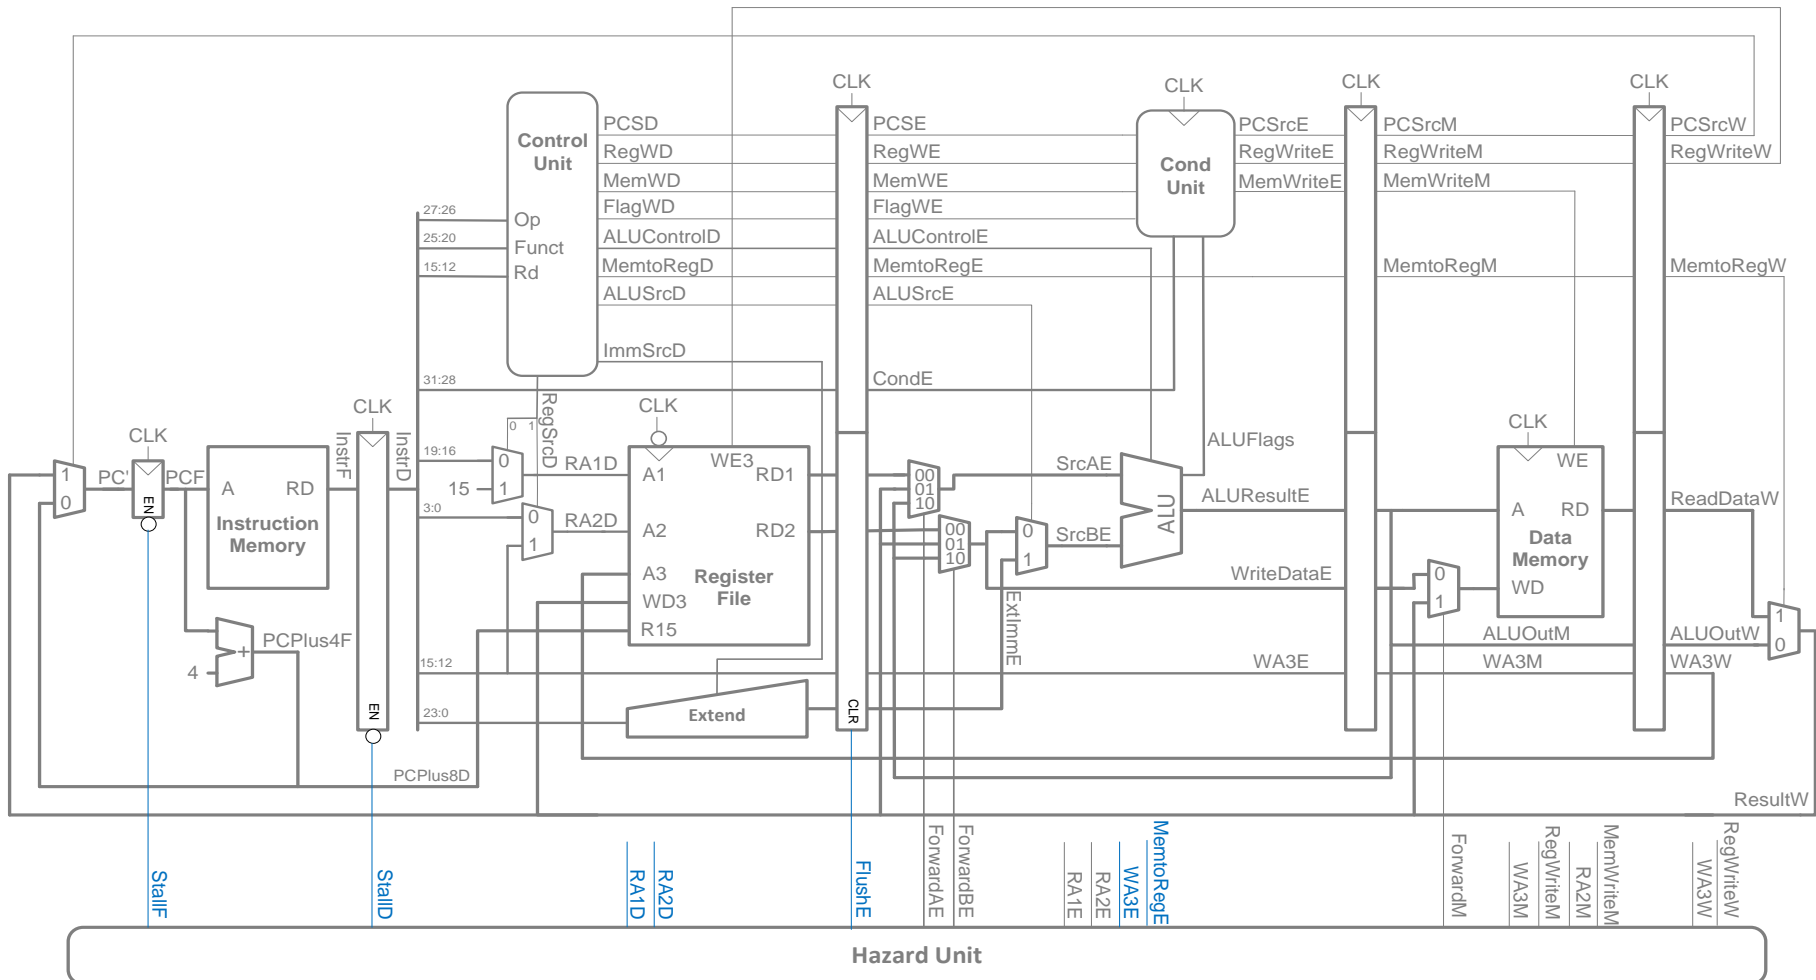
# Data Forwarding : Circuitry (Mem-Mem copy)



# Load and Use Hazard : Stalling



# Load and Use Hazard : Stalling Circuitry



## Load and Use Hazard : Condition

- Is either source register of instruction in the Decode stage the same as the destination of the instruction which is currently in the Execute stage?

$$\text{Match\_12D\_E} = (\text{RA1D} == \text{WA3E}) \parallel (\text{RA2D} == \text{WA3E})$$

- Is a `LDR` in the Execute stage AND `Match_12D_E`?

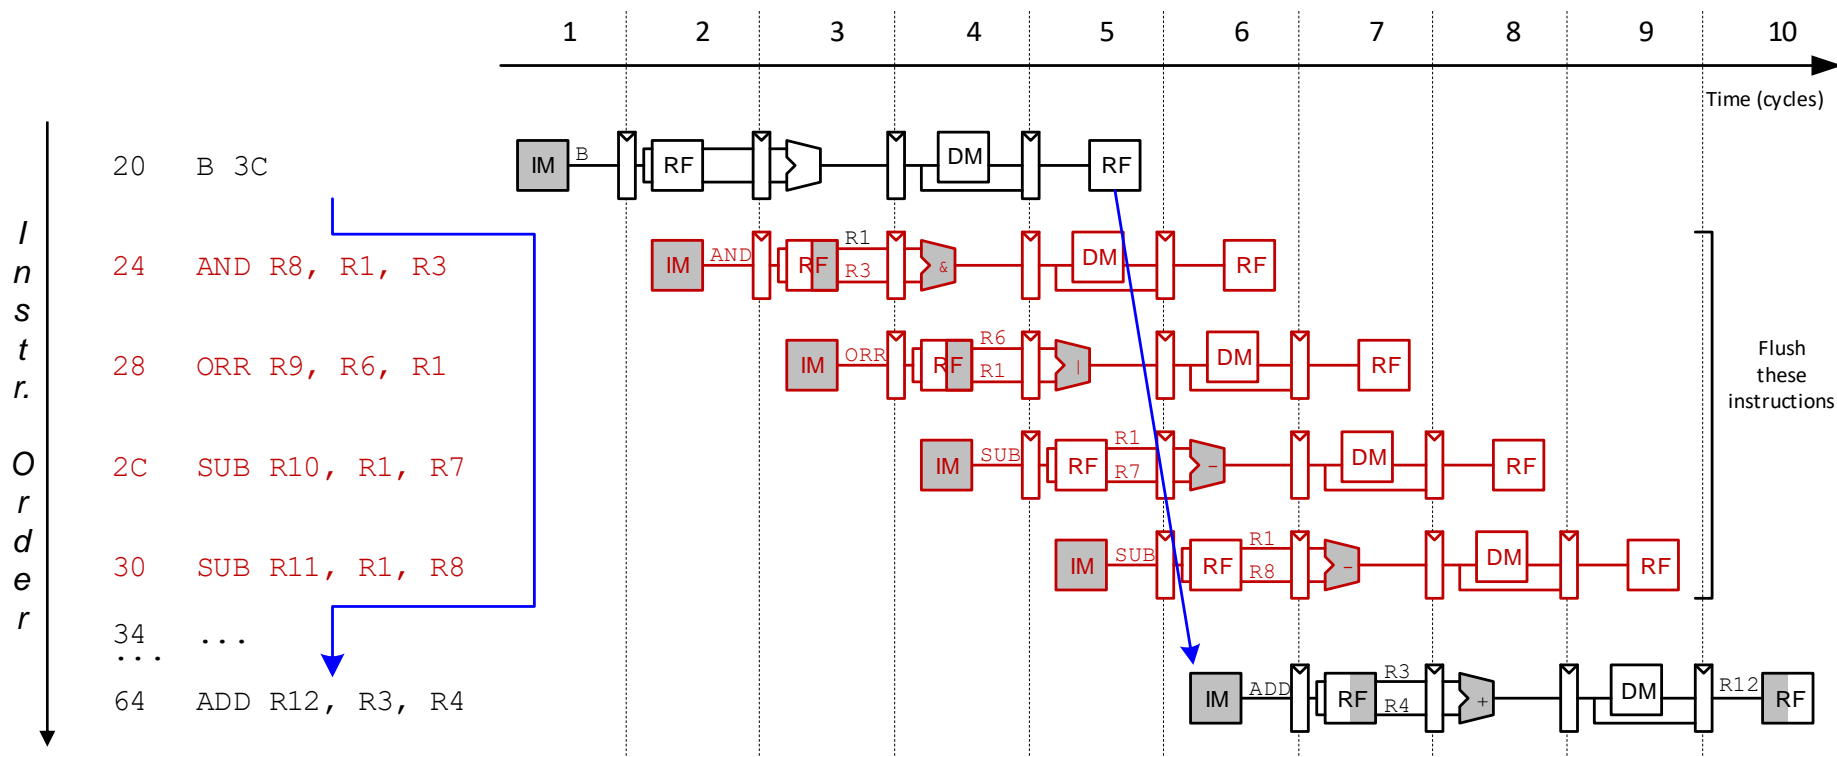
$$\text{ldrstall} = \text{Match\_12D\_E} \& \text{MemtoRegE} \& \text{RegWriteE}$$

$$\text{StallF} = \text{StallD} = \text{FlushE} = \text{ldrstall}$$

- After the 1 cycle stall, the forwarding circuitry (W to E) will take care of delivering the correct data
- `MOV` does not use `RA1`. `B`, `LDR`, `DP Imm` doesn't use `RA2`. Mem-mem copy forwarding takes care of `STR` following `LDR`. Stalls can be avoided in these cases by adding appropriate exception conditions (Omitted here for simplicity)

# Control Hazards

- Branch instructions, writes to R15 can cause control hazards



- Compile time NOPs or rearranging the code\* can help, but
  - NOPs reduce performance, even when the branch is not taken
  - Increases the code size and makes the code microarchitecture dependent

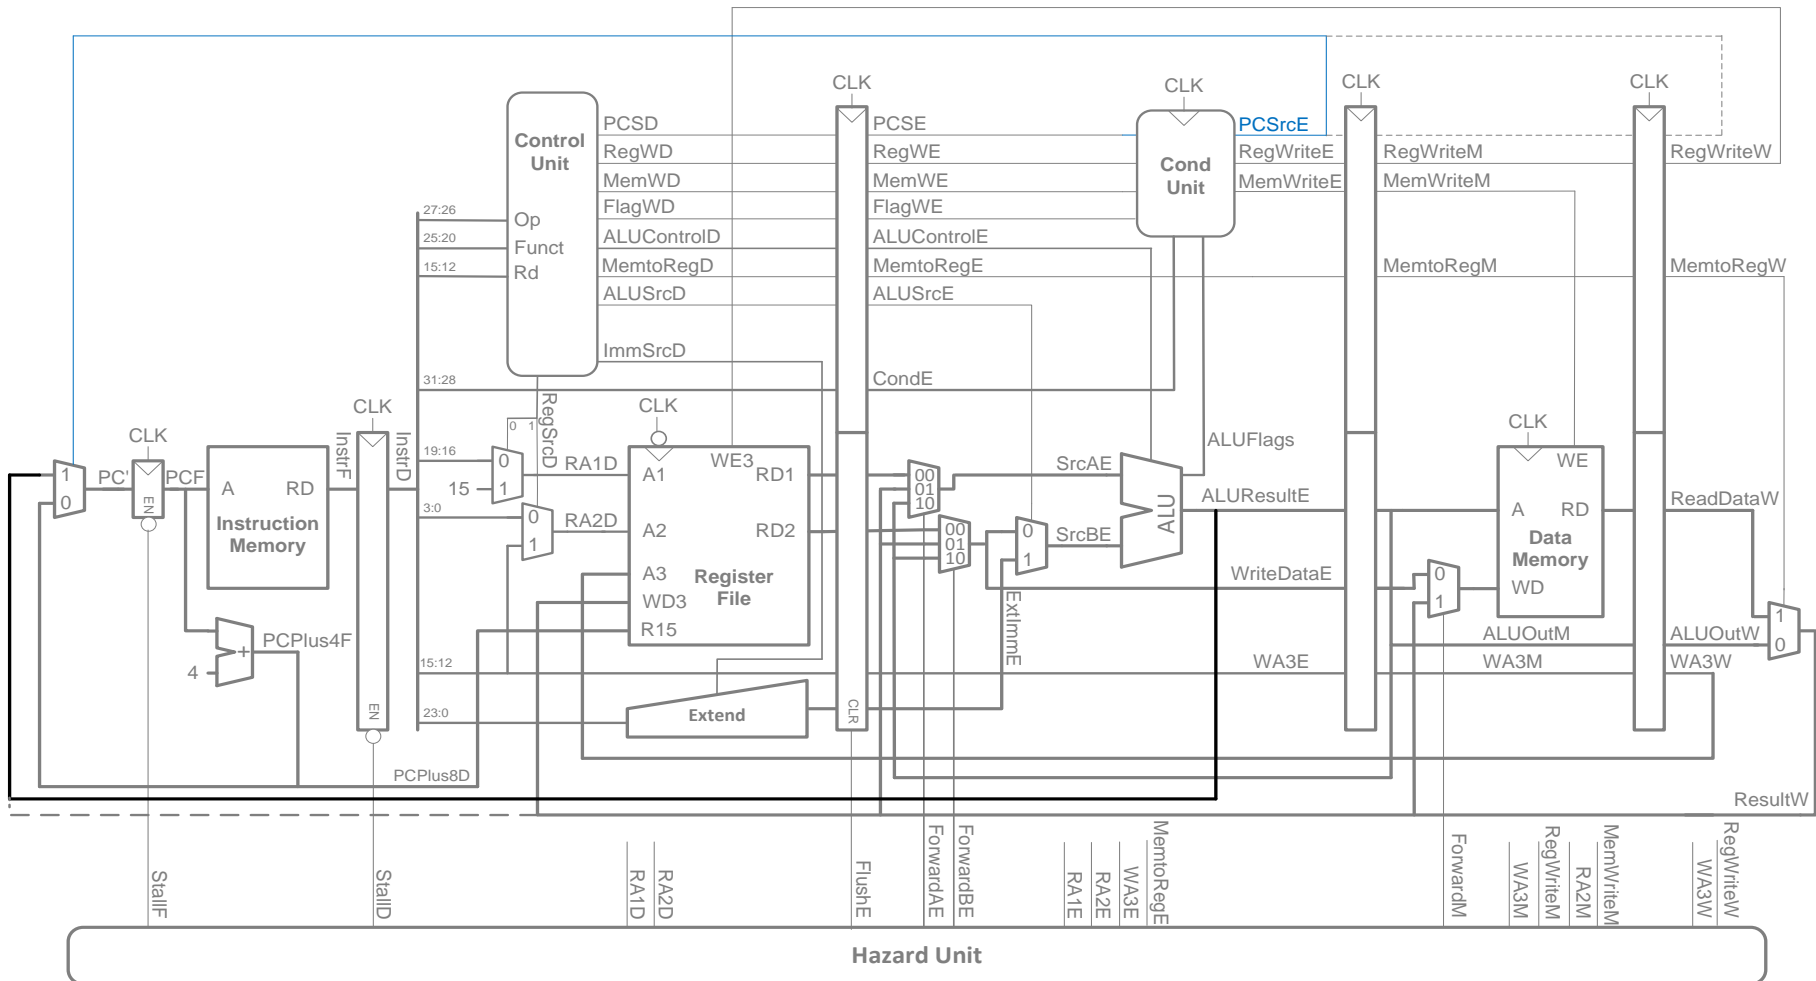
\* The instruction(s) immediately following the branch should be "safe instruction(s)", which are executed irrespective of whether the branch is taken or not (provided it is specified so in the architecture and supported by the compiler). These instructions should hopefully be useful; and shouldn't affect the correctness even if they are not. Google for **delay slot**

# Control Hazards

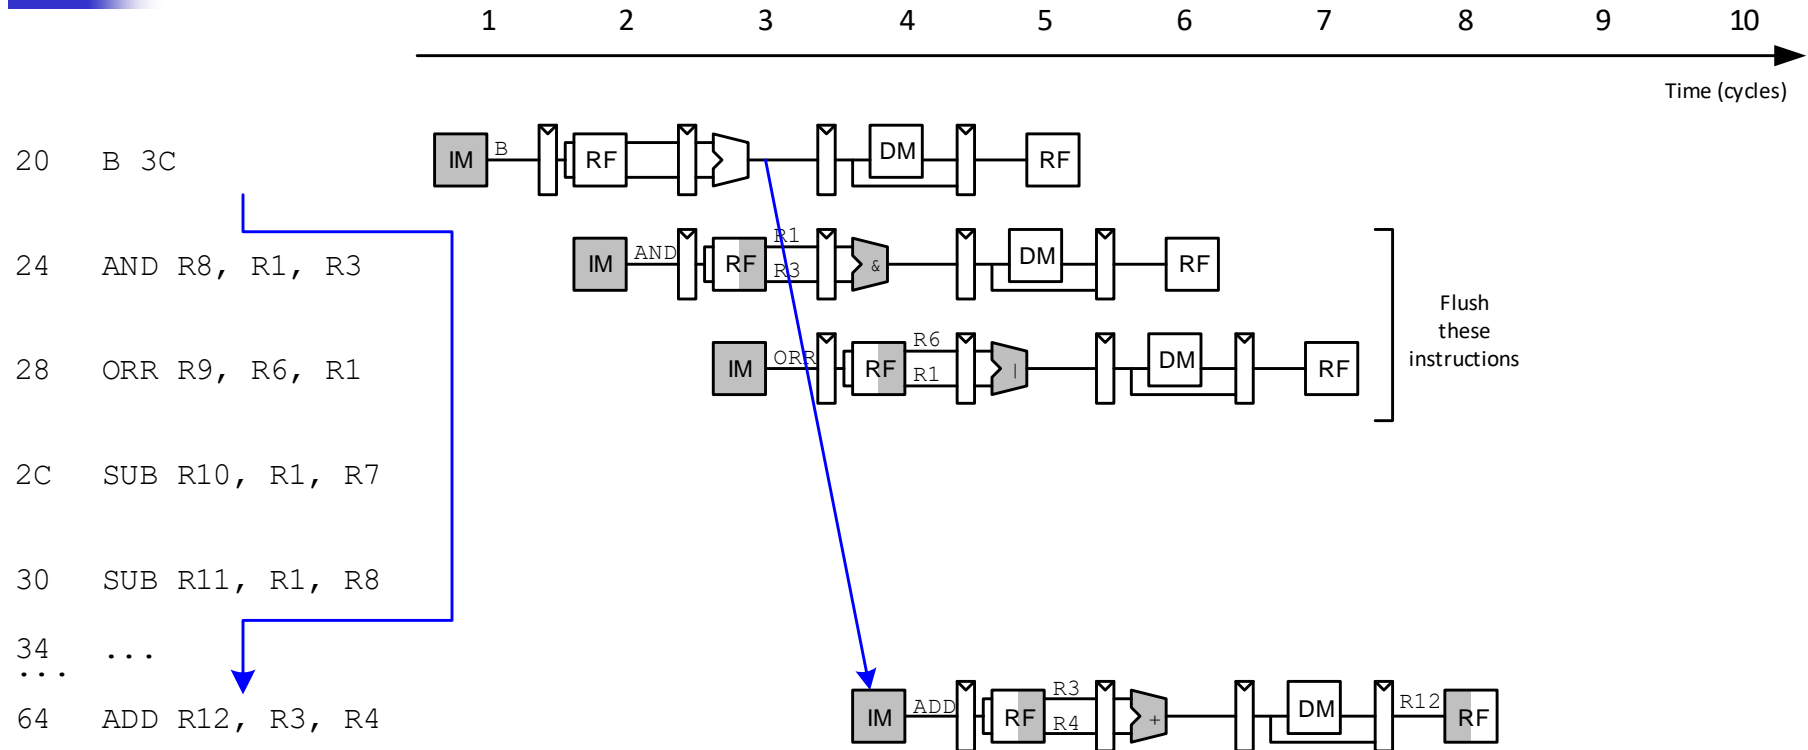
BTA = Branch Target Address

- Possible approaches
  - Stall until branch decision and BTA are available (impacts CPI)
  - Delay decision (delay slot - requires compiler support)
  - **Early BTA** - move decision point as early in the pipeline as possible, thereby reducing the number of stall cycles
  - **Branch prediction** - predict and hope for the best !
  - Fetch and execute from both branches (if BTA is known)
  - Predicated/conditional execution, fine grained multithreading,...
- Control hazards occur less frequently than data hazards, but there is *nothing* as effective against control hazards as forwarding is for data hazards
- Early BTA
  - Determine BTA in Execute stage
  - For now, ignoring the case of `LDR PC, ...` (`LDR` with `PC` as destination)
  - Branch misprediction penalty = 2 cycles
  - Could increase the critical path delay ☹

# Pipelined Processor with Early BTA



# Control Hazards with Early BTA



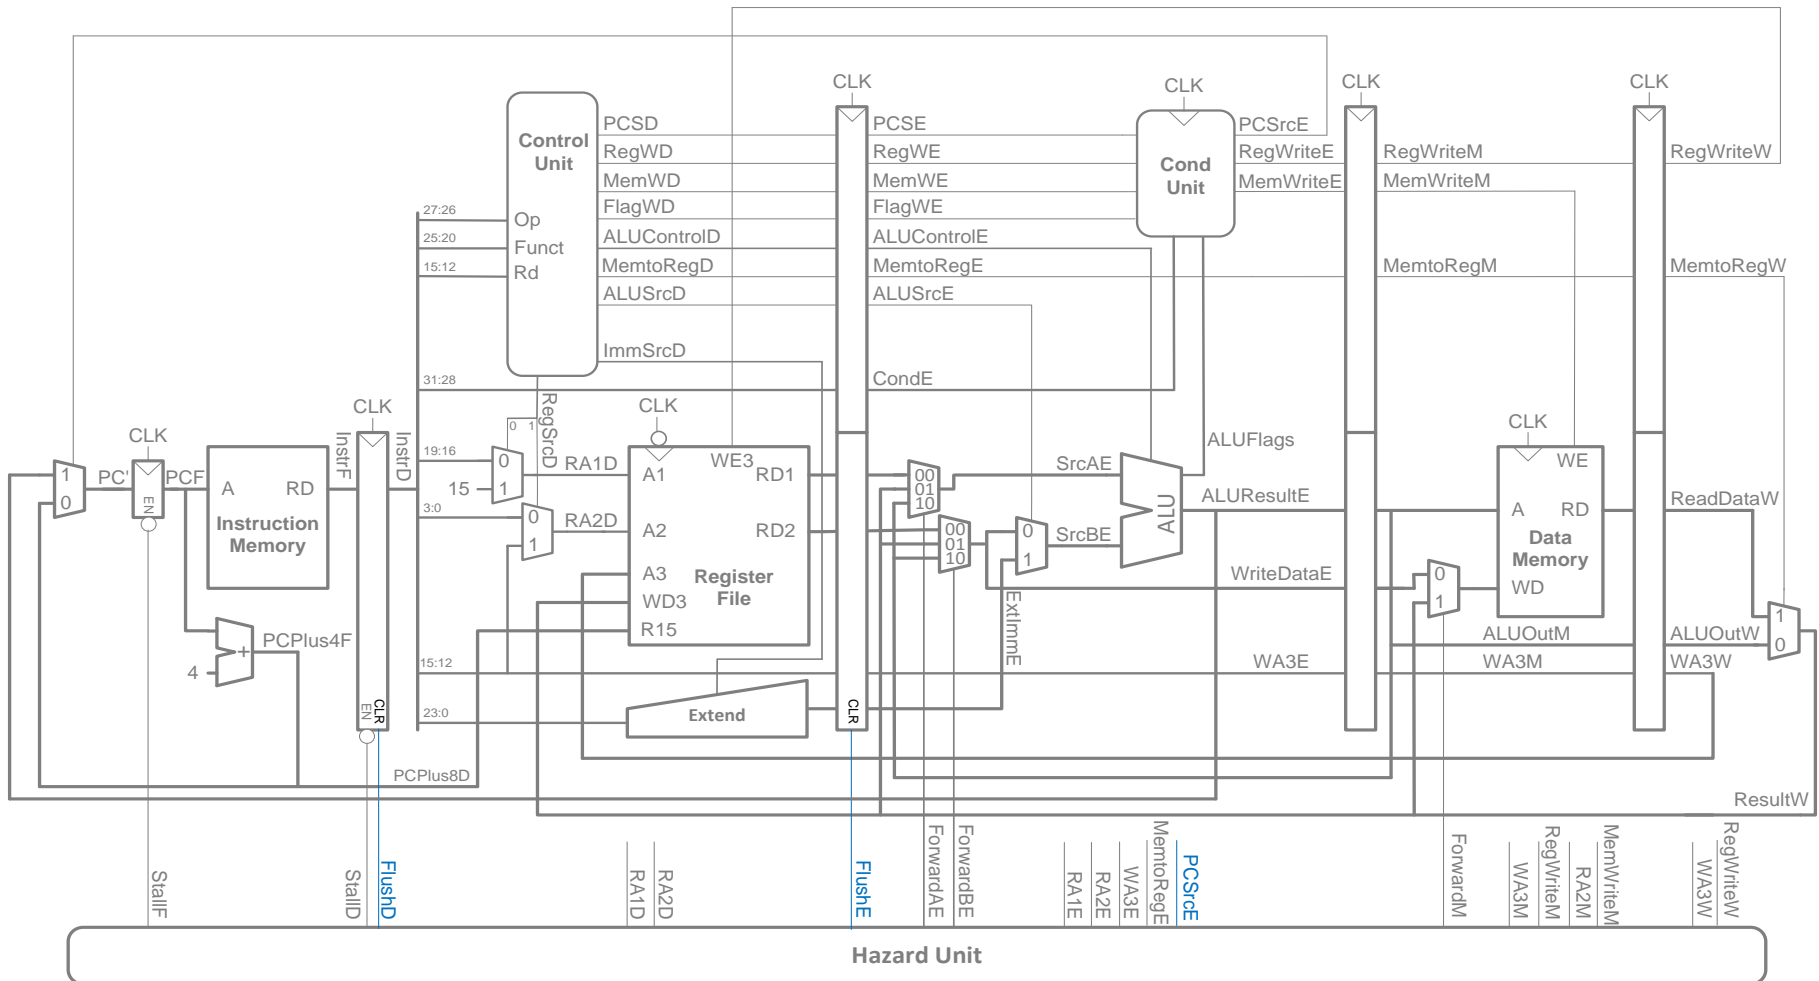
- If PCSrcE is true, flush the two earlier stages to quash the two following instructions

$$\text{FlushD} = \text{FlushE} = \text{PCSrcE}$$

$$\text{FlushE (effective)} = \text{ldrstall (load and use)} \parallel \text{PCSrcE (control)}$$



# Control Hazards : Flushing Circuitry



# Stall vs Flush

- Stall involves deasserting the write control of the pipeline register
- Flush involves asserting the 'clear' of the pipeline register
  - In practice, clear only those bits which can cause a change in state (PCSE, MemWE, RegWE, FlagWE is sufficient for E. InstrD for D)
- If any stage is stalled, (i) the previous stages should also be stalled, and (ii) the following stage should be flushed (think load and use)
  - If (ii) is not done, the instruction in the latest stage being stalled will be executed 2 times – first using incorrect data and later using correct data (this is not ok if the instruction writes to flags or cause exceptions etc.)
  - Effectively, an NOP is **inserted** between the instruction in the latest stage being stalled and the instruction in the *following* pipeline stage (*earlier* in the code)
  - Instructions in the stages being stalled are bounced in place, while the instructions in the later stages of the pipeline (earlier in the code) are allowed to proceed to completion normally
- If any stage is flushed without stalling the previous stage, the instruction which is supposed to go into the flushed stage is **replaced** with an NOP (think branch)



# Summary

---

- Almost all modern day processors use pipelining
- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Potential speedup: a CPI of 1 and a fast CCT
- Pipeline rate limited by **slowest** pipeline stage
  - Unbalanced pipe stages makes for inefficiencies
  - The time to "**fill**" pipeline and time to "**drain**" (flush) it can impact speedup for deep pipelines and short code runs
- Must detect and resolve hazards at compile time (by the compiler) or runtime (by the hardware)
  - Compile time NOPs reduce performance, increases code size and limits the portability of the code
  - Run time hazard detection requires additional hardware