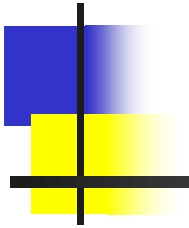

7 : Memory System Principles

Rajesh Panicker,
NUS

CG2028

Acknowledgement :

- Some slides from Prof. Bharadwaj Veeravalli
- Text by Patterson and Hennessey and companion slides
- Text and companion slides by Harris and Harris

Data Memory Interface

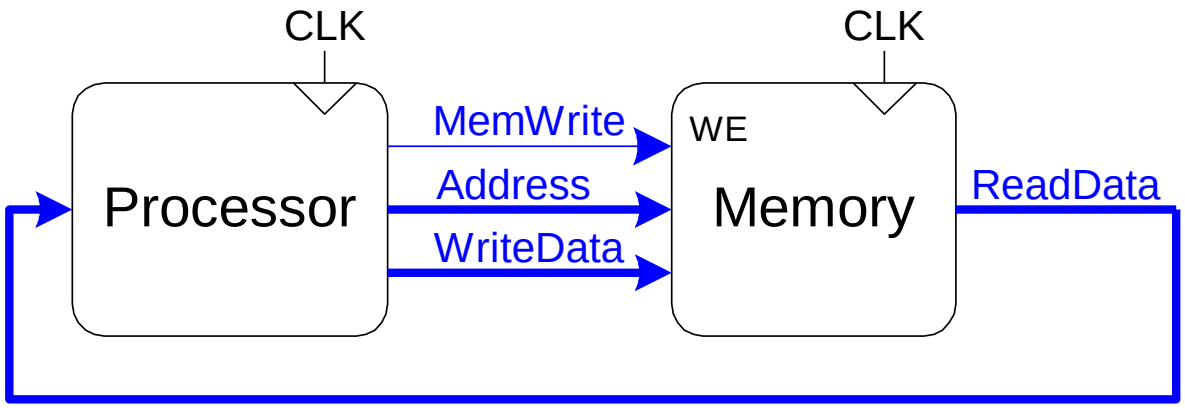

- WriteData and ReadData often combined into a single bidirectional bus interface
- Data Memory is typically composed of non-volatile memories (ROM/Flash) and volatile (RAM) memories. RAM has no contents at power on – a location should be written before it is read (i.e., there should be an STR to a location before LDR)
- Instruction memory is also external to the processor (not shown here)

Storage Media – Access Times and Costs

- Static RAM (SRAM)
  - Data stored in flip-flops
  - Usually used for cache
  - 0.5ns – 2.5ns, $2000 – $5000 per GB
- Dynamic RAM (DRAM)
  - Data stored in capacitors – refresh circuitry required
  - Usually used for Main Memory (MM)
  - 50ns – 70ns, $20 – $75 per GB
- Magnetic disk (hard disk)
  - 5ms – 20ms, $0.20 – $2 per GB
  - Getting replaced by faster and more reliable (but costlier) flash-based solid state drives (SSDs)
- Ideal memory – best of both worlds
  - Access time of SRAM; capacity and cost/GB of disk

Memory Hierarchy

registers in CPU M_0

| bytes/words

cache memories (SRAM) M_1

| blocks

Main memory (DRAM) M_2

| pages

disk storage M_3

| file segments

tape units/optical disks, etc. M_4

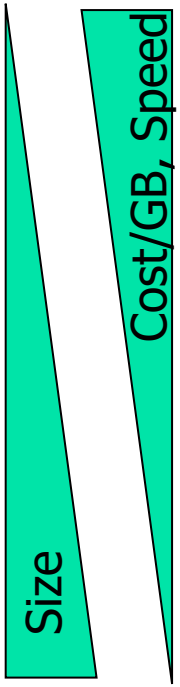

- Basic idea
  - Each level holds the most frequently accessed data from the immediate higher level
  - Reduces the effect of lower speed of the higher level without increasing the overall cost significantly

Memory Hierarchy Properties

- Coherence (consistency) Property
  - Emphasizes the need for the copies of same data to have same information at all the levels where the data is currently residing
  - If a word is modified in the cache, it must be updated at all levels
- Locality of references
  - The memory access pattern tends to be clustered in certain regions in time, space, and ordering
  - 90-10 rule by Hennessy and Patterson (1990) - a typical program may spend 90% of its execution time on only 10% of the code such as the innermost loop of a nested loop
  - Temporal: Recently referenced items are likely to be referenced in the near future - keep recently accessed data at a faster level
  - Spatial: Refers to the tendency of a process to access the items whose addresses are near to one another - when accessing data, bring nearby data also into a faster level

Memory Capacity Planning

- Hit ratios
  - When a memory M_i is accessed and if the desired word is found, it is referred to as a *hit*, otherwise *miss*
  - The hit ratio (h_i) is the probability that a word/information will be found when accessed in M_i. Miss ratio is $1-h_i$
  - The hit ratios at successive levels are a function of memory capacities, management policies, and program behaviour
  - $h_0=0$ and $h_n=1$. This means that the CPU always access M_1 first and access to the outermost level is always a hit
- Access frequency at a level i is defined as
  - $f_i = (1-h_1)(1-h_2)\ldots(1-h_{i-1})\, h_i$
  - Note that $f_1 + f_2 + \ldots + f_n = 1$ and $f_1 = h_1$
  - Due to the locality property, the access frequencies decrease rapidly from the lower levels, i.e., access freq at level i is greater than i+1
  - This means that the inner levels are accessed more often than the outer levels

Memory Capacity Planning …

- Effective Access Time is defined as

$$T_{eff} = f_1 t_1 + f_2 t_{2+\ldots} f_n t_n$$

where t_i is the access time at level i

- The total cost of a memory hierarchy is estimated as

$$C_{total} = c_1 s_1 + c_2 s_2 + \ldots + c_n s_n$$

where c_i is the cost/MB and s_i is the size (in MB) at level i

- Hierarchy optimization involves minimizing

T_{eff} given $C_{total} < C_{max}$ or C_{total} given $T_{eff} < T_{max}$

- The optimal design should result in a T_{eff} close to t_1 and a total cost close to c_n

Cache Working Principle

- When a read request is received from CPU, the contents of a block of memory words containing the location specified are transferred to the cache
  - Block is also called *cache line*, typically ~64 bytes
  - *Where* to place the incoming block in the cache is decided by the *mapping function*
- Subsequently, when the program asks for any of the locations from this block, the desired contents are read directly from cache
  - CPU need not even be aware of the presence of the cache and issues addresses meant for main memory (MM, usually DRAM)
  - Checking if the required data/block is present in the cache is performed by the cache controller
  - If yes, a *cache hit* is said to occur

Cache Working Principle …

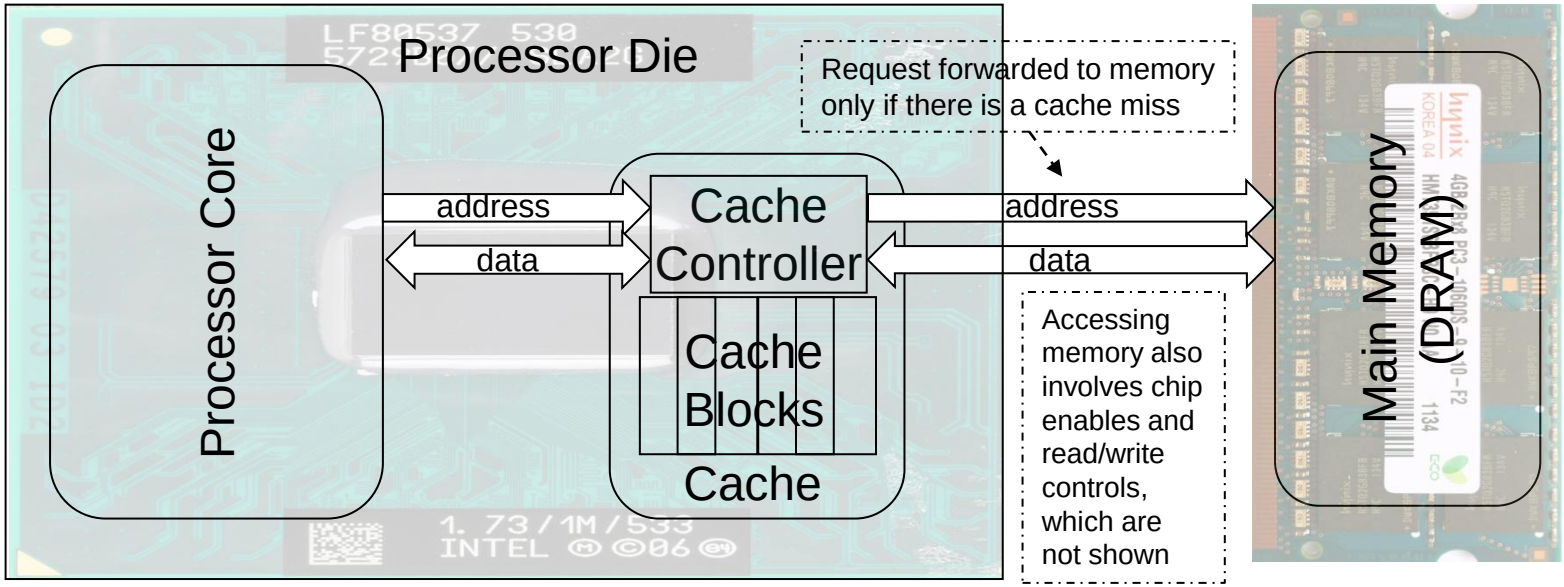

- When a block occupying cache is not referenced for a long time, it is pushed back to the MM to make space for another block
  - *Which* block to replace is decided by *replacement algorithms*
- Miss penalty: time taken to retrieve a block from slower level in the hierarchy

Read Misses

- Read miss
  - When a read miss happens, the block containing the word is loaded into the cache and then the desired word is sent to the CPU
- Load-through (*early restart*)
  - Alternatively, this word may be sent to the CPU as soon as it is read from the MM
  - Reduces CPU's waiting time, but additional circuitry needed
- Valid bit
  - If a location which is currently cached is modified in the main memory by an action which bypasses the CPU (eg : DMA), a *valid* bit for the corresponding cache block is cleared
  - The cache controller treats access to this location as a cache miss
  - Valid bits are set to 0 on power on!

DMA : A technique for moving data between memory and secondary storage / IO devices where the data transfer is managed by a separate hardware called DMA controller rather than through repeated LDR-STR by the processor

Handling Writes

- Write-through
  - In this case, the cache and MM locations are simultaneously updated
  - Simple, but results in unnecessary write operations in MM when cache is updated several times
- Write-back
  - Update only the cache location and mark it as updated with an associated flag bit, often called as *dirty* or modified bit
  - The MM word is updated later, when the block containing the word is removed from the cache by a replacement algorithm
  - May also lead to unnecessary write operations – when a cache block is written back to the memory, all the words of the block are written back, even if only a single word in that block was modified when it was in the cache

Mapping Techniques

- There are three different mapping techniques that are followed in practice
  - Direct mapping
  - Associative mapping
  - Set-Associative mapping

- The following example is used to illustrate the mapping algorithms
  - The cache consists of 128 blocks of 16 words each; a total of 2048 (2K) words
  - Assume that the MM is addressable by a 16-bit **word address** (not byte address, for simplicity)
  - MM has 64K words, which we will view as 4K blocks of 16 words each

Direct Mapping

- Direct mapping

  block j of MM -> block j modulo 128 of Cache
- MM blocks 0,128,256,… -> block 0 of cache

  MM blocks 1,129,257,… -> block 1 of cache, and so on

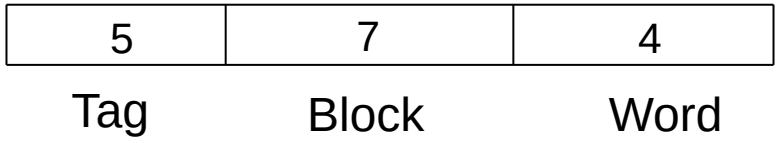

- Total of 16 bits
- lower order -> select a word within the block
- middle order -> block number in the cache
- high order -> which of the 32 blocks (4K/128 = 32 = 2^5) from MM is residing currently in the cache block

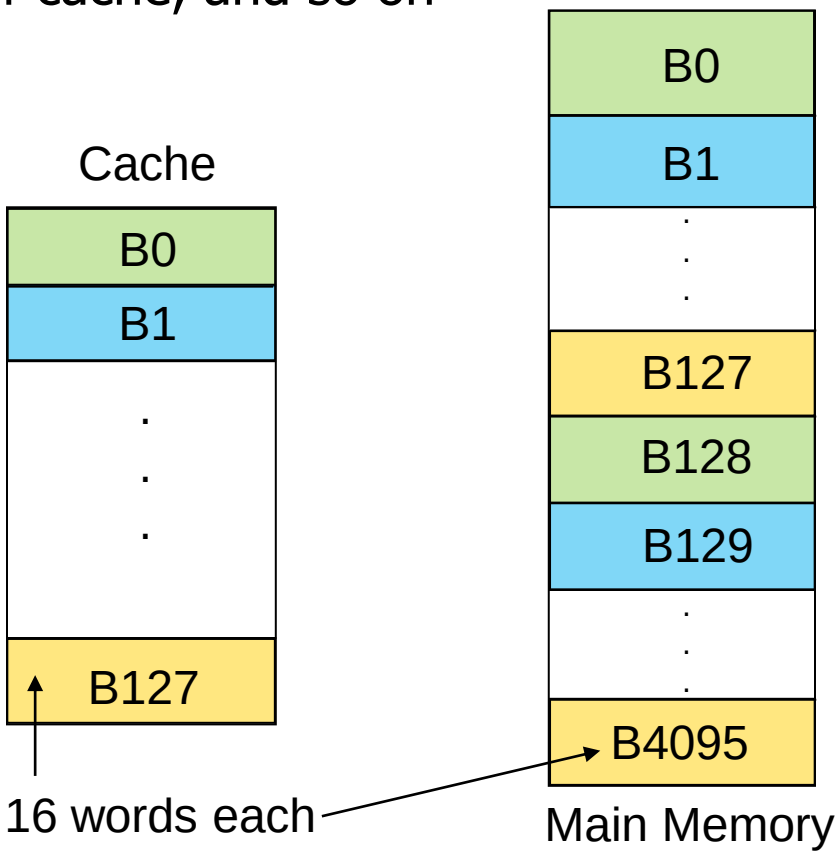

Direct Mapping …

- Note that the tag field in the above example is nothing but the higher order 5 bits of the word address
- These 5 bits are stored along with that block in the cache
- The tag field can be used to determine whether the block at this location is the required block – the tag field is unique for each block from MM which can be mapped to the same block in the cache
- Note that even when the cache is not full, contention may arise for a location
- In this case, the replacement algorithm is trivial (a main memory block is mapped to a unique cache block)

Direct Mapping ...

Note: Word access within the block is not shown

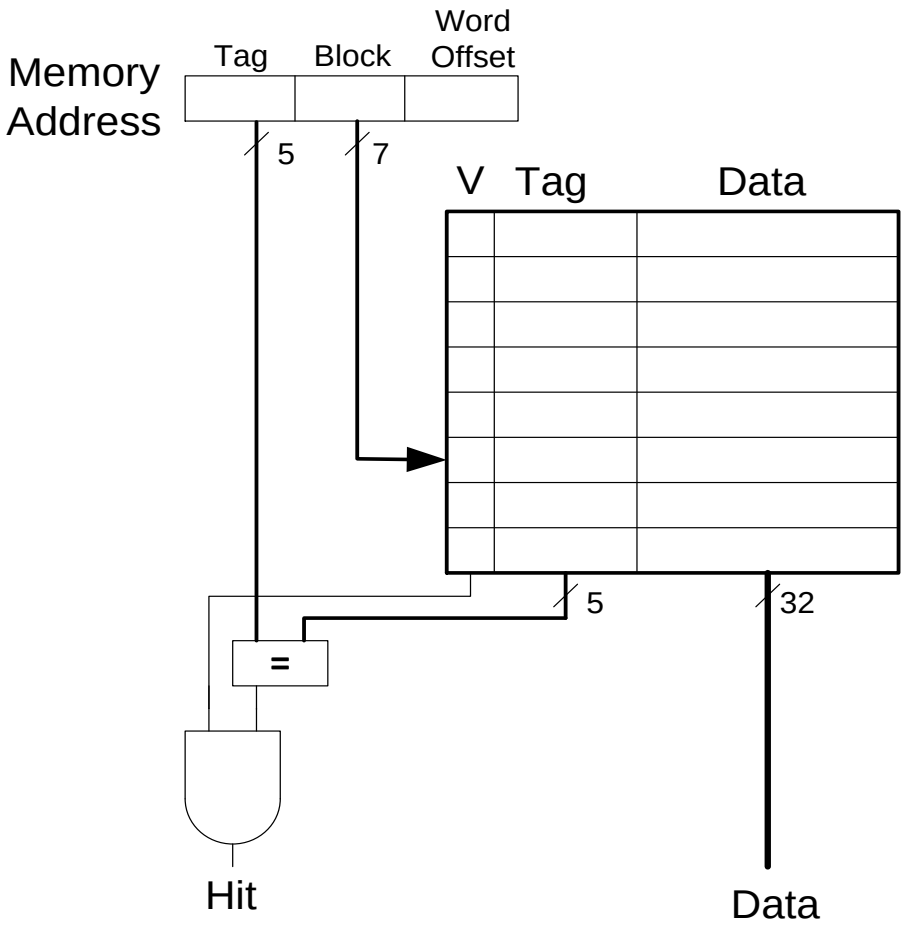

Associative Mapping

- (Fully) Associative mapping
- In this technique, a block of MM can be placed anywhere in the cache

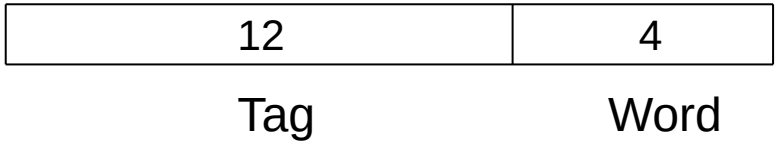

- From the CPU generated address, the higher order 12 bits are stored along with the block in the cache (which makes sense as each cache block can be from any of the 4096 = 2^{12} MM blocks)
- When the request arrives, the tag field is compared for all the blocks in the cache to see if there is a match

Associative Mapping

- This technique gives a complete freedom in choosing where in the cache a particular MM block is placed
  - Cache space is utilized more efficiently
- Disadvantage: Search 128 blocks to match for a single tag
  - This comparison has to be done for every memory access!
  - Parallel search schemes can be used
  - Still, costly and difficult to achieve high speeds
- The replacement follows one of the standard techniques such as LRU, FIFO, etc.

Set-Associative Mapping

- Set-Associative mapping
- This is a combination of / compromise between the previous techniques
- Here, blocks of cache are grouped into sets, and the mapping allows a block of the MM to reside in any block within a specific set (there is associativity within a set)
  - The contention problem of the direct method is eased by having a few choices for block placement
  - The hardware cost is reduced and speed is increased by decreasing the size of the associative search procedure
  - If there are N blocks per set, the memory is called N-way set associative

Set-Associative Mapping …

- Suppose if we allow two blocks per set in the cache. The memory blocks 0,64,128,…,4032 map into cache set 0, and they can occupy either of the two block positions within the set
- With 128 cache blocks and 2 blocks per set, we have 64 sets -> we need 6 bits to identify the right set and 4 bits for a word, leaving 6 bits for the Tag field (which makes sense as each cache block can be from any of the 4096/64 = 64 = 2^6 MM blocks)

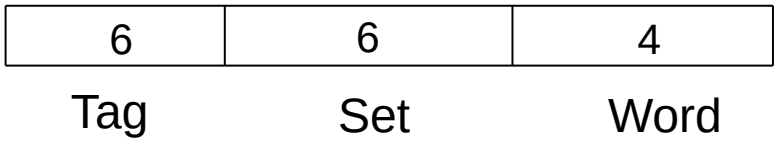

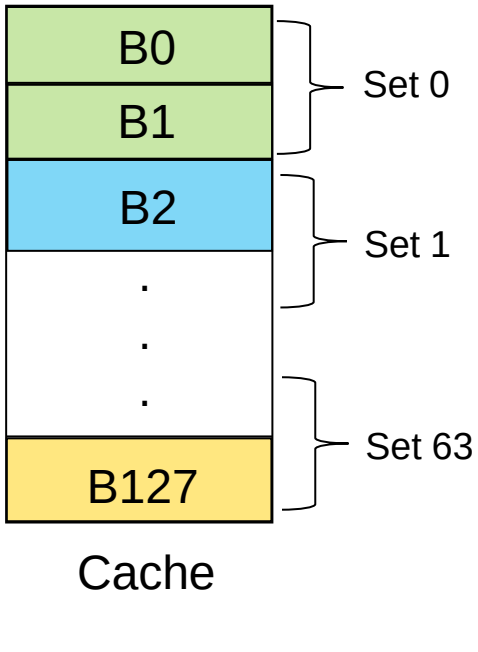

2-Way Set-Associative Mapping

Note: Word access within the block is not shown

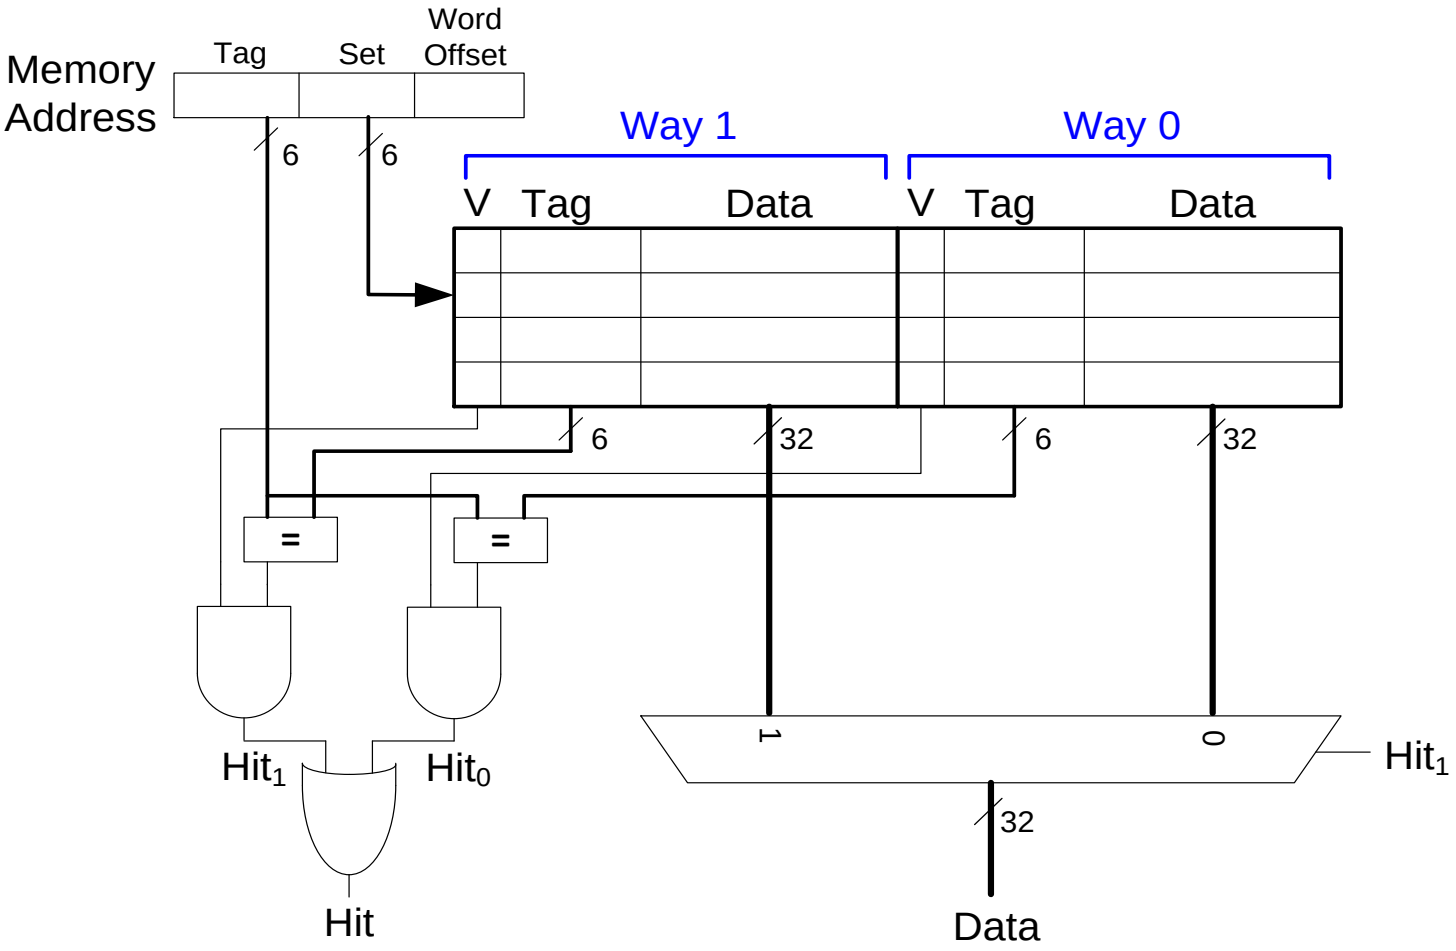

Replacement Algorithms

- In the case of associative and set-associative mapping, there is some flexibility in deciding which block should be thrown out if a new block is brought into the cache
  - Retain blocks that are likely to be referenced in the "near future"
  - Can't predict future - not easy to decide on "how long to hold a block"
- First-In-First-Out
  - Replace the oldest block in the memory
- Least Recently Used (LRU)
  - Replace the block that has not been referenced for a long time
- Optimal Algorithm: (Ideal - assumes knowing the future)
  - Replace the block that will not be used for a longest period of time
  - Cannot be implemented in practice, used only for analysis purpose

Replacement Algorithm : FIFO

- In the example, assume that the cache is fully associative and has 4 blocks
- FIFO works well if the access follows a sequential pattern (arrays etc.)

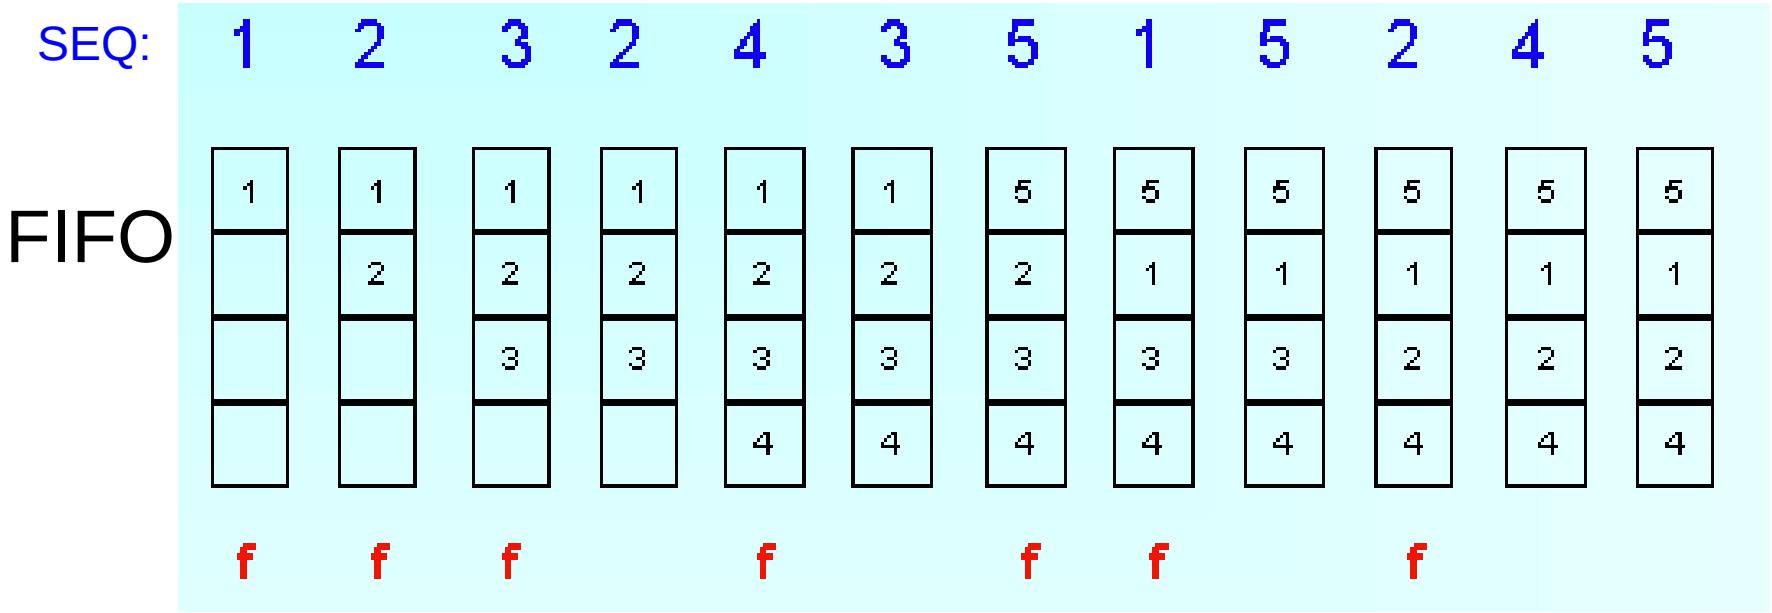

f: miss

Replacement Algorithm : LRU

- It makes sense to overwrite a block that resided in the cache for a long time (LRU block) without being referenced
  - Temporal locality of reference

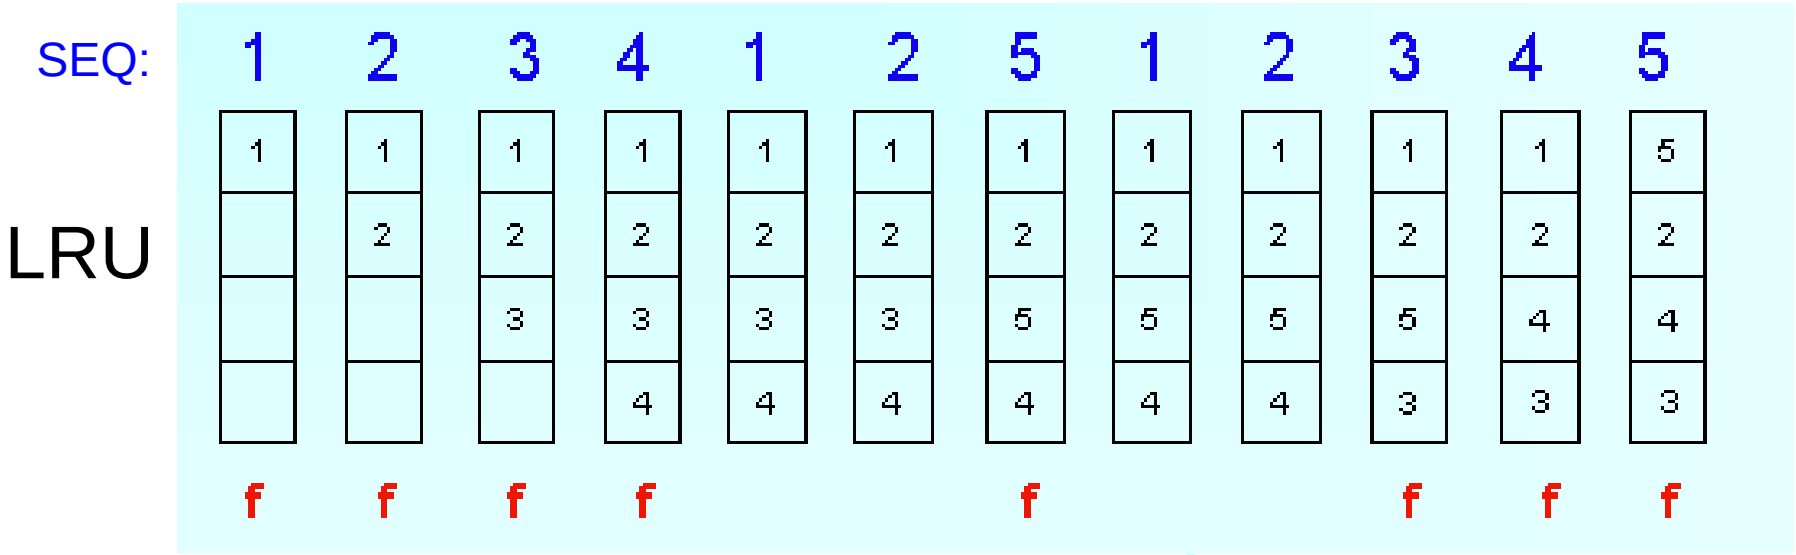

Replacement Algorithm : Optimal

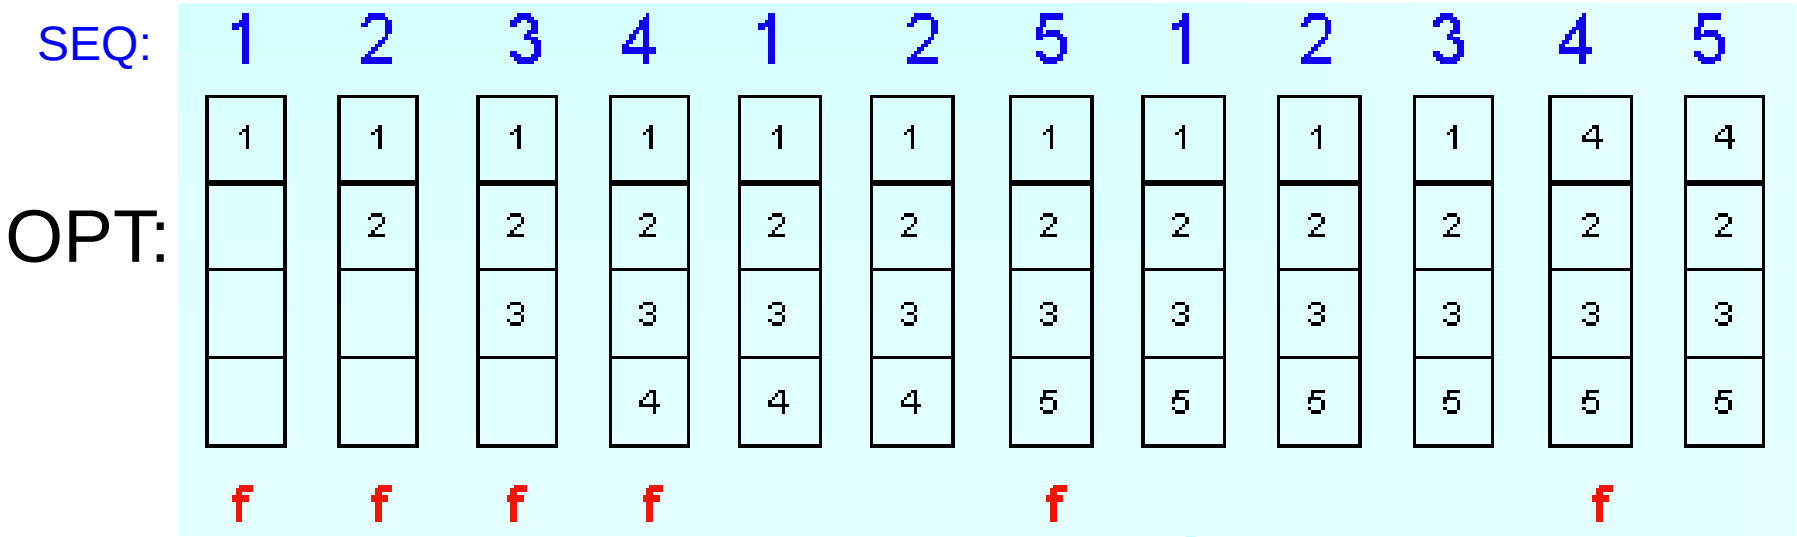

Multilevel Caches

- Larger caches have lower miss rates, but longer access times
- Expand memory hierarchy to multiple levels of caches
- Most modern PCs have L1, L2, and L3 cache
  - Intel Coffee Lake has
    - Level 1: small and fast, 32 KB, 5-6 cycles
    - Level 2: larger and slower, 256 KB, 12 cycles
    - Level 3: 2MB x number of cores, shared among cores, 42 cycles
    - Level 4: 128MB (only in some models) per package
    - In contrast, main memory has a 42 cycles + 51 ns latency (about 200 cycles in total)

Types of Cache Misses

- Compulsory: first time the data is accessed
- Capacity: cache too small to hold all data of interest (the data of interest for a particular process is called its *working set*), causing some blocks to be evicted from the cache that are required later
- Conflict: data of interest maps to same location in cache. These are misses that would not occur if the cache were fully associative with LRU replacement
- Coherence: data is changed in the main memory due to actions that bypass the cache such as DMA or another processor writing to the main memory / shared higher level cache

Types of Cache Misses : Example

- A direct mapped cache has 4 blocks. The access sequence is 0, 1, 2, 3, 4, 1, 2, 3, 0, 4, 0
- 0 (compulsory miss), 1 (compulsory miss), 2 (compulsory miss), 3 (compulsory miss), 4 (compulsory miss), 1 (hit), 2 (hit), 3 (hit), 0 (capacity miss*), 4 (capacity miss), 0 (conflict miss$)
- *A capacity miss because even if the cache were fully associative with LRU replacement, there will still be a miss because 4, 1, 2, 3 are the recently accessed blocks
- $A conflict miss because in a fully associative cache, the last 4 would have replaced 1 in the cache instead of 0

Example courtesy: https://stackoverflow.com/questions/33314115/whats-the-difference-between-conflict-miss-and-capacity-miss

Miss Rate Trends

- Bigger caches reduce capacity misses (obviously)
- Greater associativity reduces conflict misses
  - But with diminishing returns

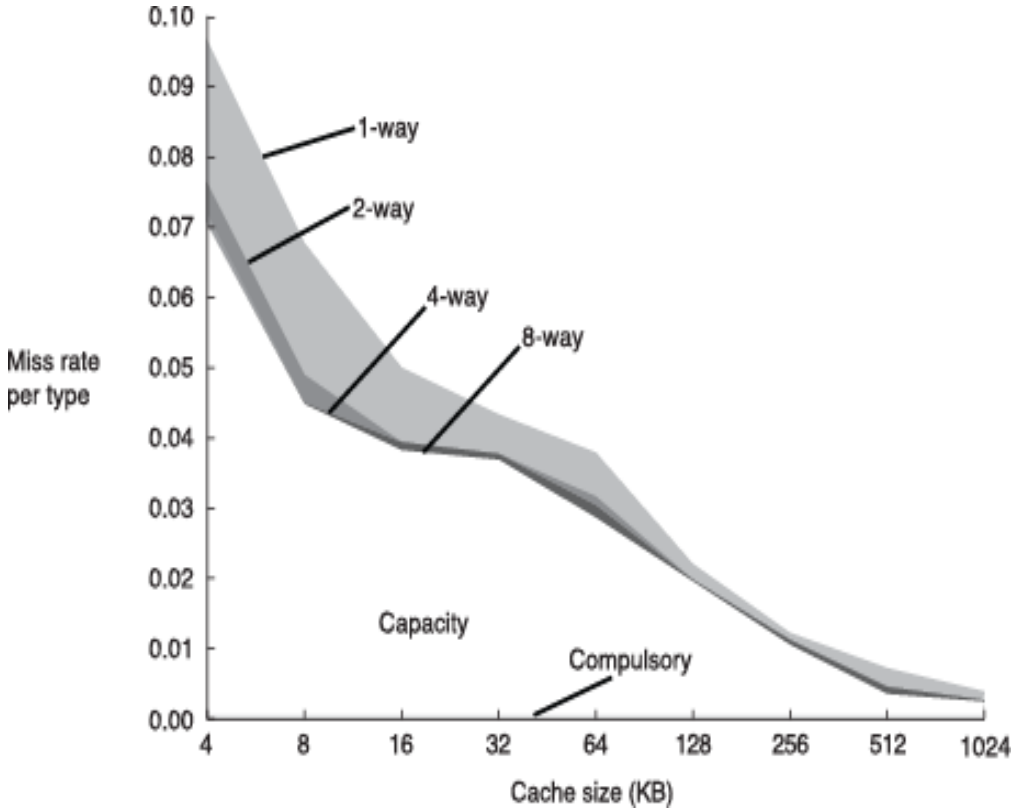

Miss Rate Trends …

- Bigger blocks reduce compulsory misses
  - Takes better advantage of spatial locality
- Bigger blocks increase conflict misses
  - Lower number of sets for a given associativity, or lower degree of associativity for a given number of sets
- Bigger blocks incur higher miss penalty as well

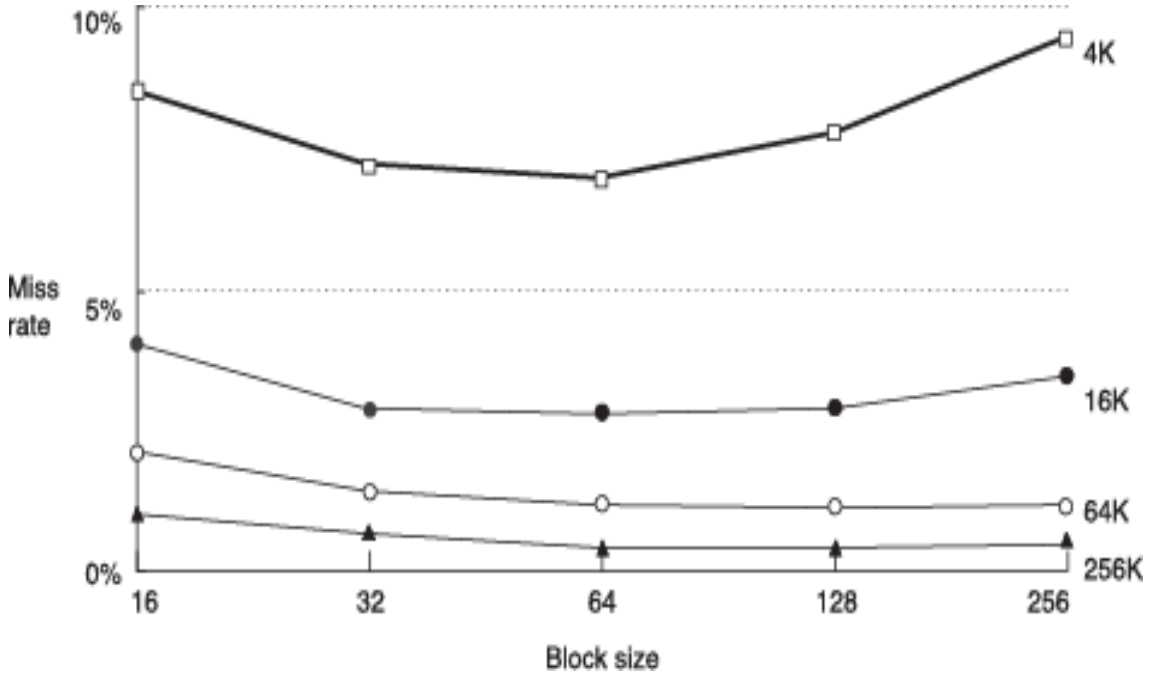

Virtual Memory Technology

- Though the modern day computers have a lot of MM (usually at least 2GB), the amount of applications that run demand a large working space
- Also, usually, the physical MM is not as large as the address space of the processor
- When a program is to be executed, it has to be brought into the MM (DRAM)
  - The OS controls the movement of data/program between the MM and the secondary storage devices
- Virtual memory gives the illusion of bigger memory*
- MM acts as cache for secondary storage (HDD/SSD)

*as opposed to cache, which gives the illusion a faster memory

Virtual Memory Technology …

- Virtual addresses
  - Programs use virtual addresses
  - Entire virtual address space stored on a secondary storage
  - Subset of it in MM
  - Memory Management Unit (MMU) inside the CPU translates virtual addresses into physical addresses (MM addresses)
  - Data not in MM fetched from secondary storage
- Memory Protection
  - Each program has own virtual to physical mapping
    - Called a Page Table (PT), managed by the OS
  - Two programs can use same virtual address for different data
  - Programs don’t need to be aware of other programs running
  - One program (or virus) can’t corrupt memory used by another

Virtual Memory Technology ...

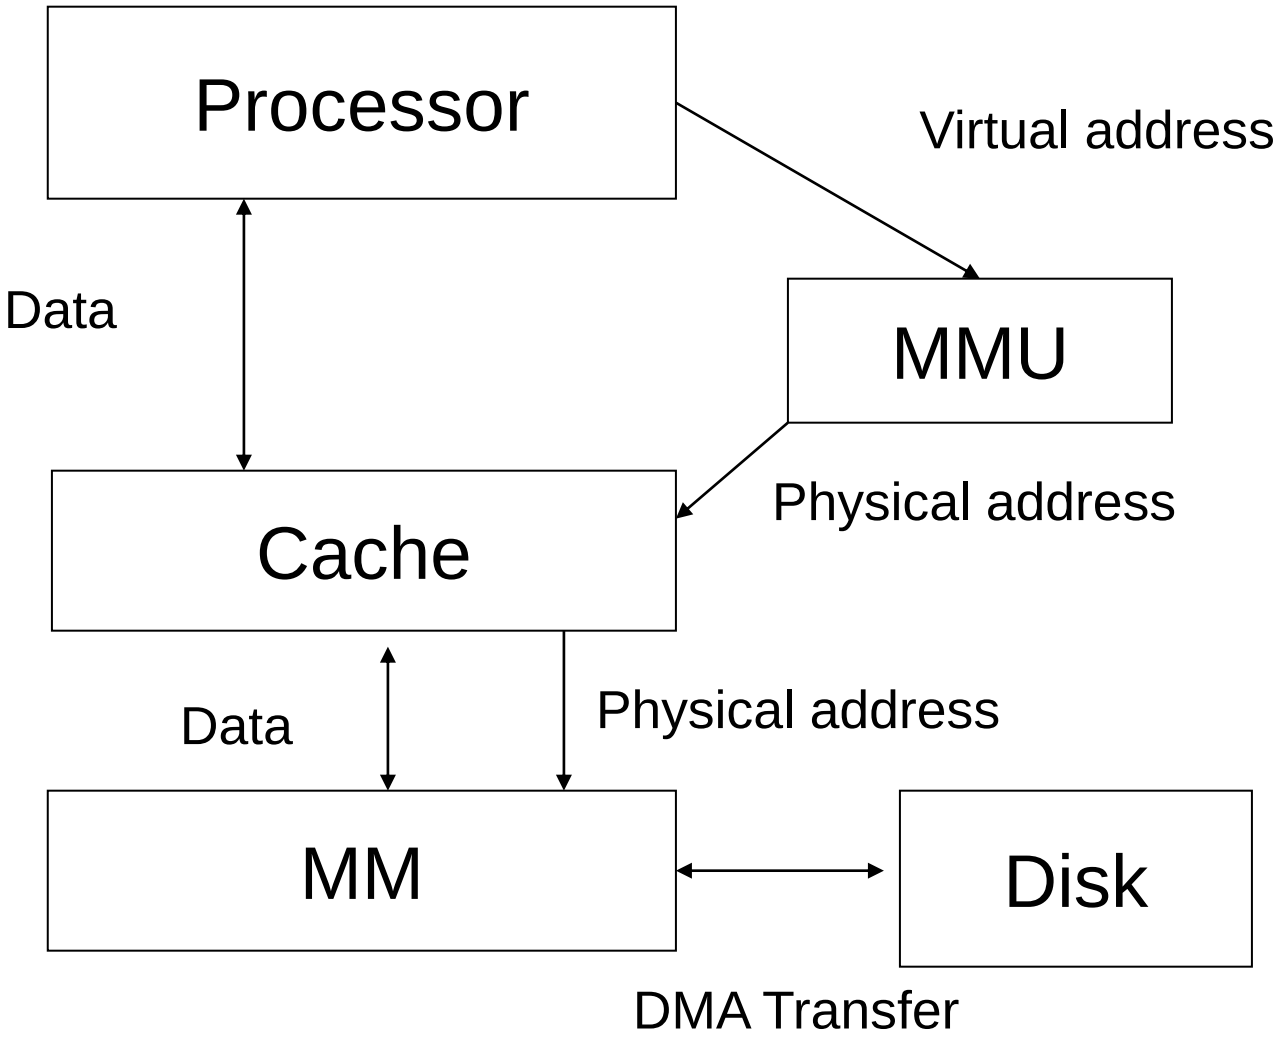

Address Translation

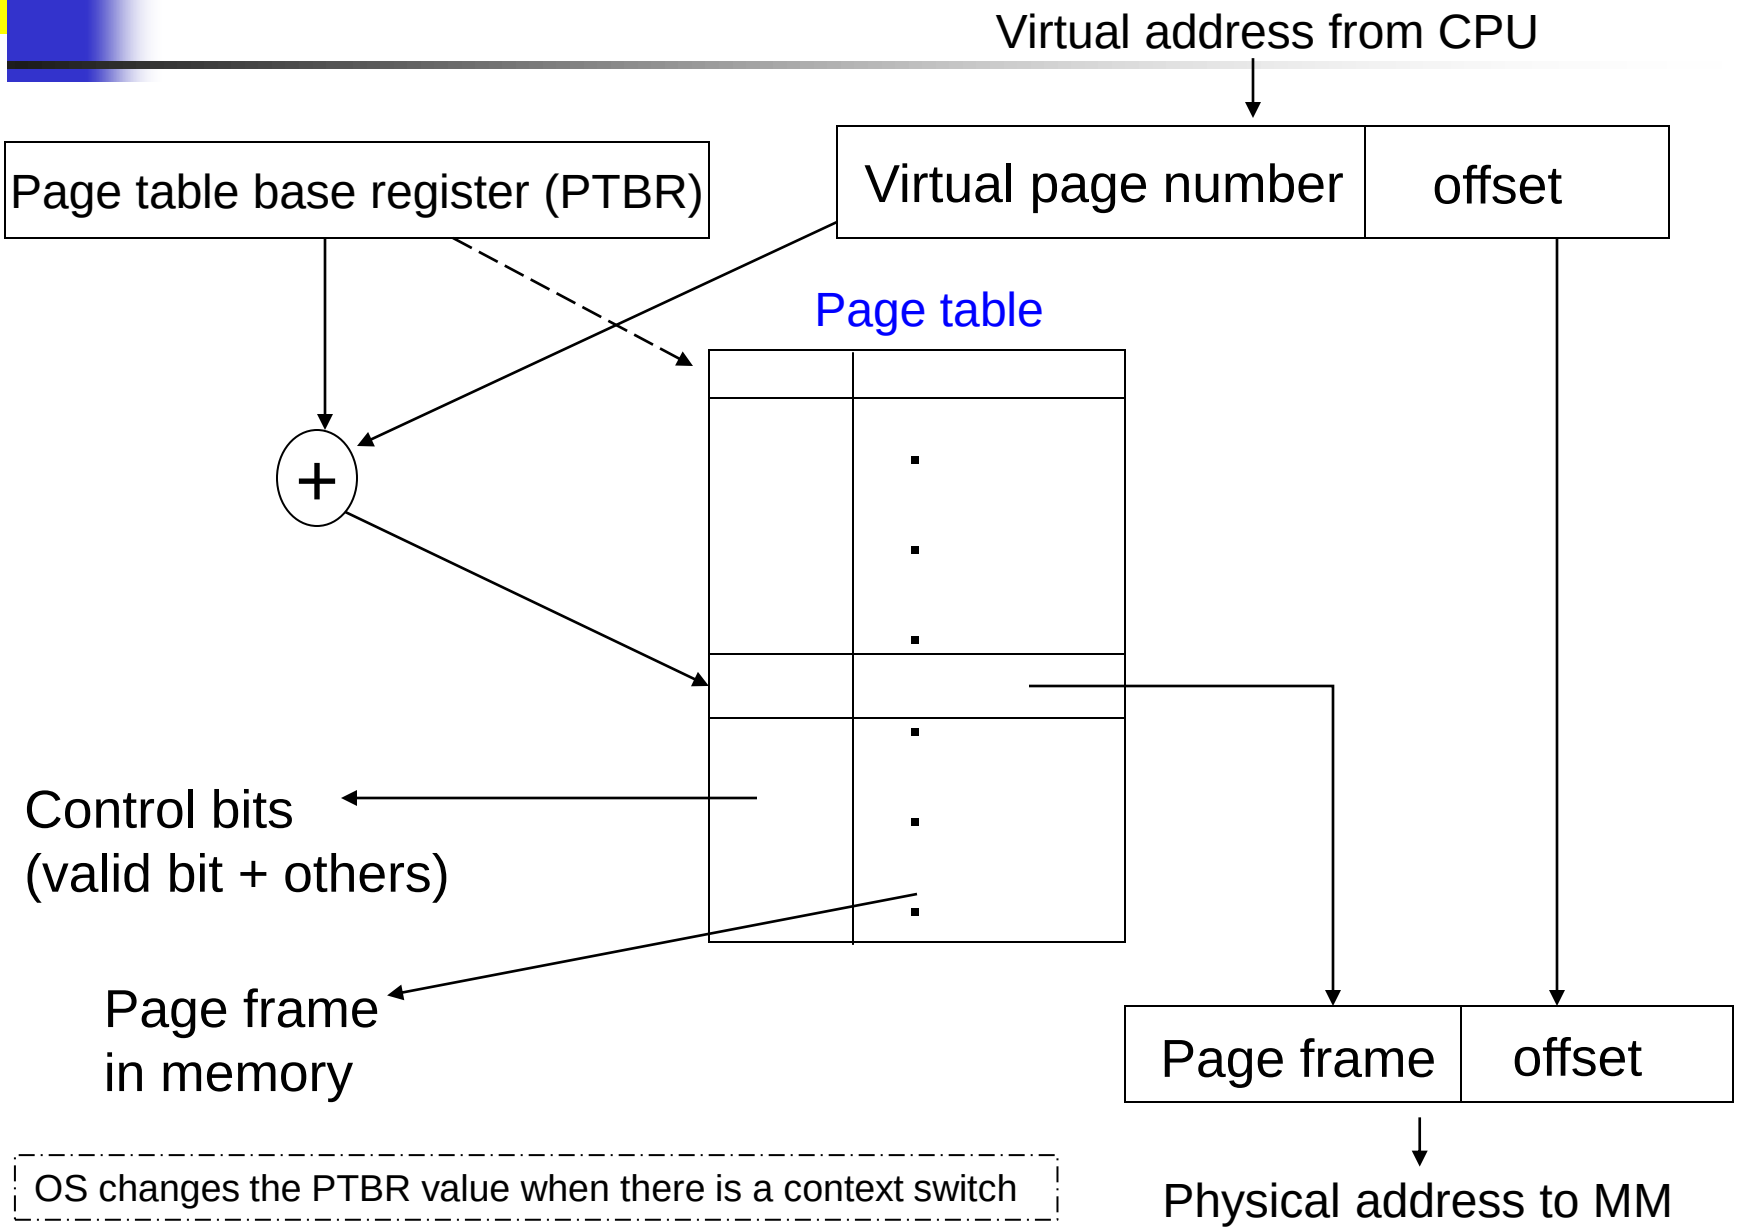

OS changes the PTBR value when there is a context switch

Translation Look-aside Buffer

- For every page that is in use, PT has an entry
  - If the required page is not present in the main memory, it causes a *page fault* exception
  - This requires OS intervention, either to initiate a read from secondary storage (valid virtual address), or program termination (invalid address)
- PT is usually big, and is kept in the MM
- For every memory access (LDR/STR) with a cache miss, we need to access the MM twice
  - One for reading the page table
  - One more for the actual data access (using the translated address)
- To minimize it, a small portion of active PT entries are kept in a tiny cache, referred to as Translation Look-aside Buffer (TLB)
  - Page size is large (2-4KB), so consecutive loads/stores are likely to access same page. >99% hit rates typical
  - Small: accessed in < 1 cycle, typically 16 - 512 entries
  - Since TLB is small, as fully-associative implementation is common

Translation Look-aside Buffer ...

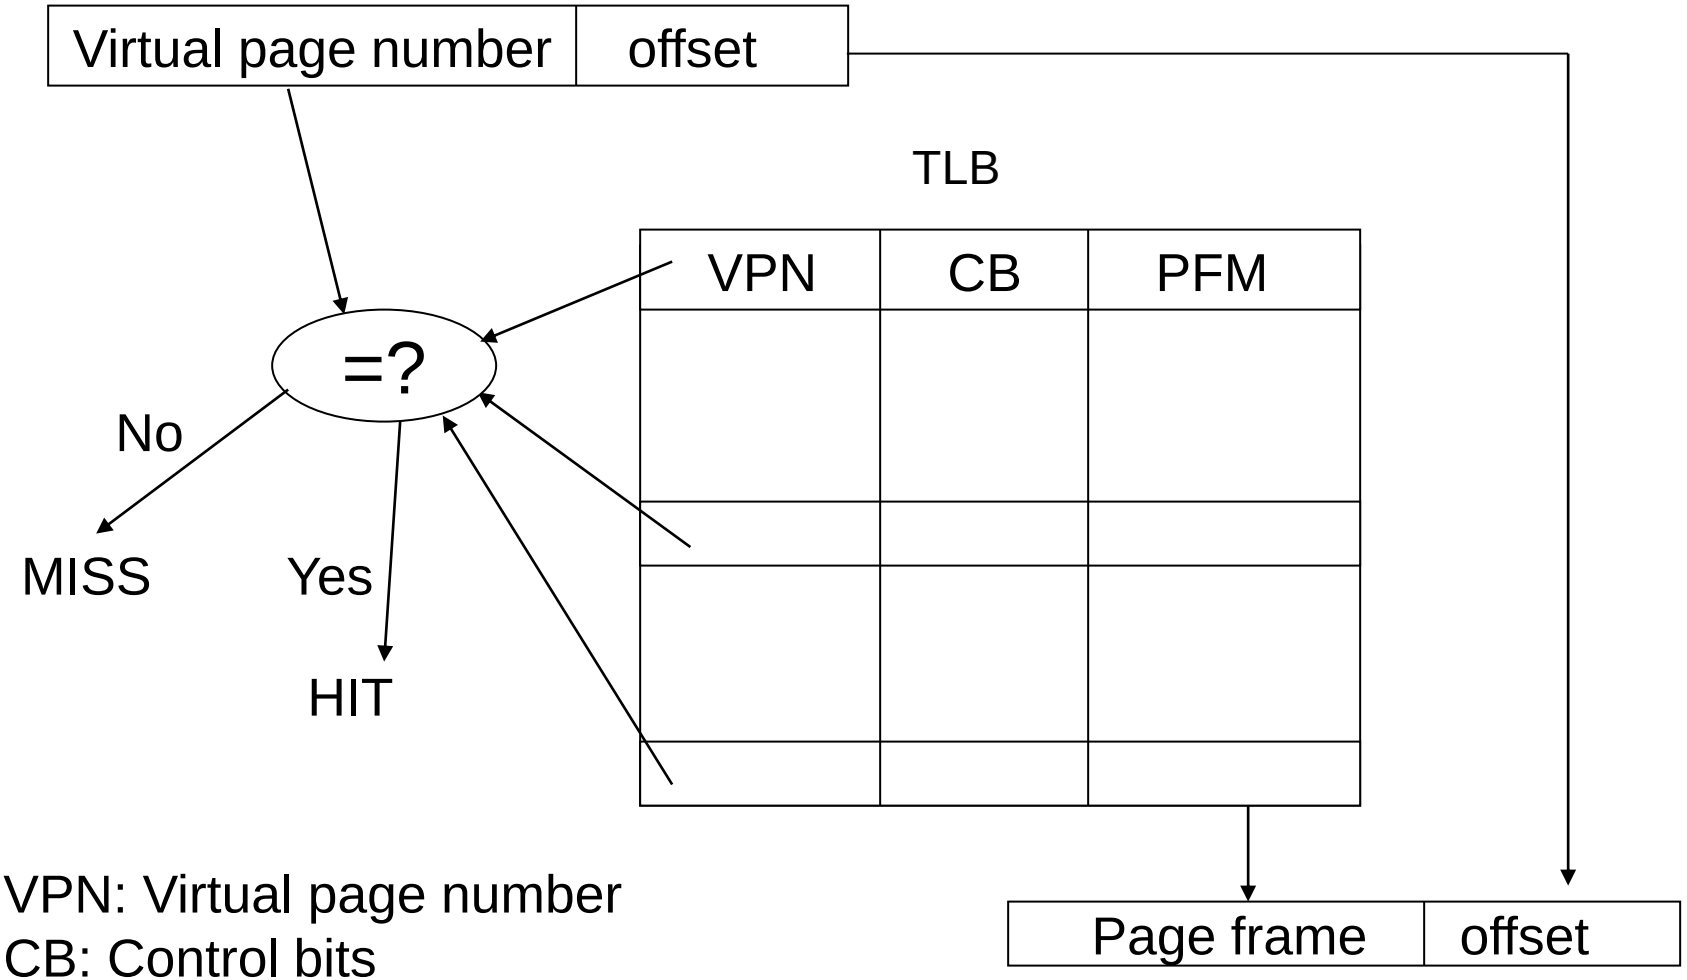

VPN: Virtual page number
CB: Control bits
PFM: Page frame in memory

Summary

- Cache – basic principles
- Hit rate and effective access times
- Handling read and write misses
- Mapping techniques
- Replacement algorithms
- Types of misses and block size considerations
- Virtual memory